

JUE Insight: Attention and the Economic Response to Immigration Enforcement

Greg C. Wright

University of California, Merced

gwright4@ucmerced.edu

September 1, 2026

Abstract

Does the community response to immigration enforcement depend on the scale of the operation or on the amount of public attention that it receives? I study U.S. Immigration and Customs Enforcement events from 2018 to 2026 and compare the community response in heavily Latino-immigrant neighborhoods to other nearby neighborhoods. More Spanish-language news coverage produces larger foot-traffic declines. An instrumental-variables design using competing news shocks from foreign disasters implies that one standard deviation of coverage reduces foot traffic by 1.6 percent. Arrests explain little and do not attenuate the coverage effect. The response is driven by national, rather than local, attention.

Keywords: public attention, immigration enforcement, chilling effects, foot traffic.

JEL Codes: J15, R23, K37.

1 Introduction

Immigration enforcement can affect the lives of individuals beyond those who are detained. In part, this is because the community response may depend as much on the information environment surrounding an event as on the event itself. A workplace raid that goes largely unnoticed may have little effect beyond the worksite, whereas a smaller raid that becomes highly visible may lead residents to avoid public places. A long literature in economics and sociology documents these “chilling effects,” which can lead to reduced use of healthcare, schooling, and public benefits due to fear of detention ([Watson, 2014](#); [Asad, 2023](#)). The intensification of U.S. Immigration and Customs Enforcement (ICE) activity under the second Trump administration provides a real-time test. For example, [Lester et al. \(2026\)](#) document an 8–10 percent foot-traffic decline and a 20–25 percent spending decline in the most heavily Latin American-born neighborhoods of Los Angeles–Orange County in the weeks following ICE’s May 14, 2025 announcement of a major regional enforcement operation.

This paper asks whether a localized enforcement action becomes a neighborhood-wide economic shock because of the operation itself or because of the public attention it receives. If the scale of an operation is central to its effect, arrest counts should help explain which events produce larger responses. If the broader *perception* of risk is what matters, the response should depend on whether news of the event reaches people who were not directly targeted. A raid affects both margins at once. It detains a limited number of individuals while also generating news coverage, social-media activity, rumor, and protest.

I study 51 ICE enforcement events from April 2018 through January 2026 across 26 states, spanning a small cluster of 2018–2019 worksite actions and a larger cluster of 2025–2026 operations. The main specification compares Census Block Groups (CBGs) in the top decile of Latin-American foreign-born (FB) share with CBGs in the bottom decile in the same event catchment and calendar weeks. In a single stacked difference-in-differences (DiD), I ask whether the relative post-event decline in high-FB neighborhoods is larger when the event receives more Spanish-language news coverage.

The main result is that the foot-traffic response tracks public attention rather than operational scale. I exploit post-event spikes in the share of Spanish-language news stories mentioning “ICE,” the agency name most salient to the population at risk. Events that receive more Spanish-language news coverage generate larger declines in foot traffic in high-FB neighborhoods. In the stacked late-window specification, one standard deviation of additional coverage predicts about a 1.5 percent larger decline in weeks 4–8 after the event. English-language news and Google search data show the same pattern.

I adopt an instrumental variables design that exploits news of major foreign disasters, which crowds out public attention to enforcement actions, as an instrument for the news coverage spike, with out-of-state U.S. disasters providing a complementary version. These news shocks strongly predict the amount of coverage given to enforcement events and occur independently of the operation’s timing. The stacked IV estimate implies that one standard deviation of additional news coverage reduces foot traffic by about 1.6 percent.

Several checks sharpen the interpretation. The effect of news coverage on foot traffic is absent in every pre-event week and emerges only after the event. The early foot-traffic response also does not predict later news coverage. The result survives controls for pre-event news coverage, operation type, and sector, and it is similar among events with no overlapping catchments and flat pre-trends. It is also not an artifact of pooling the 2018–2019 and 2025–2026 enforcement waves. The effect holds, and is if anything stronger, when the sample is restricted to the 2025–2026 wave alone.

Arrest counts do not explain the same pattern. Raw arrests, log arrests, arrests per catchment population, and arrests per high-FB population explain little of the response when entered alone. Adding these measures does not attenuate the effect of news coverage. This does not show that operation size has no causal effect, since arrests are not instrumented and may be measured with error. It does show that reported arrests do not serve empirically as a meaningful proxy for the community-relevant importance of the event.

National Google Trends and national Spanish-language news both predict the local foot-

traffic decline, while metro-level search and local Spanish-language news add little once national attention is accounted for. The response is broad, extending beyond the immediate worksite and the raided industry. National information therefore enlarges the local economic footprint of an otherwise geographically concentrated action. I decompose the response by venue and find that foot traffic around schools, parks, and museums falls more than at commercial sites such as retail and food services, a pattern more consistent with a broad withdrawal from public life than a localized demand shock.¹ The economic footprint of visible enforcement is therefore not limited to lost commercial activity but includes reduced participation in civic life.

A broader economics literature uses interior-enforcement intensity and federal programs, especially Secure Communities, to estimate effects of immigration enforcement on safety-net participation, family decisions, childcare, healthcare, and labor markets (Watson, 2014; Amuedo-Dorantes et al., 2018, 2020; Alsan and Yang, 2024; Ali et al., 2024; Herring and Barnow, 2025; East et al., 2023). A related sociology literature describes the institutional-avoidance strategies of mixed-status families (Asad, 2023). Several related papers study enforcement visibility and the recent enforcement wave. Amuedo-Dorantes and Antman (2022) link removals and raid awareness to lower immigrant labor-market engagement, Ciancio and García-Jimeno (2026) study consumption responses to learned enforcement risk, and De Balanzó et al. (2026), Hernandez (2026), and Cox and East (2026) estimate economic effects of the 2025 wave using spending, mobility, and labor-market data. These papers show that enforcement moves economic activity when risk becomes visible or learned. Two features of these designs nevertheless leave the role of attention, separate from enforcement, unmeasured. Variation built on program rollouts or enforcement intensity bundles any change in attention together with the change in enforcement itself, and single-episode studies hold the event fixed but have no variation in event visibility. I instead condition on the occurrence of a specific enforcement event and ask whether the neighborhood response is larger when

¹Religious organizations and hospitals are effectively absent from the usable POI data. See Section 4.7.

that event receives more news coverage.

These results also add to evidence on how public attention, more than operational scale, shapes the response to risk. A related COVID-mobility literature ([Goolsbee and Syver-son, 2021](#); [Baker et al., 2020](#); [Chetty et al., 2024](#)) shows that perceptions of risk changed individual behavior independently of formal COVID mandates. The case here is similar. An immigration enforcement operation produces both a worksite disruption and public attention, but the community response tracks the attention. The design shows how national information can propagate a localized government action across urban space. The same logic may extend to other salient but partly hidden risks whose economic footprint depends on how widely they are noticed. Where prior work instruments or exploits quasi-experimental variation in enforcement itself, I condition on the timing and location of the operation and instrument the news coverage it receives. This adapts the competing-news logic of [Eisensee and Strömberg \(2007\)](#) to isolate the attention component of a specific event and, to my knowledge, is new to the immigration literature.

2 Data

The analysis combines an ICE event inventory with cellphone foot traffic, SafeGraph consumer spending, American Community Survey (ACS) demographics, Google Trends search data, Spanish- and English-language news-volume series, disaster alerts, and Nielsen Designated Market Area (DMA) definitions.

2.1 Event inventory

I compile single-worksite ICE enforcement events from news sources, agency press releases, litigation records, immigration-policy trackers, and prior compilations. The inventory covers April 2018 through August 2019 and January 2025 through January 2026, the two recent windows in which worksite enforcement was especially salient. An event enters if it has

a specific worksite or commercial address, falls in one of the 26 states for which CBG-level Latin-American foreign-born data are available, and can be geocoded to the named address or, when necessary, to a city or county centroid.² I exclude multi-day or metro-wide operations, actions without a single worksite date, and actions that were not ICE-led. These screens identify 55 single-worksite ICE-led events. Four have insufficient point-of-interest (POI) coverage within the 50 km event catchment to contribute the high-FB versus low-FB comparison, leaving a main sample of 51 events in 26 states. The Online Appendix reports the sample flow, event list, and summary statistics by enforcement wave.

A publicly compiled inventory could be biased toward visible raids. No complete administrative list of single-worksite ICE actions is publicly available for these windows, so the inventory is built from several sources rather than national news alone. The final list includes many low-visibility events. Of the 55 events, 32 are coded medium or low publicity and 29 have an English-GDELT news spike at or below 0.01. These facts do not establish a complete universe of obscure raids, but they make it unlikely that the result is mechanically generated by retaining only highly visible events.

2.2 Outcomes and exposure

Foot traffic comes from Advan Research’s *Foot Traffic/Weekly Patterns Plus* dataset ([Advan Research, 2025](#)), derived from anonymized opt-in mobile devices. For each event, I collect Advan POIs within 50 km of the worksite anchor and assemble a POI $\times$ week panel spanning ± 12 weeks. POI visits are aggregated to CBG-week totals, and the sample is restricted to CBGs with at least 10 sufficiently observed POIs. The resulting panel contains about 69 million POI-week observations and 3.3 million POI–event pairs.

Spending is measured with SafeGraph’s *SpendPatterns* dataset ([SafeGraph, 2022](#)), which reports monthly POI-level records containing arrays of daily consumer spending from an opt-in financial-institution panel. I aggregate the daily values to weeks and then to the same

²Five events are geocoded to city centroids and one to a county centroid.

CBG-week catchments used for foot traffic.

Treatment is defined using the 2018–2022 5-year ACS. CBGs in the top decile of Latin-American foreign-born share within an event catchment are the high-exposure group, and bottom-decile CBGs are the comparison group. Because deciles are computed within each catchment, the design compares the most- and least-Latino-immigrant neighborhoods in the same local area rather than imposing a national cutoff.

2.3 Attention and news shocks

The headline attention measure is the post-event spike in Spanish-language news coverage of ICE. I query the Mediacloud Online News Archive ([Media Cloud, 2025](#)) for daily story counts in its curated 407-source US Spanish-language collection and normalize by all stories in the collection. The headline query is the noun “ICE.” I also construct corroborating measures using alternative Spanish terms, a state-local Spanish-news series, GDELT 2.0 English-language US news volume for “ICE raid” ([Leetaru and Schrod, 2013](#)), and Google Trends search interest for “ICE raid”, “redada”, and “deportación”. Google Trends measures are used for the 2025–2026 event wave. Their underlying search window begins in 2024, but cross-era rescaling is noisy for these low-baseline terms.

The preferred news-shock instrument comes from the Global Disaster Alert and Coordination System event feed ([Global Disaster Alert and Coordination System, 2026](#)). It counts days in the post-event window that overlap with the first seven days of a red-alert sudden disaster outside the United States, Latin America, and the Caribbean. The Latin America and Caribbean exclusion removes shocks that could affect Latino immigrant communities directly through homeland ties or Spanish-language news demand. A complementary instrument counts major U.S. natural disasters in states other than the ICE event state. The two instruments identify different margins of news-cycle access, so the IV estimates are interpreted as local effects for events whose coverage is moved by the corresponding source of competing news shocks.

3 Empirical Strategy

3.1 Stacked event-time design

The main specification stacks the CBG-week panels for all events. I use two versions of the same design. The first shows how the relationship between news coverage and foot traffic evolves in event time. Let $y_{c,t}^e$ denote weekly foot traffic in CBG c , calendar week t , and event catchment e . The event-time specification is

$$\begin{aligned} \log(1 + y_{c,t}^e) = & \alpha_c^e + \lambda_t^e + \sum_{\tau \neq -1} \gamma_\tau (\text{HighFB}_c^e \times \mathbf{1}[t - T_0^e = \tau]) \\ & + \sum_{\tau \neq -1} \delta_\tau (\text{HighFB}_c^e \times \mathbf{1}[t - T_0^e = \tau] \times \text{Spike}_e^z) + \varepsilon_{c,t}^e, \end{aligned} \quad (1)$$

where α_c^e are event $\times$ CBG fixed effects and λ_t^e are event $\times$ calendar-week fixed effects. The dummy HighFB_c^e equals one for CBGs in the top decile of Latin-American foreign-born share within the event catchment and zero for CBGs in the bottom decile. The middle deciles are omitted. Spike_e^z is the standardized post-event Spanish-language news-coverage spike. The coefficients γ_τ absorb the average high-FB versus low-FB path, while δ_τ show how that path changes with news coverage. Figure 1 plots δ_τ .

The second version is the estimating equation for the main table. It collapses the event-time path to the late window, weeks 4–8 after the event, and separates the effect of news coverage from the effect of reported arrests:

$$\begin{aligned} \log(1 + y_{c,t}^e) = & \alpha_c^e + \lambda_t^e + \theta_0 (\text{HighFB}_c^e \times \text{Late}_t^e) \\ & + \theta_S (\text{HighFB}_c^e \times \text{Late}_t^e \times \text{Spike}_e^z) \\ & + \theta_A (\text{HighFB}_c^e \times \text{Late}_t^e \times \text{LogArrests}_e^z) + \varepsilon_{c,t}^e, \end{aligned} \quad (2)$$

where Late_t^e equals one in weeks 4–8 after the event and zero in the remaining weeks of the event window, and LogArrests_e^z is standardized $\log(1 + \text{Arrests}_e)$. The coefficient of interest

is θ_S . It is the additional high-FB foot-traffic response associated with a one-standard-deviation larger news-coverage spike. Let Z_e denote foreign-disaster news pressure. In the IV version, the coverage interaction is instrumented with $\text{HighFB}_c^e \times \text{Late}_t^e \times Z_e$.

3.2 Identifying variation

The design uses within-event variation between high-FB and low-FB CBGs in the same catchment. Event $\times$ CBG fixed effects absorb time-invariant differences in population, retail density, and baseline foot traffic. Event $\times$ calendar-week fixed effects absorb shocks common to all CBGs in a catchment in a given week. The stacked coefficients are therefore differential foot-traffic responses of high-FB versus low-FB CBGs in the same metro and calendar weeks.

News coverage and arrests vary across events. The coefficient θ_S is identified by whether the within-event high-FB versus low-FB change is systematically larger in catchments associated with more heavily covered events. This relationship is estimated directly in one stacked CBG-week regression, with standard errors clustered by event.

The design does not require raids to be unanticipated, and some initiatives are clearly partially anticipated. The initial question is whether the news coverage generated by each event explains why otherwise similar raids produce different responses. Section 4.3 implements an instrumental-variables design to isolate plausibly exogenous variation in coverage from the surrounding news cycle, and the Online Appendix adds pre-event coverage as a control.

3.3 News-coverage and event severity measures

To distinguish news coverage from operational scale, I interact the high-FB late-window indicator with two measures that vary across events. The headline news-coverage measure is the post-event spike in the Mediacloud US Spanish-language news-volume series for the query “ICE” (Section 2.3), measured as the share of stories in the 407-source US Spanish-language collection that mention the term. The spike is computed as the peak over $\tau \in [0, 8]$

minus the baseline over $\tau \in [-8, -5]$.

The headline measure is complemented with three corroborating measures: the analogous GDELT national US English-language news-volume spike for “ICE raid,” the national Google Trends spike for “ICE raid,” and the DMA-level Google Trends spike for the local media market in which the event occurred. The two Trends measures are available for the 2025–2026 event wave and are defined as within-geography pre-vs-post changes, since Google Trends rescales each query within a geography and timeframe.

To capture the severity of the enforcement event, the main specification uses the reported worksite arrest count in log form, $\log(1 + \text{Arrests}_e)$, to accommodate the right-skewed distribution.³ The Online Appendix reports specifications using raw arrests and arrests scaled by the catchment and high-FB populations. Arrest counts are available for 49 of the 51 analysis events.⁴

3.4 News-shock instruments

The OLS estimates treat realized post-event news coverage as the explanatory variable. I use competing-news shocks that occur independently of enforcement timing as instruments for coverage generated by the surrounding news cycle. The preferred instrument uses major foreign disaster news shocks. It counts days in the post-event window that overlap with the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. The disasters are contemporaneous with post-event coverage but unrelated to the timing or scale of the enforcement operation. The exclusion of Latin America and the Caribbean is intended to remove shocks that could directly affect Latino immigrant communities through homeland ties rather than through news-cycle competition only. The identifying assumption is that these shocks move news coverage of the enforcement action without directly changing the relative foot-traffic response in high-FB versus low-FB

³No sampled event has zero reported arrests (the minimum is 2), so using $\log(\text{Arrests}_e)$ instead leaves the results unchanged.

⁴The two events with missing arrest counts in the inventory are Newark NJ (an unnamed business) and Baldwin County AL (an unnamed construction site).

neighborhoods. Section 4.3 presents the IV estimates. The Online Appendix reports the corresponding balance and robustness checks.

Out-of-state U.S. disasters also serve as a complementary instrument. This measure captures domestic disaster news shocks in states other than the ICE event state. The foreign- and U.S.-disaster instruments identify different local margins of news-cycle access. The IV estimates should therefore be read as local effects for events whose public attention is shifted by a particular source of competing news shocks.

4 Results


The relationship between news coverage and foot traffic is shown in event time first, pooling the 51 event panels so that its timing and pre-event path are visible. The main estimates then use the stacked late-window design, with OLS and instrumented effects of news coverage side by side. Timing checks address reverse causality, and the remaining sections trace the geography, venues, and spending margins of the response.

4.1 News coverage and the response in event time

Figure 1 shows when the effect of news coverage emerges. It plots the additional high-FB versus low-FB foot-traffic gap associated with a one-standard-deviation larger Spanish-language news-coverage spike. Estimates are shown for the full sample and for the subsample of 23 events with no overlapping-catchment contamination and statistically flat pre-trends.

The relationship is flat and near zero through the pre-event period in both samples. It turns negative only after the event and reaches about -0.02 log points per standard deviation of news coverage in the late window. The similar point estimates indicate that the result is not driven by overlapping catchments or by events with non-flat pre-trends. The corresponding average high-FB versus low-FB path and additional timing checks appear in the Online Appendix.

Figure 1: News coverage and the foot-traffic response in event time



Notes. Pooled distributed-lag DiD with event $\times$ CBG and event $\times$ calendar-week fixed effects. The figure plots the additional high-FB versus low-FB foot-traffic gap associated with a one-standard-deviation larger Spanish-language news-coverage spike. The clean subsample has no overlapping-catchment contamination and flat pre-trends. Shaded bands are 95 percent confidence intervals based on standard errors clustered by event.

4.2 The effect of news coverage

Table 1 reports the main estimates from the stacked late-window specification. The coefficient on high-FB $\times$ weeks 4–8 $\times$ news coverage is negative in OLS. Events that receive a one-standard-deviation larger Spanish-language news-coverage spike produce about a 1.5 percent larger decline in foot traffic in high-FB neighborhoods relative to low-FB neighborhoods in the same event catchment. Adding log arrests leaves the coverage coefficient essentially unchanged, while the arrests interaction is small. When arrests are entered without news coverage, they explain little of the post-event response.

Table 1: Pooled late-window effect of news coverage on foot traffic

	OLS			2SLS	
	(1) Coverage	(2) + arrests	(3) Arrests only	(4) Coverage IV	(5) + arrests
High-FB $\times$ late $\times$ coverage	−0.0149*** (0.0041)	−0.0140*** (0.0041)		−0.0161*** (0.0061)	−0.0168*** (0.0063)
High-FB $\times$ late $\times$ log arrests		−0.0032 (0.0051)	−0.0053 (0.0045)		−0.0028 (0.0054)
High-FB $\times$ late	−0.0076* (0.0042)	−0.0072 (0.0046)	−0.0061 (0.0052)	−0.0076* (0.0042)	−0.0074 (0.0048)
First-stage F				37.4	33.0
Events	51	49	49	51	49
CBG-week observations	219,555	195,194	195,194	219,555	195,194

Notes: Stacked CBG-week DiD. Outcome is log weekly foot traffic. All columns include event $\times$ CBG and event $\times$ calendar-week fixed effects. Late window is weeks 4–8. Coverage is the standardized Mediacloud Spanish-ICE spike, and log arrests is standardized. Columns (4)–(5) instrument coverage with foreign-disaster news pressure. Standard errors are clustered by event. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

The instrumented estimates are similar in sign and magnitude. The IV estimate implies that one standard deviation of additional news coverage reduces foot traffic by about 1.6 percent in the late window. Including arrests as a control leaves the estimate unchanged.

The Online Appendix reports additional robustness checks. The English news spike from GDELT has the same sign and statistical significance, and the US national Google Trends spike gives the same answer for the 2025–2026 event wave. The effect survives controlling for operation type and sector. It is also robust to using raw arrests, log arrests, arrests per catchment population, and arrests per high-FB population as alternative scale proxies. These checks make it unlikely that the effect of news coverage simply proxies for operational scale. Controlling for pre-event news coverage in the four weeks before the raid does not explain away the post-event effect. Estimating the outcome by Poisson pseudo-maximum likelihood rather than in logs gives similar results, and the effect builds over the first eight post-event weeks.

4.3 Instrumenting news coverage with news shocks

Realized news coverage is not randomly assigned. Features of an operation that make it salient to editors could also make it salient to residents, so the OLS association could overstate the causal effect of news coverage. This concern is addressed with an instrumental-variables strategy following the competing-news logic of [Eisensee and Strömberg \(2007\)](#). In short, the idea is that independent national news shocks can shift coverage of an enforcement event without altering the operation itself.

The preferred instrument uses major foreign disaster news shocks. It counts the number of days in the post-event window $\tau \in [0, 8]$ that fall within the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. These disasters occur independently of the enforcement event and, to varying degrees, occupy national news attention. The outcome is a within-catchment differential between heavily Latino-immigrant neighborhoods and otherwise similar neighborhoods. A national news shock that moves attention equally in both groups nets out of this differential. The residual threat is therefore a home-region disaster that differentially affects the treated population. Latin America and the Caribbean are therefore excluded, since disasters there could directly affect Latino immigrant communities through family ties, remittances, or Spanish-language news demand. Out-of-state U.S. disasters are also reported as a complementary margin of domestic disaster news shocks. These instruments need not identify the same local average treatment effect. Each isolates events whose public attention is shifted by a particular source of competing news shocks.

Columns (4) and (5) of Table 1 report the second-stage estimates. Foreign-disaster pressure has the expected crowd-out sign, with first-stage F statistics above 30 in the pooled late-window specification with event-clustered standard errors. The second-stage estimates are slightly larger in magnitude than the OLS estimates. Reported arrests do not predict the foot-traffic response in the main specifications, and the IV estimate is unchanged when arrests are included as a control. These results make an arrest channel less likely, although

arrest counts remain an imperfect measure of operational scale. The IV identifies the effect of news coverage for mid-visibility events at the margin of the news cycle. These are not the largest, most visible operations. The Online Appendix reports instrument balance, exclusion-restriction checks, and complementary specifications using out-of-state U.S. disasters.

4.4 Timing and design validity

Because news coverage and foot traffic are measured after the same event, the estimated effect could in principle reflect coverage responding to the neighborhood response rather than the reverse. Table 2 reports three tests. First, the effect is re-estimated using only news coverage from the first two post-event weeks, which predates the late-window response by at least three weeks and so cannot be caused by it. This temporally ordered estimate remains negative but is smaller than the corresponding full-window estimate, and it is essentially unchanged by a control for arrests. Second, the reverse direction is tested directly. The early foot-traffic response does not predict late-window news coverage conditional on early news coverage. Third, a weekly cross-lag panel with event and event-time fixed effects shows the same asymmetry. Lagged news coverage predicts the next week’s response differential, while the lagged response does not predict the next week’s news coverage. These tests show no evidence of feedback from the observed response to measured news coverage, though they do not rule out every common cause. The arrests control addresses operational severity, the leading candidate common cause, and the news-shock instrument addresses common causes in the news environment.

Table 2: Temporal ordering of news coverage and the foot-traffic response

	News coverage → response	Response → news coverage
<i>Panel A: Cross-event ordering ($N = 51$ events)</i>		
Early vs. late window	−0.0103** (0.0051)	−0.0392 (0.1256)
+ arrests control	−0.0098* (0.0055)	
<i>Panel B: Weekly cross-lag, event and event-time FE ($N = 459$ event-weeks, 51 events)</i>		
One-week lag	−0.0586* (0.0295)	0.0728 (0.0848)

Notes: Panel A uses the separate-event robustness estimates. The first column regresses the weeks 4–8 foot-traffic response, in log points, on the standardized peak in Spanish-language news coverage during weeks 0–1. The second column reverses the ordering by regressing standardized late coverage on the standardized early response, controlling for early coverage. Panel B pools weeks 0–8 with event and event-time fixed effects. Its variables are standardized, and each specification controls for the once-lagged dependent variable. Standard errors use Knapp–Hartung inference in the first Panel A column, HC1 in the second, and two-way clustering by event and calendar week in Panel B. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Another concern is that high-coverage events may occur in catchments where the event-study comparison is less credible. Figure 1 addresses this concern directly. The coverage path is similar when the sample is restricted to the 23 events with no overlapping-catchment contamination and statistically flat pre-trends. The Online Appendix also shows that news coverage is unrelated to pre-trend flatness, the number of block groups in the catchment, estimate precision, or exposure to overlapping events.

The effect is also a within-wave phenomenon rather than an artifact of the difference between the 2018–2019 and 2025–2026 enforcement waves. Estimated on the 2025–2026

wave alone, which contains 44 of the 51 events, it is if anything stronger (Online Appendix).

4.5 National vs. local attention

Holding the search platform and search terms fixed and varying only geography, national Google Trends predicts larger foot-traffic declines than metro-level (DMA) Trends, and DMA Trends adds little once national coverage is included. National coverage therefore provides the more predictive variation, and reported arrests remain null. The Online Appendix reports the national-versus-DMA comparison for the 2025–2026 event wave.

4.6 Mechanisms

The evidence so far shows that news coverage predicts the size of the neighborhood response. I next ask what kind of response it produces. If direct disruption is the main channel, the effect should be strongest close to the worksite and in the raided sector. If perceived risk is the main channel, it should extend across geography and across sectors. The negative response is not confined to the immediate worksite area. It appears in the 5–15 km and 15–30 km bands and remains negative, though smaller, in the 30–50 km band. It is also not concentrated in the raided sector. The same-sector estimate is not negative, while other-sector POIs and both consumer- and nonconsumer-facing POIs decline. The response therefore does not look like a same-sector demand shock or a purely retail channel. The Online Appendix reports the distance- and sector-split estimates.

Tests of the timing of coverage point the same way. Weekly news peaks, national search, and an eight-week English-news area-under-the-curve measure predict the response, while one-day peaks and days-above-threshold measures are weaker (Online Appendix). The mechanism evidence is most consistent with a broad perceived-risk response among heavily Latino-immigrant neighborhoods.

4.7 Civic vs. commercial response

I next identify the broad activity types most affected by the enforcement actions. The decomposition uses two POI bins. The “civic/community” bin includes NAICS-3 codes 712 and 611, which cover parks, museums, and educational services. The “commercial” bin includes NAICS-2 codes 44, 45, and 72, which cover retail and food services. NAICS-3 code 813 (religious organizations and civic associations) is omitted from the civic bin because the category is heavily contaminated by community-bank branches rather than civic organizations.

Twelve events have both a civic and a commercial estimate. Across these events, civic foot traffic in high-FB CBGs falls by an average of 0.120 log points relative to low-FB CBGs, while the corresponding commercial decline is smaller and not distinguishable from zero. The within-event civic-minus-commercial differential averages -0.111 log points and is statistically significant. This larger estimate can coexist with the small aggregate response because civic POIs are a limited share of observed traffic and the civic comparison is available for only 12 events. The activity that falls most sharply is therefore institutional contact at schools, parks, and museums, rather than purely commercial activity. The civic sample is small, and I cannot determine whether the gap is itself moderated by news coverage. The point estimates run opposite to that hypothesis, but the coverage subgroups contain only six events each and are not statistically distinguishable. The civic result is therefore reported as an average contrast rather than a coverage-moderated effect. The Online Appendix reports the inference and venue-coverage limits.

4.8 Spending

I also estimate the spending response using $\log(1 + \text{weekly spending})$ from the exact-geocoded SafeGraph SpendPatterns panel. Thirty events have enough high- and low-FB CBG coverage for this exercise. The spending estimates are much noisier than the foot-traffic estimates, and the aggregate is statistically indistinguishable from zero. Spending appears to fall more after

events with larger Spanish-ICE and English-GDELT news spikes. Above-median-coverage events show a differential decline of about 3 percent, while below-median-coverage events show a small increase of about 4 percent, but the cross-event gradient is not statistically significant (Online Appendix). Because the panel records only card spending at covered merchants, it does not capture the full spending response.

5 Conclusion

This paper studies how public attention changes the local economic footprint of immigration enforcement. Across 51 single-worksite ICE events in 26 states, events that receive more national Spanish-language news coverage generate larger declines in foot traffic in heavily Latino-immigrant neighborhoods. English-language news and search data show the same pattern. A stacked difference-in-differences design and an instrumental-variables strategy based on competing news shocks from major foreign disasters both point to the same conclusion. Media coverage of enforcement has an independent effect on local activity.

The evidence also helps interpret what kind of response this is. Reported arrests, whether raw, logged, or scaled by local population, explain little of the response on their own and do not attenuate the coverage effect. The response is more closely tied to national than local attention. It extends beyond the immediate worksite and beyond the raided sector, and it appears especially in civic and institutional venues such as schools, parks, and museums. The cost of visible enforcement is therefore not limited to lost transactions. It also includes reduced participation in public and institutional life.

One implication is that ambiguity itself can be costly. A raid may be localized in operational terms, but coverage can turn it into a broader signal of risk. When enforcement actions become highly visible, agencies and local officials can reduce harmful spillovers by clearly communicating the scope and limits of those actions. The broader lesson is that the economic footprint of enforcement depends not only on what the government does, but also

on how widely that action is noticed and understood.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work, the author used OpenAI's ChatGPT and Codex to assist with coding, analysis checks, proofreading, and manuscript editing. After using these tools, the author reviewed and edited the content as needed and takes full responsibility for the content of the article.

References

- Advan Research. Foot traffic / weekly patterns plus [dataset]. Dewey Data, 2025. doi: 10.82551/C103-N851.
- Umair Ali, Jessica H. Brown, and Chris M. Herbst. Secure communities as immigration enforcement: How secure is the child care market? *Journal of Public Economics*, 233: 105101, 2024. doi: 10.1016/j.jpubeco.2024.105101.
- Marcella Alsan and Crystal S. Yang. Fear and the safety net: Evidence from secure communities. *Review of Economics and Statistics*, 106(6):1427–1441, 2024. doi: 10.1162/rest_a_01250.
- Catalina Amuedo-Dorantes and Francisca M. Antman. De facto immigration enforcement, ICE raid awareness, and worker engagement. *Economic Inquiry*, 60(1):373–391, 2022. doi: 10.1111/ecin.13041.
- Catalina Amuedo-Dorantes, Esther Arenas-Arroyo, and Almudena Sevilla. Immigration enforcement and economic resources of children with likely unauthorized parents. *Journal of Public Economics*, 158:63–78, 2018. doi: 10.1016/j.jpubeco.2017.12.004.
- Catalina Amuedo-Dorantes, Esther Arenas-Arroyo, and Chunbei Wang. Is immigration enforcement shaping immigrant marriage patterns? *Journal of Public Economics*, 190: 104242, 2020. doi: 10.1016/j.jpubeco.2020.104242.
- Asad L. Asad. *Engage and Evade: How Latino Immigrant Families Manage Surveillance in Everyday Life*. Princeton University Press, 2023. doi: 10.1515/9780691249049.
- Scott R. Baker, Robert A. Farrokhnia, Steffen Meyer, Michaela Pagel, and Constantine Yannelis. How does household spending respond to an epidemic? Consumption during the 2020 COVID-19 pandemic. *Review of Asset Pricing Studies*, 10(4):834–862, 2020. doi: 10.1093/rapstu/raaa009.

- Raj Chetty, John N. Friedman, and Michael Stepner. The economic impacts of COVID-19: Evidence from a new public database built using private sector data. *Quarterly Journal of Economics*, 139(2):829–889, 2024. doi: 10.1093/qje/qjad048.
- Alberto Ciancio and Camilo García-Jimeno. Anticipating state action: Risk perceptions and consumption under immigration enforcement. Discussion Paper 18446, IZA Institute of Labor Economics, 2026.
- Elizabeth Cox and Chloe N. East. Labor market impacts of ICE activity in Trump 2.0. Working Paper 35129, National Bureau of Economic Research, 2026.
- Uma De Balanzó, Núria Rodríguez-Planas, and Jennifer Roff. Immigration enforcement visibility and consumer spending. Discussion Paper 18620, IZA Institute of Labor Economics, 2026.
- Chloe N. East, Annie Laurie Hines, Philip Luck, Hani Mansour, and Andrea Velásquez. The labor market effects of immigration enforcement. *Journal of Labor Economics*, 41(4): 957–996, 2023. doi: 10.1086/721152.
- Thomas Eisensee and David Strömberg. News droughts, news floods, and U.S. disaster relief. *Quarterly Journal of Economics*, 122(2):693–728, 2007. doi: 10.1162/qjec.122.2.693.
- Global Disaster Alert and Coordination System. GDACS event feed [database]. <https://www.gdacs.org>, 2026. Accessed via GDACS API.
- Austan Goolsbee and Chad Syverson. Fear, lockdown, and diversion: Comparing drivers of pandemic economic decline 2020. *Journal of Public Economics*, 193:104311, 2021. doi: 10.1016/j.jpubeco.2020.104311.
- Exequiel Hernandez. ICE-ing the economy: Immigration enforcement under Trump 2.0 and local economic activity. Technical report, SSRN, 2026.

- Jordan Herring and Burt S. Barnow. Indirect effects of immigration enforcement on health care utilization among lawfully present older Hispanics. *Social Science & Medicine*, 384: 118540, 2025. doi: 10.1016/j.socscimed.2025.118540.
- Kalev Leetaru and Philip A. Schrodtt. GDELT: Global data on events, location, and tone, 1979–2012. In *International Studies Association Annual Convention*, San Francisco, CA, 2013.
- T. William Lester, Matthew Wilson, and Eli Knaap. The “chilling effect” of ICE enforcement: Evidence from high-frequency mobility and spending data in the Los Angeles region. Working paper, January 26, 2026, 2026. URL <https://bpb-us-e2.wpmucdn.com/sites.uci.edu/dist/6/5832/files/2026/03/Chilling-Effect-of-ICE-in-LA-LWK.pdf>.
- Media Cloud. Media Cloud online news archive [database]. <https://search.mediacloud.org>, 2025. Accessed via Python API client (`mediacloud>=4.3.0`); US Spanish Language Collection, id 196136637.
- SafeGraph. Spend patterns [dataset]. Dewey Data, 2022. doi: 10.82551/NSF5-R186.
- Tara Watson. Inside the refrigerator: Immigration enforcement and chilling effects in Medicaid participation. *American Economic Journal: Economic Policy*, 6(3):313–338, 2014. doi: 10.1257/pol.6.3.313.