

# Attention and the Economic Response to Immigration Enforcement

Greg C. Wright

University of California, Merced

[gwright4@ucmerced.edu](mailto:gwright4@ucmerced.edu)

July 2, 2026

## Abstract

ICE enforcement reduces foot traffic in some immigrant neighborhoods but not others. Across 51 single-worksite ICE events in 26 states (2018–2026), a per-event difference-in-differences compares heavily Latino-immigrant neighborhoods with similar nearby ones. The average response is near zero, yet about 70 percent of the cross-event variation reflects true differences rather than noise. Post-event Spanish-language news coverage predicts larger declines and explains roughly 30 percent of the variation, while arrests explain none. The coverage effect is absent in pre-event weeks and emerges only after the event, and the early foot-traffic response does not predict later coverage. English-language news and search show the same pattern. A news-pressure instrument supports a causal interpretation. Civic and institutional visits fall more than commercial activity.

**Keywords:** public attention, immigration enforcement, chilling effects, foot traffic.

**JEL Codes:** J15, R23, K37.

*Acknowledgments:* I thank the UC Merced Community and Labor Center for support and Albert Kochaphum for helpful conversations.

# 1 Introduction

Immigration enforcement can affect the lives of individuals beyond those who are detained. In part, this is because the community response may depend as much on the information environment surrounding an event as on the event itself. A workplace raid that goes largely unnoticed may have little effect beyond the worksite, whereas a smaller raid that becomes highly visible may lead residents to avoid public places. A long literature in economics and sociology documents these “chilling effects,” which can lead to reduced use of healthcare, schooling, and public benefits due to fear of detention [Watson, 2014, Asad, 2023]. The intensification of U.S. Immigration and Customs Enforcement (ICE) activity under the second Trump administration provides a real-time test. For example, Lester et al. [2026] document an 8–10 percent foot-traffic decline and a 20–25 percent spending decline in the most heavily Latin American-born neighborhoods of Los Angeles–Orange County in the weeks following ICE’s May 14, 2025 announcement of a major regional enforcement operation.

In this paper I ask: are the observed chilling effects due to the operations themselves, or their visibility to the surrounding community? If the scale of an operation is central to its effect, arrest counts should help explain which events produce larger responses. On the other hand, if the broader *perception* of risk is what matters, the response will depend on whether news of the event reaches people who were not directly targeted. A raid affects both margins at once, detaining a limited number of individuals while also generating news coverage, social-media activity, rumor, and protest.

I exploit cross-event variation to separate these channels. The sample comprises 51 ICE enforcement events from April 2018 through January 2026 across 26 states, spanning a small cluster of high-profile 2018–2019 worksite actions and a much larger cluster of 2025–2026 operations. For each event a distributed-lag difference-in-differences (DiD) specification is estimated on Advanced cellphone foot-traffic data, comparing Census Block Groups (CBGs) in the top decile of Latin American foreign-born (FB) share to those in the bottom decile, and a summary estimate of the average response is extracted.<sup>1</sup> Because an average can hide very different event-level responses, I first measure how much the true response varies. A random-effects decomposition rejects a common effect and attributes roughly 70 percent of the cross-event variation to genuine differences rather than estimation noise, with the response to a new event ranging from a roughly 5 percent decline to little or no response. At matched distances from the reference week the dispersion is uniformly larger after the event than before, consistent with heterogeneity triggered by the event rather than pre-existing differences. I then ask which features of the event predict where in that range a neighborhood falls.

The strongest predictor of foot traffic is the public attention garnered by the event, not its operational scale. I exploit post-event spikes in Spanish-language news-volume queries for “ICE”, the agency name most salient to the population at risk, constructed from a collection of US

---

<sup>1</sup>This is the specification Lester et al. [2026] use for their single LA event, which is implemented across each event in the sample for comparison.

Spanish-language outlets with daily coverage across the sample period. Events that receive more Spanish-language attention generate larger declines in foot traffic, and the same pattern appears in English-language news and Google search data. In a precision-weighted meta-regression, the Spanish-language attention spike accounts for roughly 30 percent of the true cross-event variation in responses, while arrest counts account for essentially none. The absence of an effect due to arrests is consistent with the information mechanism.

The pattern is robust to several checks. It holds when attention and arrests are included together, when noisier event estimates are down-weighted, when pre-event media attention is accounted for, and when I estimate within-event changes rather than levels. The attention effect is absent in every pre-event week and holds when coverage is measured only in the first two post-event weeks, well before the late-window response it predicts. The early foot-traffic response also does not predict later coverage. News of major foreign disasters, which crowd out public attention on enforcement actions, also serves as an instrument for the attention spike, with out-of-state U.S. disasters providing a complementary version. These sources of news pressure strongly predict the amount of attention given to enforcement events and are predetermined with respect to the operation itself.<sup>2</sup> Balance checks show no systematic relationship with pre-event foot-traffic trends or placebo responses, and the estimates are robust to a variety of controls. The effect is stronger within the recent 2025–2026 enforcement wave, suggesting that it does not reflect differences between the two enforcement eras.

National Google Trends and national Spanish-language news both predict the local foot-traffic decline, while metro-level search and local Spanish-language news add little once national attention is accounted for. The response is broad, extending beyond the immediate worksite and the raided industry. I decompose the response by venue and find that foot traffic around schools, parks, and museums falls more than at commercial sites such as retail and food services, a pattern more consistent with a broad withdrawal from public life than a localized demand shock.<sup>3</sup> The welfare cost of visible enforcement is therefore not limited to lost commercial activity but includes reduced participation in civic life.

A broader economics literature uses interior-enforcement intensity and federal programs, especially Secure Communities, to estimate effects of immigration enforcement on safety-net participation, family decisions, childcare, healthcare, and labor markets [Watson, 2014, Amuedo-Dorantes et al., 2018, 2020, Alsan and Yang, 2024, Ali et al., 2024, Herring and Barnow, 2025, East et al., 2023], and a related sociology literature describes the institutional-avoidance strategies of mixed-status families [Asad, 2023]. Recent work studies the current wave: Amuedo-Dorantes and Antman [2022] link removals and raid awareness to lower immigrant labor-market engagement, Ciano and García-Jimeno [2026] study consumption responses to learned enforcement risk, and De Balanzó et al. [2026], Hernandez [2026], and Cox and East [2026] estimate economic effects of the 2025 wave using spending, mobility, and labor-market data. These papers show that enforcement moves eco-

---

<sup>2</sup>Since competing news shifts attention only for events whose coverage is sensitive to the news cycle, the instrument identifies an effect for mid-visibility events rather than for the largest, most visible operations.

<sup>3</sup>Religious organizations and hospitals are not present in the data; see Section 4.8.

conomic activity when risk becomes visible or learned. I instead hold the enforcement event fixed and ask why two raids of similar operational severity produce very different neighborhood responses, varying the event’s public visibility rather than the event itself.

These results also add to evidence on how public attention, more than operational scale, shapes the response to risk. A related COVID-mobility literature [Goolsbee and Syverson, 2021, Baker et al., 2020, Chetty et al., 2024] shows that perceptions of risk changed individual behavior independently of the formal COVID mandates. This is similar to the case studied here, where an immigration enforcement operation produces both a worksite disruption and public attention, but the community response tracks the attention. The same logic may extend to other salient but partly hidden risks whose economic footprint depends on how widely they are noticed. Where prior work instruments or exploits quasi-experimental variation in enforcement itself, I hold the event fixed and instrument its visibility, adapting the news-pressure logic of Eisensee and Strömberg [2007] to isolate the attention component of a specific event, which to my knowledge is new to the immigration literature.

## 2 Data

The analysis combines the event inventory with outcome, exposure, and attention data: cellphone foot traffic, SafeGraph consumer spending, and American Community Survey (ACS) demographics for outcomes and exposure; and Google Trends, Mediacloud US Spanish-language and GDELT 2.0 English-language news volume, Global Disaster Alert and Coordination System (GDACS) disaster alerts, and Nielsen Designated Market Area (DMA) definitions for attention and news pressure.

### 2.1 Event inventory

A list of single-worksite ICE enforcement events is compiled from news sources, agency press releases, and litigation records. The inventory covers two distinct enforcement waves from April 2018 through August 2019 and January 2025 through January 2026 (the second term’s worksite operations).<sup>4</sup> The inventory restricts attention to events that satisfy four criteria. First, the event date must fall within these two windows. Second, the event must correspond to a specific worksite or commercial address. Third, the worksite must be in one of 26 states for which CBG-level Latin American foreign-born data from the ACS are available. Finally, the worksite lat/lon must be geocoded either to the named business address or, when that is not feasible, to the city or county centroid.<sup>5</sup> These screens identify 55 single-worksite ICE-led events. Four have insufficient catchment point-of-interest (POI) coverage to estimate the event-study specification, leaving a

---

<sup>4</sup>The first Trump term’s worksite operations included Bean Station TN, Salem OH Fresh Mark, Mt. Pleasant IA Precast, Allen TX CVE Technology, Sumner TX Load Trail, and the August 2019 Mississippi poultry operation, which hit seven plants across five companies on a single day and which is included as a single Morton-MS-anchored compound event.

<sup>5</sup>Four events are explicitly coded as city-centroid locations, one additional event uses the earlier city-query fallback, and one event uses a county-level fallback.

main analysis sample of 51 events spanning 26 states. The Supplemental Appendix summarizes the event-level variables by enforcement wave. The 2018–2019 events report more arrests on average but generate much smaller news spikes than the 2025–2026 events, which motivates the era-robustness check in Section 4.2.

A public inventory could be selected toward visible raids. No complete administrative list of single-worksite ICE actions is publicly available for these windows, so the inventory is built from several approaches rather than national news alone: agency releases, litigation records, immigration-policy trackers, state and local reporting, and prior compilations. The result includes a set of events that received little attention. Of the 55 events, 32 are coded medium or low publicity and 29 have an English-GDELT news spike at or below 0.01 (Section 2.6), and the four POI-coverage exclusions include three medium-publicity events and one high-publicity event. These facts do not establish a complete universe of obscure raids, but they make it unlikely that the attention effect is mechanically generated by retaining only highly visible events. The Supplemental Appendix reports the sample flow, event enumeration, and source-type audit.

## 2.2 Foot traffic

Weekly visitor counts at point-of-interest level are obtained from Advan Research’s *Foot Traffic/Weekly Patterns Plus* dataset [Advan Research, 2025], derived from anonymized opt-in mobile devices. Each POI is geocoded to a latitude and longitude, assigned a North American Industry Classification System (NAICS) code, and assigned to its CBG. For each event, all Advan POIs within 50 km of the worksite anchor are collected and assembled into a POI  $\times$  week panel spanning  $\pm 12$  weeks of the event date.

The resulting multi-event panel contains approximately 69 million POI-week observations and 3.3 million distinct POI–event pairs across 1.5 million unique POIs. Within each event’s catchment, POI visits are aggregated to CBG-week totals and restricted to CBGs containing at least 10 POIs. This restriction concentrates the analysis on retail-active CBGs.

## 2.3 Consumer spending

Spending is measured with SafeGraph’s *SpendPatterns* dataset [SafeGraph, 2022], which reports monthly POI-level aggregates of consumer spending from a large opt-in financial-institution panel. Each row contains an array of daily spending values that are aggregated to weeks, producing a POI  $\times$  week spending panel comparable to the foot-traffic data. POI-week values are then aggregated to CBG-week using the same catchment definitions as the foot-traffic build.

## 2.4 ACS exposure measure

Treatment is defined as CBGs in the top decile of share of residents born in Latin America (“high exposure”), with bottom-decile CBGs as the comparison group. The 2018–2022 5-year ACS provides the Latin-American foreign-born share, computed at the tract level and allocated to all CBGs

within each tract.

The deciles are computed within each event’s catchment, so the threshold for “top decile” varies across events. In heavily-Latino metros, the top-decile cutoff is roughly 32 percent foreign-born from Latin America, while in less-Latino metros the threshold is much lower in absolute terms. This within-event normalization means the cross-event comparison is between “most-Latino” and “least-Latino” CBGs in each respective metro.

## 2.5 Google Trends salience

For each event, salience is measured using weekly Google Trends search interest for three keywords (“ICE raid”, “redada”, “deportación”) at three geographies: U.S. national, the event’s state, and its Nielsen Designated Market Area (DMA). Because Trends rescales each query within a geography and timeframe, within-geography pre-vs-post changes are used. The baseline is the four-week average over  $\tau \in [-8, -5]$  relative to the event Monday. Trends are restricted to the 2024–2026 wave because cross-era rescaling is noisy for these low-baseline terms.

## 2.6 News-volume series

**Mediacloud US Spanish-language news.** The primary news-volume series measures attention in the Spanish-language news environment the affected population consumes, on a daily scale consistent across both enforcement waves. The Mediacloud Online News Archive [Media Cloud, 2025] is queried for daily story counts in its curated US Spanish-language collection of 407 sources, including national, state, and local outlets. The headline query is the noun “ICE” restricted to this collection. Counts are normalized by the daily total of all stories in the collection to produce a share-of-coverage time series that is comparable across years. Complementary Spanish keywords are also pulled, and a state-local Spanish-news series is constructed by restricting the same query to sources tagged as published in the event’s home state.

**GDELT English-language US news.** As a corroborating measure on a different platform, language, and methodology, the GDELT 2.0 Document API [Leetaru and Schrodt, 2013] is queried for daily article counts matching “ICE raid” in US English-language sources, in normalized “Volume Intensity” units, giving a national daily series spanning January 2018 through March 2026. GDELT closely tracks the Mediacloud series at the per-event spike level.

## 2.7 Competing-news instruments

My preferred news-pressure instrument uses the Global Disaster Alert and Coordination System (GDACS) event feed [Global Disaster Alert and Coordination System, 2026]. Red-alert sudden disasters (earthquakes, tropical cyclones, floods, volcanic eruptions, wildfires) outside the United States, Latin America, and the Caribbean are obtained. The Latin America and Caribbean exclusion is deliberate, since homeland shocks could affect Latino immigrant communities directly

through family ties, remittances, or Spanish-language news demand rather than only through news-cycle competition. For each ICE event, the instrument counts days in the post-event window overlapping the first seven days of such an alert.

A complementary news-pressure measure is also constructed using a hand-coded calendar of major U.S. natural disasters in states other than the ICE event state. The foreign- and U.S.-disaster instruments identify different margins of news-cycle access, so the instrumental-variables (IV) estimates are interpreted as local effects for events whose attention is moved by the corresponding source of competing news pressure.

### 3 Empirical Strategy

#### 3.1 Per-event distributed-lag DiD

For each event  $e$  in the 51-event sample (single-worksite ICE enforcement events from April 2018 through January 2026, see Section 2.1), the same distributed-lag difference-in-differences model is estimated:

$$\log(1 + y_{c,t}^e) = \alpha_c^e + \lambda_t^e + \sum_{\tau \neq -1} \gamma_\tau^e (\text{HighFB}_c^e \times \mathbf{1}[t - T_0^e = \tau]) + \varepsilon_{c,t}^e, \quad (1)$$

where  $y_{c,t}^e$  is weekly visitor count (foot traffic) at CBG  $c$  in week  $t$ , restricted to the catchment of event  $e$ .  $T_0^e$  is the Monday of the week containing event  $e$ 's date. The dummy  $\text{HighFB}_c^e$  equals one for CBGs in the top decile of share Latin American foreign-born within event  $e$ 's catchment and zero for CBGs in the bottom decile (middle deciles are dropped).  $\alpha_c^e$  and  $\lambda_t^e$  are CBG and calendar-week fixed effects, both event-specific because Equation (1) is estimated separately for each event. Standard errors are clustered at the CBG level.

The event-time window is  $\tau \in [-8, +8]$  weeks and all seventeen weeks enter the estimation. The omitted category is  $\tau = -1$ , the last full pre-event week, so each  $\gamma_\tau^e$  measures the high-FB versus low-FB visit gap in week  $\tau$  relative to that reference week rather than a comparison of the event week to all other weeks. The per-event summary estimate  $\beta_{\text{late}}^e$  is the average of  $\gamma_\tau^e$  over the late window  $\tau \in [4, 8]$ , with a standard error computed as the standard error of that linear combination from the CBG-clustered covariance matrix. Because most events do not fall on a Monday (60 of 66 in the inventory), the  $\tau = 0$  week mixes pre- and post-event days.<sup>6</sup>

#### 3.2 Identifying variation

The design has two steps. First, Equation (1) uses within-event variation between high-FB and low-FB CBGs in the same catchment. CBG fixed effects absorb time-invariant differences in population, retail density, and baseline foot traffic. Calendar-week fixed effects absorb shocks common to all

---

<sup>6</sup>Re-anchoring  $\tau = 0$  to the first Monday on or after the event date, or dropping the partial event week entirely, leaves the results unchanged. The late-window estimate excludes the partial week by construction, and the attention effect below moves from  $-0.018$  to  $-0.016$  log points per standard deviation under re-anchoring (Supplemental Appendix).



CBGs in the catchment in a given week. The event estimate is therefore the differential foot-traffic response of heavily Latin-American-foreign-born CBGs relative to low-FB CBGs in the same metro and calendar weeks. Second, the 51 event estimates are used to test whether cross-event differences in the attention spike predict cross-event differences in the estimated response.

The design does not require raids to be unanticipated, and some initiatives are clearly partially anticipated. The initial question is whether the coverage generated by each event explains why otherwise similar raids produce different responses. A later section implements an instrumental-variables design to isolate plausibly exogenous variation in coverage from the surrounding news cycle, and Section 4.3 adds pre-event coverage as a control.

### 3.3 Information-shock and event severity measures

To test whether the chilling response is driven by information or by operational scale, two categories of cross-event measures are constructed. The headline information-shock measure is the post-event spike in the Mediacloud US Spanish-language news-volume series for the query “ICE” (Section 2.6), measured as the share of stories in the 407-source US Spanish-language collection that mention the term. The spike is computed as the peak over  $\tau \in [0, 8]$  minus the baseline over  $\tau \in [-8, -5]$ .

The headline measure is complemented with three corroborating measures: the analogous GDELT national US English-language news-volume spike for “ICE raid,” the national Google Trends spike for “ICE raid,” and the DMA-level Google Trends spike for the local media market in which the event occurred. The two Trends measures are available on the 2024–2026 subsample and are defined as within-geography pre-vs-post changes, since Google Trends rescales each query within a geography and timeframe.

With respect to the severity of the enforcement event, two measures of underlying operational scale are used. These are the count of arrests reported by ICE at the worksite,  $\text{Arrests}_e$ , and the same measure in log form,  $\log(1 + \text{Arrests}_e)$ , to accommodate the right-skewed distribution. The arrests measure has coverage of 49 of 51 fitted events.<sup>7</sup>

### 3.4 Horse-race specification

To assess whether attention rather than severity predicts chilling, a cross-event ordinary least squares (OLS) regression relates the per-event late-window coefficient to a measure of attention and, in the severity specification, the log-arrests control:

$$\beta_{\text{late}}^e = \pi_0 + \pi_S \text{Spike}_e + \pi_A \log(1 + \text{Arrests}_e) + u_e, \quad (2)$$

where  $\text{Spike}_e$  is the post-event spike in one of the four information-shock measures defined above. Heteroskedasticity-robust (HC1) standard errors are reported. The full specification table in the Supplemental Appendix reports the headline Mediacloud Spanish-ICE specification alone and with

---

<sup>7</sup>The two events with missing arrest counts in the inventory are Newark NJ (an unnamed business) and Baldwin County AL (an unnamed construction site).



the arrests control on the full 51-event panel, a standalone arrests specification, and single-regressor specifications using English-GDELT and national Google Trends as corroborating measures on the same windows. The Trends specification uses the 2024–2026 subsample. Comparing the single-regressor estimates across sources shows whether the same attention effect appears outside the headline Spanish-language measure.

### 3.5 News-pressure instruments

The OLS horse race treats realized post-event attention as the explanatory variable. To isolate variation in attention generated by the surrounding news cycle, predetermined competing-news shocks are used as instruments. The preferred instrument is major foreign disaster news pressure. It counts days in the post-event window that overlap with the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. The exclusion of Latin America and the Caribbean is intended to remove shocks that could directly affect Latino immigrant communities through homeland ties rather than through news-cycle competition only. The identifying assumption is that these shocks move enforcement action media coverage without directly changing the relative foot-traffic response in high-FB versus low-FB neighborhoods. Section 4.6 and the Supplemental Appendix report the corresponding balance and robustness checks.

Out-of-state U.S. disasters also serve as a complementary instrument. This measure captures domestic disaster-news pressure in states other than the ICE event state. The foreign- and U.S.-disaster instruments identify different local margins of news-cycle access. The IV estimates should therefore be read as local effects for events whose public attention is shifted by the corresponding source of competing news pressure.

## 4 Results


The response is shown in event time first, pooling the 51 event panels so that the average path, its pre-trends, and the scaling of the response with coverage are all visible. The cross-event analysis then quantifies the heterogeneity and asks which event features predict it. Timing tests address whether coverage could instead be responding to the neighborhood response, and an instrumental-variables strategy uses variation in news volume from major foreign disasters, with out-of-state U.S. disasters as a complementary margin.

### 4.1 The response in event time

Figure 1 pools the 51 event panels and traces the response in event time. Panel A plots the level path  $\gamma_\tau$ , the high-FB versus low-FB visit gap for an event of average attention, and Panel B plots the interaction path  $\delta_\tau$ , the additional gap per standard deviation of the Spanish-language coverage spike. Each is shown for the full sample and for the 23 events with no overlapping-catchment contamination and statistically flat pre-trends. Three features stand out. First, the level path shows no average response. The path  $\gamma_\tau$  moves around zero after the event in both samples, and

its late-window average is  $-0.003$  log points in each. Second, the full-sample pre-event level path is not flat, drifting down from about  $+0.013$  at  $\tau = -8$ , while the clean subsample's pre-event path stays near zero. The full-sample level path is therefore read as descriptive, and a formal sensitivity analysis of the pooled path to pre-trend violations appears in the Supplemental Appendix. Third, the interaction path is flat and near zero through the entire pre-period in both samples and turns negative only after the event, reaching about  $-0.02$  log points per standard deviation of coverage in the late window. The clean-subsample interaction,  $-0.018$  (SE  $0.006$ ), is nearly identical to the full-sample estimate of  $-0.019$  (SE  $0.003$ ), so the coverage effect is not driven by contaminated catchments or by the events with non-flat pre-trends.

Figure 1: The pooled response in event time




*Notes.* Pooled distributed-lag DiD across the 51 event panels with event  $\times$  CBG and event  $\times$  calendar-week fixed effects, standard errors clustered at the CBG level. Panel A plots the level path  $\gamma_\tau$ , the high-FB versus low-FB visit gap for an event of average attention. Panel B plots the interaction path  $\delta_\tau$ , the additional gap per standard deviation of the Mediacloud Spanish-language ICE coverage spike, standardized at the event level on the full sample. The clean subsample contains the 23 events with no overlapping-catchment contamination and statistically flat pre-trends ( $p > 0.10$ ). Shaded bands are 95 percent confidence intervals. The reference week  $\tau = -1$  is omitted and marked at zero.

Figure 2 shows the same structure at the level of dispersion. At each event time, a random-effects meta-analysis of the 51 per-event coefficients estimates the between-event standard deviation of true effects net of sampling noise. Because each  $\gamma_\tau^e$  is measured relative to the  $\tau = -1$  reference week, dispersion accumulates mechanically with distance from the reference in both directions,

so the informative comparison is at matched distance. Post-event dispersion exceeds pre-event dispersion at all seven matched distances, by a factor of 1.4 on average, and rises to 0.035 log points by  $\tau = +8$ . The lower panel estimates the attention meta-regression separately at each event time. The attention effect is indistinguishable from zero in every pre-event week and emerges after the event, stabilizing near  $-0.02$  log points per standard deviation over the late window. The cross-event heterogeneity that the rest of the paper studies is therefore consistent with being triggered by the event rather than pre-existing, and its alignment with coverage has no counterpart before the event.

Figure 2: Between-event dispersion and the attention effect over event time



*Notes.* The top panel reports, at each event time  $\tau$ , the estimated between-event standard deviation of true effects from a random-effects meta-analysis of the 51 per-event coefficients  $\gamma_\tau^e$ , each weighted by its CBG-clustered sampling variance, with REML point estimates and Q-profile 95 percent confidence bands. The bottom panel reports the coefficient from a random-effects meta-regression of  $\gamma_\tau^e$  on the standardized Spanish-language coverage spike at each  $\tau$ , with Knapp–Hartung 95 percent confidence intervals. The reference week  $\tau = -1$  is zero for every event by construction and is marked separately. Estimates are pointwise across  $\tau$ . The underlying numbers appear in the Supplemental Appendix.

## 4.2 Heterogeneity in the chilling response

The average neighborhood response to an enforcement event is close to zero, but the response varies far more across events than sampling error alone can explain. Across the 51 events the late-window (weeks 4 through 8) estimates average  $-0.008$  log points and are not distinguishable from zero on average, yet a test of effect homogeneity is decisively rejected (Cochran’s  $Q = 164.9$  on 50 degrees of freedom,  $p < 10^{-13}$ ). A random-effects decomposition attributes about 70 percent of the cross-event variance in the estimates to true differences in the underlying response rather than to estimation noise ( $I^2 = 69.7$  percent), with an estimated standard deviation of true effects of  $\tau = 0.029$  log points. The implied 95 percent prediction interval for the response to a new event spans  $[-0.057, +0.047]$  log points, about a 5.5 percent decline to a 4.8 percent increase. Table 1 reports the decomposition, and a forest plot of the 51 event estimates against the pooled mean and prediction interval appears in the Supplemental Appendix. The random-effects model puts about 43 percent of true effects above zero, in line with the 22 of 51 positive point estimates, and the sign is ordered by attention: negative estimates concentrate among high-attention events and positive estimates among low-attention events (see the Supplemental Appendix). The positive cases are therefore events where little public attention produced little chilling, rather than evidence that enforcement raises local activity.

This dispersion is structured by public attention. A random-effects meta-regression of the per-event estimate on the standardized post-event Spanish-language ICE news spike explains roughly 30 percent of the true between-event variance, and a one standard deviation larger spike predicts a 1.8 percent larger foot-traffic decline ( $-0.018$  log points,  $p < 0.001$  with Knapp-Hartung inference, Table 1, Panel B). Adding the log arrest count leaves the attention coefficient essentially unchanged, while the arrest coefficient is small and statistically insignificant.<sup>8</sup> The English-language GDELT news spike has the same sign and statistical significance and explains about 19 percent of the variance. The effect survives controlling for operation type and sector and is robust to plausible corrections for measurement error in arrests, making it unlikely that the attention effect simply proxies for operational scale (Supplemental Appendix).

---

<sup>8</sup>Arrests are a coarse proxy for operational scale; a small coefficient need not imply that no dimension of scale matters, only that the dimension captured by arrest counts does not predict the response.

Table 1: Direct estimation of cross-event heterogeneity

|                                                           |                   |
|-----------------------------------------------------------|-------------------|
| <i>Panel A: Heterogeneity of the late-window response</i> |                   |
| Events                                                    | 51                |
| Cochran’s $Q$                                             | 164.9             |
| $Q$ $p$ -value                                            | 0.0000            |
| $I^2$                                                     | 69.7%             |
| $\tau$ (REML, log points)                                 | 0.0287            |
| $\tau$ (DL, log points)                                   | 0.0254            |
| RE pooled mean                                            | -0.0049           |
| 95% CI                                                    | [-0.0145, 0.0047] |
| 95% prediction interval                                   | [-0.0568, 0.0470] |
| <i>Panel B: Meta-regression on attention (per SD)</i>     |                   |
| Spanish-ICE spike                                         | -0.0184           |
| (KH SE)                                                   | (0.0051)          |
| KH $p$                                                    | 0.0006            |
| $R^2_{\text{meta}}$                                       | 30.1%             |
| Spanish-ICE   arrests                                     | -0.0192           |
| (KH SE)                                                   | (0.0053)          |
| $R^2_{\text{meta}}$                                       | 25.1%             |
| English-GDELT spike                                       | -0.0154           |
| (KH SE)                                                   | (0.0050)          |
| $R^2_{\text{meta}}$                                       | 18.5%             |

*Notes:* Random-effects meta-analysis of the 51 per-event late-window (weeks 4–8) estimates, using the correct standard error of the late-window average (SE of the linear combination  $c'\beta$  from the CBG-clustered covariance matrix, not the mean of the weekly SEs).  $\tau$  is the estimated standard deviation of true effects across events, shown under restricted maximum likelihood (REML) and DerSimonian–Laird (DL). The 95% prediction interval is the Higgins–Thompson–Spiegelhalter interval for the true effect of a new event. Panel B reports random-effects meta-regressions of the per-event estimate on standardized attention spikes with Knapp–Hartung (KH) adjusted inference;  $R^2_{\text{meta}}$  is the share of between-event variance ( $\tau^2$ ) explained, with  $\tau^2$  estimated by DerSimonian–Laird for both the intercept-only and meta-regression models on the same sample. Spikes are standardized to the 51-event SD; the arrests specification has  $N = 49$ .

The heterogeneity is not an artifact of the event-study design varying across contexts. Per-event design quality, measured by the flatness of pre-event trends, the number of block groups in the catchment, the precision of the estimate, and exposure to overlapping nearby events, is uncorrelated with the attention spike, and the attention effect is unchanged when these design characteristics are added as controls (Supplemental Appendix). Restricting the sample to the 34 events with no overlapping-catchment contamination, and further to the 23 events with both no contamination and

statistically flat pre-trends, leaves the measured heterogeneity intact, with  $I^2$  of 61 and 63 percent, and preserves the attention effect, although precision falls in the smallest subsample. The effect is also a within-wave phenomenon rather than an artifact of the difference between the 2018–2019 and 2025–2026 enforcement eras. Estimated on the 2025–2026 wave alone, which contains 44 of the 51 events, it is if anything stronger. As Section 4.1 shows, the pooled level path is descriptive rather than identifying since its late-window estimate is sensitive to modest departures from parallel pre-trends. The cross-event design and the instrumental-variables estimates below, not the pooled level path, carry the identification (Supplemental Appendix).

### 4.3 The information-shock horse race

The cross-event OLS horse race of Equation (2) corroborates the meta-regression, with the full specification table in the Supplemental Appendix.<sup>9</sup> The specifications anchor on the Mediacloud US Spanish-language news-volume spike for the query “ICE” (407-source Spanish-language collection, Section 2.6), add a standalone arrests specification, and replace the headline measure with two corroborating attention proxies built on different platforms, the English-language GDELT national news-volume spike and the US national Google Trends spike for “ICE raid.” A larger Spanish-ICE spike predicts a larger decline in foot traffic. Adding  $\log(1 + \text{Arrests}_e)$  leaves the Spanish-ICE coefficient essentially unchanged, while arrests alone explain very little of the cross-event variation. English-GDELT and national Google Trends show the same negative association with foot traffic.

Several robustness checks support this interpretation. Spanish-language news, English-language news, and national search measures move together across events. The post-event attention result is not explained by public news attention in the four weeks before the raid. Precision-weighted estimates, event-level bootstrap intervals, pre-event coefficients, and placebo estimates give the same conclusion. The Supplemental Appendix reports these checks, along with an early-window table showing that the attention effect builds over the first eight post-event weeks.

### 4.4 National vs. local attention

Holding the search platform and query fixed and varying only geography, national Google Trends predicts larger foot-traffic declines than metro-level (DMA) Trends, and DMA Trends adds little once national attention is included. National rather than local attention appears to carry the signal, and reported arrests remain null. The national-vs.-DMA Google Trends horse race for the 2024–2026 subsample appears in the Supplemental Appendix.

### 4.5 Timing and reverse causality

Because coverage and foot traffic are measured after the same event, the estimated effect could in principle reflect coverage responding to the neighborhood response rather than the reverse.

---

<sup>9</sup>The horse-race and instrumental-variables coefficients are scaled per unit of the unstandardized attention spike, so they are larger in magnitude than the per-standard-deviation meta-regression coefficient in Section 4.2.



Timing restricts that account, and Table 2 reports three tests. First, the effect is re-estimated using only coverage from the first two post-event weeks, which predates the late-window response by at least three weeks and so cannot be caused by it. This ordered effect is  $-0.010$  log points per standard deviation ( $p = 0.048$ ), about half the full-window effect, and the point estimate is essentially unchanged by an arrests control ( $-0.010$ ,  $p = 0.08$ ). Second, the reverse direction is tested directly. The early foot-traffic response (weeks 0 through 2) does not predict late-window coverage conditional on early coverage, with a coefficient near zero ( $p = 0.76$ ). Third, a weekly cross-lag panel with event and event-time fixed effects shows the same asymmetry: lagged coverage predicts the next week’s response differential ( $-0.059$  standard deviations,  $p = 0.052$ ), while the lagged response does not predict the next week’s coverage ( $p = 0.39$ ). Temporal precedence rules out feedback from the measured response to the measured coverage, though not every common cause. The arrests control addresses operational severity, the leading candidate common cause, and the instrumental-variables strategy below addresses common causes in the news environment.

Table 2: Temporal ordering of attention and the foot-traffic response

|                                                                                                         |  |          |
|---------------------------------------------------------------------------------------------------------|--|----------|
| <i>Panel A: Cross-event ordering (<math>N = 51</math> events)</i>                                       |  |          |
| Late response ( $\tau \in [4, 8]$ ) on early attention ( $\tau \in [0, 1]$ )                            |  | -0.0103  |
| (KH SE)                                                                                                 |  | (0.0051) |
| KH $p$                                                                                                  |  | 0.0476   |
| ...controlling for arrests (operational scale)                                                          |  | -0.0098  |
| (KH SE)                                                                                                 |  | (0.0055) |
| KH $p$                                                                                                  |  | 0.0775   |
| ...controlling for the early response ( $\tau \in [0, 2]$ )                                             |  | -0.0035  |
| (KH SE)                                                                                                 |  | (0.0035) |
| KH $p$                                                                                                  |  | 0.3152   |
| Late attention ( $\tau \in [4, 8]$ ) on early response ( $\tau \in [0, 2]$ )                            |  | -0.0392  |
| (SE)                                                                                                    |  | (0.1256) |
| $p$                                                                                                     |  | 0.7563   |
| <i>Panel B: Weekly cross-lag, event and event-time FE (<math>N = 459</math> event-weeks, 51 events)</i> |  |          |
| $\gamma_{e,\tau}$ on $\text{attention}_{e,\tau-1}$ (and $\gamma_{e,\tau-1}$ )                           |  | -0.0586  |
| (SE)                                                                                                    |  | (0.0295) |
| $p$                                                                                                     |  | 0.0524   |
| $\text{Attention}_{e,\tau}$ on $\gamma_{e,\tau-1}$ (and $\text{attention}_{e,\tau-1}$ )                 |  | 0.0728   |
| (SE)                                                                                                    |  | (0.0848) |
| $p$                                                                                                     |  | 0.3947   |

*Notes:* Attention is the weekly US Spanish-language Mediacloud ICE coverage ratio aligned to each event's Monday anchor, as a deviation from the event's pre-period baseline (mean over  $\tau \in [-8, -5]$ );  $\gamma_{e,\tau}$  is the per-event distributed-lag DiD coefficient (high- vs. low-foreign-born CBGs, reference  $\tau = -1$ ). All regressors and outcomes are standardized, so coefficients are per SD. Panel A rows 1–2 are random-effects meta-regressions weighted by the CBG-clustered sampling variance of the late-window mean, with Knapp–Hartung inference; row 3 is OLS with HC1 standard errors. Panel B pools post-event weeks ( $\tau \in [0, 8]$ ) with event and event-time fixed effects and standard errors two-way clustered by event and ISO calendar week. Where the estimated  $\gamma$  appears as a regressor, sampling noise attenuates its coefficient toward zero, so the early-response control is imperfect. Temporal precedence rules out backward-in-time causation, not common causes; the arrests control addresses the leading common-cause candidate (operational severity). The Panel B lagged-dependent-variable specifications with unit fixed effects in a short panel are subject to Nickell bias.

#### 4.6 Instrumenting attention with news pressure

The timing tests rule out feedback from the measured response to the measured coverage, but coverage and the response could still share a common cause, such as features of an operation that make it salient to editors and residents alike. This concern is addressed with an instrumental-variables strategy following the news-pressure logic of [Eisensee and Strömberg \[2007\]](#). In short, the idea is that predetermined national news-cycle conditions can shift attention to an enforcement event without altering the operation itself.

The preferred instrument is major foreign disaster news pressure. It counts the number of days in the post-event window  $\tau \in [0, 8]$  that fall within the first seven days of a GDACS red-alert

sudden disaster outside the United States, Latin America, and the Caribbean. These disasters are predetermined with respect to the enforcement event and to varying degrees they occupy national news attention. Because the outcome is a within-catchment differential between heavily Latino-immigrant neighborhoods and otherwise similar neighborhoods, a national news shock that moves attention equally in both groups nets out of this differential, so the residual threat is a home-region disaster that differentially affects the treated population. Latin America and the Caribbean are therefore excluded, where shocks could directly affect Latino immigrant communities through family ties, remittances, or Spanish-language news demand. Out-of-state U.S. disasters are also reported as a complementary margin of domestic disaster-news pressure. These instruments need not identify the same local average treatment effect. Each isolates events whose enforcement attention is moved by a particular source of competing news pressure.

Table 3: OLS and IV estimates of the effect of ICE media coverage on foot traffic

|                                                                 | (1) OLS              | (2) IV: Foreign disasters | (3) IV: US disasters<br>(cross-state) | (4) Both IVs<br>(over-id) |
|-----------------------------------------------------------------|----------------------|---------------------------|---------------------------------------|---------------------------|
| <i>Panel A: Mediacloud Spanish-ICE spike (endogenous)</i>       |                      |                           |                                       |                           |
| Coverage coefficient                                            | −0.919***<br>(0.256) | −1.059***<br>(0.338)      | −0.911<br>(0.796)                     | −1.052***<br>(0.340)      |
| <i>t</i> -statistic                                             | −3.59                | −3.13                     | −1.14                                 | −3.09                     |
| First-stage <i>F</i> on instruments                             | —                    | 52.66                     | 28.44                                 | 25.78                     |
| Sargan over-id ( <i>p</i> )                                     | —                    | —                         | —                                     | 0.830                     |
| <i>Panel B: English-GDELT spike (endogenous, corroborating)</i> |                      |                           |                                       |                           |
| Coverage coefficient                                            | −1.338***<br>(0.438) | −1.699***<br>(0.565)      | −2.272<br>(1.779)                     | −1.682***<br>(0.563)      |
| <i>t</i> -statistic                                             | −3.05                | −3.01                     | −1.28                                 | −2.99                     |
| First-stage <i>F</i> on instruments                             | —                    | 53.12                     | 12.85                                 | 26.07                     |
| Sargan over-id ( <i>p</i> )                                     | —                    | —                         | —                                     | 0.748                     |
| <i>N</i>                                                        | 51                   | 51                        | 51                                    | 51                        |

*Notes:* Each panel reports the coefficient on the post-event news-attention spike in a regression of per-event  $\beta_{\text{late}}^e$  on attention, with no other covariates. Column (1) is OLS. Columns (2)–(4) are two-stage least squares (2SLS) specifications using predetermined news-pressure instruments. The foreign-disaster instrument is the number of days in the post-event window  $\tau \in [0, 8]$  that fall within the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. The US-disaster instrument counts major natural disasters in a different US state from the ICE event. Column (4) uses both instruments jointly. First-stage statistics are robust Wald *F* tests for the excluded instruments. Heteroskedasticity-robust standard errors in parentheses. Stars: \*  $p < 0.10$ , \*\*  $p < 0.05$ , \*\*\*  $p < 0.01$ .

Table 3 reports the IV estimates against the OLS reference. Foreign-disaster pressure has the expected crowd-out sign in the first stage, and Table 3 reports first-stage *F* statistics for each specification. The second-stage estimates are similar to, and if anything slightly larger in magnitude than, the OLS estimates, consistent with classical measurement error in the attention spike attenuating the OLS association toward zero. The foreign-disaster IV is strong on its own, and the overidentified specification using both disaster instruments gives a similar result, although

the U.S.-disaster instrument is uninformative in isolation. Reported arrests do not predict the foot-traffic response in any specification, so the instrument cannot transmit to the outcome through an arrest channel. The estimate is also unchanged when arrests are included as a control. The IV is consistent with a causal interpretation of the effect of media coverage on foot traffic for mid-visibility events at the margin of the news cycle. These are not the largest, most visible operations. The Supplemental Appendix reports the first-stage coefficients, robustness checks, and balance checks.

## 4.7 Mechanisms

The evidence so far shows that media attention predicts the size of the neighborhood response. I next ask what kind of response it produces. If direct disruption is the main channel, the effect should be strongest close to the worksite and in the raided sector; if perceived risk is the main channel, it should extend across geography and across sectors. When the event study is estimated separately by distance band and by sector, the negative response is not confined to the immediate worksite area, appearing in the 5–15 km and 15–30 km bands and remaining negative though smaller in the 30–50 km band. It is also not concentrated in the raided sector: the same-sector estimate is not negative, while other-sector POIs and both consumer- and nonconsumer-facing POIs decline. The response therefore does not look like a same-sector demand shock or a purely retail channel. These distance- and sector-split estimates appear in the Supplemental Appendix.

Tests of the timing of coverage point the same way: weekly news peaks, national search, and an eight-week English-news area-under-the-curve measure predict the response, while one-day peaks and days-above-threshold measures are weaker. The mechanism evidence is most consistent with a broad perceived-risk response among heavily Latino-immigrant neighborhoods.

## 4.8 Civic vs. Commercial Response

The broad activity types most affected by the enforcement actions can also be identified. Equation (1) is re-estimated separately for two POI bins. The “civic/community” bin includes NAICS-3 codes 712 and 611, which cover parks, museums, and educational services. The “commercial” bin includes NAICS-2 codes 44, 45, and 72, which cover retail and food services. NAICS-3 = 813 (religious organizations and civic associations) is omitted from the civic bin because the category is heavily contaminated by community-bank branches rather than civic organizations.

Twelve events have both a civic and a commercial estimate. Across these events, civic visits in heavily-Latino CBGs fall by an average of 0.120 log points, statistically distinguishable from zero ( $p = 0.02$ ), while the commercial decline is smaller and not distinguishable from zero. The within-event civic-minus-commercial differential averages  $-0.111$  log points and is also statistically significant ( $p = 0.02$ ; Supplemental Appendix). The activity that falls most sharply is therefore institutional contact at schools, parks, and museums, rather than purely commercial activity. The civic sample is small, and I cannot determine whether the gap is itself moderated by attention; the point estimates run opposite to that hypothesis, but the attention subgroups contain only six events each and are not statistically distinguishable. The civic result is therefore reported as an

average contrast rather than an attention-moderated effect. The Supplemental Appendix reports the venue-coverage limits.

## 4.9 Spending

Equation (1) is re-estimated using  $\log(1 + \text{weekly spending})$  from the SafeGraph SpendPatterns panel. Forty-three events have enough high- and low-FB CBG coverage to estimate a spending response. Individual-event estimates are much noisier than for foot traffic, and the aggregate is statistically indistinguishable from zero. Despite the noise, spending falls more in events with larger Spanish-ICE and English-GDELT news spikes: above-median-attention events show a differential decline of about 3 percent in heavily Latino-immigrant neighborhoods over the following eight weeks, while below-median-attention events show none. Because the panel records only card spending at covered merchants, this understates the total economic cost of the response.

## 5 Conclusion

This paper asks why ICE enforcement events generate large neighborhood responses in some cases and little response in others. Across 51 single-worksite events in 26 states, the average response is near zero, but the effect varies far more across events than chance allows, and about 70 percent of the cross-event variation reflects true differences in the response rather than estimation noise. That variation is structured by public attention. Events that become visible in national Spanish-language news, English-language news, and web search generate larger declines in foot traffic in heavily Latino-immigrant neighborhoods, and attention explains roughly 30 percent of the cross-event variation while arrest counts explain essentially none.

Enforcement operates through an information environment as well as through direct detention. A raid disrupts a worksite, but it also sends a public signal about enforcement risk, and whether that signal reaches people outside the worksite depends on news coverage, search, protest, rumor, and the broader attention cycle. The foreign-disaster instrumental-variables design supports this reading for events at the margin of news-cycle access, with out-of-state U.S. disasters a complementary margin. It does not say all enforcement news shocks are identical or that the largest events obey the same margin, but predetermined news pressure moves coverage, and that movement predicts the neighborhood response.

The venue decomposition also changes the welfare interpretation. Commercial activity contracts modestly, and the commercial response is the part most naturally connected to retail-spending projections. The larger response appears in civic and institutional venues that include schools, parks, museums, and related public-facing institutions. Since the foot traffic panel does not cover churches or hospitals, this result is best read as evidence of institutional withdrawal in the venues observed, with an important measurement gap in other civic settings. The cost of enforcement visibility is therefore not only lost transactions. It is also reduced participation in public and institutional life.

Importantly, the most visible events, including the Los Angeles cluster, are unlikely to be IV compliers because their attention spikes are large enough to persist through any reasonable amount of competing coverage. The spending leg of the analysis also has limited power. The venue decomposition is silent on foot traffic associated with religious congregations or healthcare venues, both of which are central to the January 2025 sensitive-locations rescission and to the broader immigration-fear literature. The civic estimate is consequently restricted to schools, parks, and museums. Additionally, the event inventory is built from public sources rather than a complete administrative list.

Taken together, the consequences of enforcement depend not only on the operation but also on the information environment it creates. Media coverage can turn a localized action into a broader signal of risk, and communities respond by reducing activity well beyond the worksite, the raided sector, and ordinary commerce. The same lesson applies wherever a salient event’s economic footprint runs through public attention rather than the event itself.

## **Declaration of generative AI and AI-assisted technologies in the manuscript preparation process**

During the preparation of this work, the author used OpenAI’s ChatGPT and Codex to assist with coding, analysis checks, proofreading, and manuscript editing. After using these tools, the author reviewed and edited the content as needed and takes full responsibility for the content of the article.

## References

- Advan Research. Foot traffic / weekly patterns plus [dataset]. Dewey Data, 2025. doi: 10.82551/C103-N851.
- Umair Ali, Jessica H. Brown, and Chris M. Herbst. Secure communities as immigration enforcement: How secure is the child care market? *Journal of Public Economics*, 233:105101, 2024. doi: 10.1016/j.jpubeco.2024.105101.
- Marcella Alsan and Crystal S. Yang. Fear and the safety net: Evidence from secure communities. *Review of Economics and Statistics*, 106(6):1427–1441, 2024. doi: 10.1162/rest\_a\_01250.
- Catalina Amuedo-Dorantes and Francisca Antman. De facto immigration enforcement, ICE raid awareness, and worker engagement. *Economic Inquiry*, 60(1):373–391, 2022. doi: 10.1111/ecin.13041.
- Catalina Amuedo-Dorantes, Esther Arenas-Arroyo, and Almudena Sevilla. Immigration enforcement and economic resources of children with likely unauthorized parents. *Journal of Public Economics*, 158:63–78, 2018. doi: 10.1016/j.jpubeco.2017.12.004.
- Catalina Amuedo-Dorantes, Esther Arenas-Arroyo, and Chunbei Wang. Is immigration enforcement shaping immigrant marriage patterns? *Journal of Public Economics*, 190:104242, 2020. doi: 10.1016/j.jpubeco.2020.104242.
- Asad L. Asad. *Engage and Evade: How Latino Immigrant Families Manage Surveillance in Everyday Life*. Princeton University Press, 2023. doi: 10.1515/9780691249049.
- Scott R. Baker, Robert A. Farrokhnia, Steffen Meyer, Michaela Pagel, and Constantine Yannelis. How does household spending respond to an epidemic? consumption during the 2020 COVID-19 pandemic. *Review of Asset Pricing Studies*, 10(4):834–862, 2020. doi: 10.1093/rapstu/raaa009.
- Raj Chetty, John N. Friedman, and Michael Stepner. The economic impacts of COVID-19: Evidence from a new public database built using private sector data. *Quarterly Journal of Economics*, 139(2):829–889, 2024. doi: 10.1093/qje/qjad048.
- Alberto Ciancio and Camilo García-Jimeno. Anticipating state action: Risk perceptions and consumption under immigration enforcement. Discussion Paper 18446, IZA Institute of Labor Economics, 2026.
- Elizabeth Cox and Chloe N. East. Labor market impacts of ICE activity in trump 2.0. Working Paper 35129, National Bureau of Economic Research, 2026.
- Uma De Balanzó, Núria Rodríguez-Planas, and Jennifer Roff. Immigration enforcement visibility and consumer spending. Discussion Paper 18620, IZA Institute of Labor Economics, 2026.

- Chloe N. East, Annie Laurie Hines, Philip Luck, Hani Mansour, and Andrea Velasquez. The labor market effects of immigration enforcement. *Journal of Labor Economics*, 41(4):957–996, 2023. doi: 10.1086/721152.
- Thomas Eissensee and David Strömberg. News droughts, news floods, and U.S. disaster relief. *Quarterly Journal of Economics*, 122(2):693–728, 2007. doi: 10.1162/qjec.122.2.693.
- Global Disaster Alert and Coordination System. Gdacs event feed [database]. <https://www.gdacs.org>, 2026. Accessed via GDACS API.
- Austan Goolsbee and Chad Syverson. Fear, lockdown, and diversion: Comparing drivers of pandemic economic decline 2020. *Journal of Public Economics*, 193:104311, 2021. doi: 10.1016/j.jpubeco.2020.104311.
- Exequiel Hernandez. ICE-ing the economy: Immigration enforcement under trump 2.0 and local economic activity. Technical report, SSRN, 2026.
- Jordan Herring and Burt S. Barnow. Indirect effects of immigration enforcement on health care utilization among lawfully present older Hispanics. *Social Science & Medicine*, 384:118540, 2025. doi: 10.1016/j.socscimed.2025.118540.
- Kalev Leetaru and Philip A. Schrodtt. GDELT: Global data on events, location, and tone, 1979–2012. In *International Studies Association Annual Convention*, San Francisco, CA, 2013.
- T. William Lester, Matthew Wilson, and Eli Knaap. The “chilling effect” of ice enforcement: Evidence from high-frequency mobility and spending data in the los angeles region. Working paper, January 26, 2026, 2026. URL <https://bpb-us-e2.wpmucdn.com/sites.uci.edu/dist/6/5832/files/2026/03/Chilling-Effect-of-ICE-in-LA-LWK.pdf>.
- Media Cloud. Media cloud online news archive [database]. <https://search.mediacloud.org>, 2025. Accessed via Python API client (mediacloud>=4.3.0); US Spanish Language Collection, id 196136637.
- SafeGraph. Spend patterns [dataset]. Dewey Data, 2022. doi: 10.82551/NSF5-R186.
- Tara Watson. Inside the refrigerator: Immigration enforcement and chilling effects in Medicaid participation. *American Economic Journal: Economic Policy*, 6(3):313–338, 2014. doi: 10.1257/pol.6.3.313.