

# Attention and the Economic Response to Immigration Enforcement

Greg C. Wright

University of California, Merced  
[gwright4@ucmerced.edu](mailto:gwright4@ucmerced.edu)

July 7, 2026

## Abstract

Does the community response to immigration enforcement depend on the operation itself or on the public attention it receives? Across 51 single-worksite ICE events in 26 states (2018–2026), a per-event difference-in-differences compares heavily Latino-immigrant neighborhoods with similar nearby ones. Events that receive more Spanish-language news coverage see larger foot-traffic declines, while arrest counts explain none of the variation. Two designs support a causal reading. The coverage effect is absent in every pre-event week and emerges only after the event. Competing news pressure from major foreign disasters shifts the coverage an event receives, and the response shifts with it. Instrumented estimates imply that one standard deviation of additional coverage reduces foot traffic in these neighborhoods by about 2.6 percent. The average response is near zero, and attention explains roughly 30 percent of the true cross-event variation. English-language news and search show the same pattern. Civic and institutional visits fall more than commercial activity.

**Keywords:** public attention, immigration enforcement, chilling effects, foot traffic.

**JEL Codes:** J15, R23, K37.

*Acknowledgments:* I thank the UC Merced Community and Labor Center for support and Albert Kochaphum for helpful conversations.

# 1 Introduction

Immigration enforcement can affect the lives of individuals beyond those who are detained. In part, this is because the community response may depend as much on the information environment surrounding an event as on the event itself. A workplace raid that goes largely unnoticed may have little effect beyond the worksite, whereas a smaller raid that becomes highly visible may lead residents to avoid public places. A long literature in economics and sociology documents these “chilling effects,” which can lead to reduced use of healthcare, schooling, and public benefits due to fear of detention [??]. The intensification of U.S. Immigration and Customs Enforcement (ICE) activity under the second Trump administration provides a real-time test. For example, ? document an 8–10 percent foot-traffic decline and a 20–25 percent spending decline in the most heavily Latin American–born neighborhoods of Los Angeles–Orange County in the weeks following ICE’s May 14, 2025 announcement of a major regional enforcement operation.

In this paper I ask whether the observed chilling effects are due to the operations themselves or to their visibility to the surrounding community. If the scale of an operation is central to its effect, arrest counts should help explain which events produce larger responses. On the other hand, if the broader *perception* of risk is what matters, the response will depend on whether news of the event reaches people who were not directly targeted. A raid affects both margins at once, detaining a limited number of individuals while also generating news coverage, social-media activity, rumor, and protest.

I exploit cross-event variation to separate these channels. The sample comprises 51 ICE enforcement events from April 2018 through January 2026 across 26 states, spanning a small cluster of high-profile 2018–2019 worksite actions and a much larger cluster of 2025–2026 operations. For each event I estimate a distributed-lag difference-in-differences (DiD) specification on Advan cellphone foot-traffic data, comparing Census Block Groups (CBGs) in the top decile of Latin-American foreign-born (FB) share to those in the bottom decile.<sup>1</sup> I then extract a summary estimate of the average response. Because an average can hide very different event-level responses, I first measure how much the true response varies. Roughly 70 percent of the cross-event variation reflects genuine differences rather than estimation noise, and the foot traffic response to a new event ranges from a roughly 5 percent decline to little or no response. I then ask which features of the event predict where in that range a neighborhood falls.

The strongest predictor of foot traffic is the public attention garnered by the event, not

---

<sup>1</sup>This is the specification ? use for their single LA event, which is implemented across each event in the sample for comparison.

its operational scale. I exploit post-event spikes in Spanish-language news-volume queries for “ICE”, the agency name most salient to the population at risk. This measure is constructed from a collection of US Spanish-language outlets with daily coverage across the sample period. Events that receive more Spanish-language attention generate larger declines in foot traffic, and the same pattern appears in English-language news and Google search data. In a regression across events that weights each event by the precision of its estimate, the Spanish-language attention spike accounts for roughly 30 percent of the true cross-event variation in responses, while arrest counts account for essentially none. The absence of an effect due to arrests is consistent with the information mechanism.

Two designs support a causal reading. First, the attention effect is absent in every pre-event week and emerges only after the event, and it holds when coverage is measured only in the first two post-event weeks, well before the late-window response it predicts. The early foot-traffic response also does not predict later coverage. Second, news of major foreign disasters, which crowds out public attention to enforcement actions, serves as an instrument for the attention spike, with out-of-state U.S. disasters providing a complementary version. These sources of news pressure strongly predict the amount of attention given to enforcement events and are predetermined with respect to the operation itself.<sup>2</sup> The instrumented estimates imply that one standard deviation of additional coverage reduces foot traffic by about 2.6 percent, slightly larger than the OLS association.

The pattern is robust to several further checks. It holds when attention and arrests are included together, when noisier event estimates are down-weighted, when pre-event media attention is accounted for, and when I estimate within-event changes rather than levels. Balance checks show no systematic relationship with pre-event foot-traffic trends or placebo responses. The effect is stronger within the recent 2025–2026 enforcement wave, suggesting that it does not reflect differences between the two enforcement eras.

National Google Trends and national Spanish-language news both predict the local foot-traffic decline, while metro-level search and local Spanish-language news add little once national attention is accounted for. The response is broad, extending beyond the immediate worksite and the raided industry. I decompose the response by venue and find that foot traffic around schools, parks, and museums falls more than at commercial sites such as retail and food services, a pattern more consistent with a broad withdrawal from public life than a localized demand shock.<sup>3</sup> The welfare cost of visible enforcement is therefore not limited to lost commercial activity but includes reduced participation in civic life.

---

<sup>2</sup>Since competing news shifts attention only for events whose coverage is sensitive to the news cycle, the instrument identifies an effect for mid-visibility events rather than for the largest, most visible operations.

<sup>3</sup>Religious organizations and hospitals are not present in the data. See Section 4.7.

A broader economics literature uses interior-enforcement intensity and federal programs, especially Secure Communities, to estimate effects of immigration enforcement on safety-net participation, family decisions, childcare, healthcare, and labor markets [???????]. A related sociology literature describes the institutional-avoidance strategies of mixed-status families [?]. Recent work studies the current wave. ? link removals and raid awareness to lower immigrant labor-market engagement, ? study consumption responses to learned enforcement risk, and ?, ?, and ? estimate economic effects of the 2025 wave using spending, mobility, and labor-market data. These papers show that enforcement moves economic activity when risk becomes visible or learned. Two features of these designs nevertheless leave the information margin unmeasured. Variation built on program rollouts or enforcement intensity bundles any change in salience together with the change in enforcement itself, and single-episode studies hold the event fixed but leave its visibility without variation. I instead hold the enforcement event fixed and ask why two raids of similar operational severity produce very different neighborhood responses, varying the event’s public visibility rather than the event itself.

These results also add to evidence on how public attention, more than operational scale, shapes the response to risk. A related COVID-mobility literature [???] shows that perceptions of risk changed individual behavior independently of the formal COVID mandates. The case here is similar. An immigration enforcement operation produces both a worksite disruption and public attention, but the community response tracks the attention. The same logic may extend to other salient but partly hidden risks whose economic footprint depends on how widely they are noticed. Where prior work instruments or exploits quasi-experimental variation in enforcement itself, I hold the event fixed and instrument its visibility instead, adapting the news-pressure logic of ? to isolate the attention component of a specific event. To my knowledge, this approach is new to the immigration literature. The event-by-event design also changes what can be learned from a null result. Because the 51 events are estimated separately and then modeled as a distribution, the analysis can distinguish a world in which enforcement rarely moves behavior from one in which it moves behavior strongly but only where it is widely noticed. An average effect cannot make this distinction.

## 2 Data

The analysis combines the event inventory with outcome, exposure, and attention data. Outcome and exposure data include cellphone foot traffic, SafeGraph consumer spending, and American Community Survey (ACS) demographics. Attention and news-pressure data include Google Trends, Mediacloud US Spanish-language and GDELT 2.0 English-language

news volume, Global Disaster Alert and Coordination System (GDACS) disaster alerts, and Nielsen Designated Market Area (DMA) definitions.

## 2.1 Event inventory

A list of single-worksite ICE enforcement events is compiled from news sources, agency press releases, and litigation records. The inventory covers two distinct enforcement waves from April 2018 through August 2019 and January 2025 through January 2026 (the second term’s worksite operations).<sup>4</sup> It restricts attention to events that satisfy four criteria. First, the event date must fall within these two windows. Second, the event must correspond to a specific worksite or commercial address. Third, the worksite must be in one of 26 states for which CBG-level Latin American foreign-born data from the ACS are available. Finally, the worksite lat/lon must be geocoded either to the named business address or, when that is not feasible, to the city or county centroid.<sup>5</sup> Two classes of candidate events fall outside these criteria. Multi-day or metro-wide operations, such as National Guard deployments and sustained urban surges, lack the single discrete enforcement date at a single worksite that the event-time design requires. A lead-agency audit additionally removes actions that were not ICE-led, including one Border Patrol worksite operation and one FBI labor-trafficking investigation in which the workers were treated as victims rather than removal targets, so the sample is cleanly ICE worksite enforcement. These screens identify 55 single-worksite ICE-led events. Four have insufficient catchment point-of-interest (POI) coverage to estimate the event-study specification, leaving a main analysis sample of 51 events spanning 26 states. The Online Appendix summarizes the event-level variables by enforcement wave. The 2018–2019 events report more arrests on average but generate much smaller news spikes than the 2025–2026 events, which motivates the era-robustness check in Section 4.4.

A publicly compiled inventory could be biased toward visible raids. No complete administrative list of single-worksite ICE actions is publicly available for these windows, so the inventory is built from several approaches rather than national news alone: agency releases, litigation records, immigration-policy trackers, state and local reporting, and prior compilations. The result includes a set of events that received little attention. Of the 55 events, 32 are coded medium or low publicity and 29 have an English-GDELT news spike at or below 0.01 (Section 2.6), and the four POI-coverage exclusions include three medium-publicity

---

<sup>4</sup>The first Trump term’s worksite operations included Bean Station TN, Salem OH Fresh Mark, Mt. Pleasant IA Precast, Allen TX CVE Technology, Sumner TX Load Trail, and the August 2019 Mississippi poultry operation, which hit seven plants across five companies on a single day and which is included as a single Morton-MS-anchored compound event.

<sup>5</sup>Four events are explicitly coded as city-centroid locations, one additional event uses the earlier city-query fallback, and one event uses a county-level fallback.

events and one high-publicity event. These facts do not establish a complete universe of obscure raids, but they make it unlikely that the attention effect is mechanically generated by retaining only highly visible events. The Online Appendix reports the sample flow, event enumeration, and source-type audit.

## 2.2 Foot traffic

Weekly visitor counts at point-of-interest level are obtained from Advan Research’s *Foot Traffic/Weekly Patterns Plus* dataset [?], derived from anonymized opt-in mobile devices. Each POI is geocoded to a latitude and longitude, assigned a North American Industry Classification System (NAICS) code, and assigned to its CBG. For each event, all Advan POIs within 50 km of the worksite anchor are collected and assembled into a  $\text{POI} \times \text{week}$  panel spanning  $\pm 12$  weeks of the event date.

The resulting multi-event panel contains approximately 69 million POI-week observations and 3.3 million distinct POI–event pairs across 1.5 million unique POIs. Within each event’s catchment, POI visits are aggregated to CBG-week totals and restricted to CBGs containing at least 10 POIs. A CBG’s POI count is computed from POIs that report visits in at least half of the event’s sample weeks, so short-lived or intermittently reported POIs do not determine which CBGs enter. This restriction concentrates the analysis on retail-active CBGs.

## 2.3 Consumer spending

Spending is measured with SafeGraph’s *SpendPatterns* dataset [?], which reports monthly POI-level aggregates of consumer spending from a large opt-in financial-institution panel. Each row contains an array of daily spending values that are aggregated to weeks, producing a  $\text{POI} \times \text{week}$  spending panel comparable to the foot-traffic data. POI-week values are then aggregated to CBG-week using the same catchment definitions as the foot-traffic panel.

## 2.4 ACS exposure measure

Treatment is defined as CBGs in the top decile of share of residents born in Latin America (“high exposure”), with bottom-decile CBGs as the comparison group. The 2018–2022 5-year ACS provides the Latin-American foreign-born share, computed at the tract level and allocated to all CBGs within each tract.

The deciles are computed within each event’s catchment, so the threshold for “top decile” varies across events. In heavily-Latino metros, the top-decile cutoff is roughly 32 percent foreign-born from Latin America, while in less-Latino metros the threshold is much lower

in absolute terms. This within-event normalization means the cross-event comparison is between “most-Latino” and “least-Latino” CBGs in each respective metro.

## 2.5 Google Trends salience

For each event, salience is measured using weekly Google Trends search interest for three keywords (“ICE raid”, “redada”, “deportación”) at the U.S. national, event-state, and Nielsen Designated Market Area (DMA) levels. Because Trends rescales each query within a geography and timeframe, within-geography pre-vs-post changes are used. The baseline is the four-week average over  $\tau \in [-8, -5]$  relative to the event Monday. Trends are restricted to the 2024–2026 wave because cross-era rescaling is noisy for these low-baseline terms.

## 2.6 News-volume series

**Mediacloud US Spanish-language news.** The primary news-volume series measures attention in the Spanish-language news environment the affected population consumes, on a daily scale consistent across both enforcement waves. The Mediacloud Online News Archive [?] is queried for daily story counts in its curated US Spanish-language collection of 407 sources, including national, state, and local outlets. The headline query is the noun “ICE” restricted to this collection. Counts are normalized by the daily total of all stories in the collection to produce a share-of-coverage time series that is comparable across years. Complementary Spanish keywords are also pulled, and a state-local Spanish-news series is constructed by restricting the same query to sources tagged as published in the event’s home state.

**GDELT English-language US news.** As a corroborating measure on a different platform, language, and methodology, the GDELT 2.0 Document API [?] is queried for daily article counts matching “ICE raid” in US English-language sources, in normalized “Volume Intensity” units, giving a national daily series spanning January 2018 through March 2026. GDELT closely tracks the Mediacloud series at the per-event spike level.

## 2.7 Competing-news instruments

The preferred news-pressure instrument comes from the Global Disaster Alert and Coordination System (GDACS) event feed [?]. Red-alert sudden disasters (earthquakes, tropical cyclones, floods, volcanic eruptions, wildfires) outside the United States, Latin America, and the Caribbean are obtained. The Latin America and Caribbean exclusion is deliberate, since

homeland shocks could affect Latino immigrant communities directly through family ties, remittances, or Spanish-language news demand rather than only through news-cycle competition. For each ICE event, the instrument counts days in the post-event window overlapping the first seven days of such an alert.

A complementary news-pressure measure is also constructed using a hand-coded calendar of major U.S. natural disasters in states other than the ICE event state. The foreign- and U.S.-disaster instruments identify different margins of news-cycle access, so the instrumental-variables (IV) estimates are interpreted as local effects for events whose attention is moved by the corresponding source of competing news pressure.

### 3 Empirical Strategy

#### 3.1 Per-event distributed-lag DiD

For each event  $e$  in the 51-event sample (single-worksite ICE enforcement events from April 2018 through January 2026, see Section 2.1), the same distributed-lag difference-in-differences model is estimated:

$$\log(1 + y_{c,t}^e) = \alpha_c^e + \lambda_t^e + \sum_{\tau \neq -1} \gamma_\tau^e (\text{HighFB}_c^e \times \mathbf{1}[t - T_0^e = \tau]) + \varepsilon_{c,t}^e, \quad (1)$$

where  $y_{c,t}^e$  is weekly visitor count (foot traffic) at CBG  $c$  in week  $t$ , restricted to the catchment of event  $e$ .  $T_0^e$  is the Monday of the week containing event  $e$ 's date. The dummy  $\text{HighFB}_c^e$  equals one for CBGs in the top decile of share Latin American foreign-born within event  $e$ 's catchment and zero for CBGs in the bottom decile (middle deciles are dropped).  $\alpha_c^e$  and  $\lambda_t^e$  are CBG and calendar-week fixed effects, both event-specific because Equation (1) is estimated separately for each event. Standard errors are clustered at the CBG level.

The event-time window is  $\tau \in [-8, +8]$  weeks and all seventeen weeks enter the estimation. The omitted category is  $\tau = -1$ , the last full pre-event week, so each  $\gamma_\tau^e$  measures the high-FB versus low-FB visit gap in week  $\tau$  relative to that reference week rather than a comparison of the event week to all other weeks. The per-event summary estimate  $\beta_{\text{late}}^e$  is the average of  $\gamma_\tau^e$  over the late window  $\tau \in [4, 8]$ . Its standard error is the standard error of that linear combination, computed from the CBG-clustered covariance matrix. Because most events do not fall on a Monday (60 of 66 in the inventory), the  $\tau = 0$  week mixes pre- and post-event days.<sup>6</sup>

---

<sup>6</sup>Re-anchoring  $\tau = 0$  to the first Monday on or after the event date, or dropping the partial event week entirely, leaves the results unchanged. The late-window estimate excludes the partial week by construction,

### 3.2 Identifying variation

The design has two steps. First, Equation (1) uses within-event variation between high-FB and low-FB CBGs in the same catchment. CBG fixed effects absorb time-invariant differences in population, retail density, and baseline foot traffic. Calendar-week fixed effects absorb shocks common to all CBGs in the catchment in a given week. The event estimate is therefore the differential foot-traffic response of CBGs with a high share of Latin-American foreign-born residents relative to low-FB CBGs in the same metro and calendar weeks. Second, the 51 event estimates test whether differences in the attention spike across events predict differences in the estimated response.

The design does not require raids to be unanticipated, and some initiatives are clearly partially anticipated. The initial question is whether the coverage generated by each event explains why otherwise similar raids produce different responses. A later section implements an instrumental-variables design to isolate plausibly exogenous variation in coverage from the surrounding news cycle, and Section 4.2 adds pre-event coverage as a control.

### 3.3 Information-shock and event severity measures

To test whether the chilling response is driven by information or by operational scale, two categories of cross-event measures are constructed. The headline information-shock measure is the post-event spike in the Mediacloud US Spanish-language news-volume series for the query “ICE” (Section 2.6), measured as the share of stories in the 407-source US Spanish-language collection that mention the term. The spike is computed as the peak over  $\tau \in [0, 8]$  minus the baseline over  $\tau \in [-8, -5]$ .

The headline measure is complemented with three corroborating measures: the analogous GDELT national US English-language news-volume spike for “ICE raid,” the national Google Trends spike for “ICE raid,” and the DMA-level Google Trends spike for the local media market in which the event occurred. The two Trends measures are available on the 2024–2026 subsample and are defined as within-geography pre-vs-post changes, since Google Trends rescales each query within a geography and timeframe.

To capture the severity of the enforcement event, I use two measures of underlying operational scale. These are the count of arrests reported by ICE at the worksite,  $\text{Arrests}_e$ , and the same measure in log form,  $\log(1 + \text{Arrests}_e)$ , to accommodate the right-skewed distribution. The arrests measure has coverage of 49 of 51 fitted events.<sup>7</sup>

---

and the attention effect below moves from  $-0.018$  to  $-0.016$  log points per standard deviation under re-anchoring (Online Appendix).

<sup>7</sup>The two events with missing arrest counts in the inventory are Newark NJ (an unnamed business) and Baldwin County AL (an unnamed construction site).

### 3.4 Horse-race specification

To assess whether attention rather than severity predicts chilling, a cross-event ordinary least squares (OLS) regression relates the per-event late-window coefficient to a measure of attention and, in the severity specification, the log-arrests control:

$$\beta_{\text{late}}^e = \pi_0 + \pi_S \text{Spike}_e + \pi_A \log(1 + \text{Arrests}_e) + u_e, \quad (2)$$

where  $\text{Spike}_e$  is the post-event spike in one of the four information-shock measures defined above. Heteroskedasticity-robust (HC1) standard errors are reported. Table 1 reports the headline per-standard-deviation estimates alongside the instrumented versions. The full per-unit specification table, with a standalone arrests specification and single-regressor specifications using English-GDELT and national Google Trends as corroborating measures, appears in the Online Appendix. All specifications use the same event windows. The Trends specification uses the 2024–2026 subsample. Comparing the single-regressor estimates across sources shows whether the same attention effect appears outside the headline Spanish-language measure.

### 3.5 News-pressure instruments

The OLS horse race treats realized post-event attention as the explanatory variable. To isolate variation in attention generated by the surrounding news cycle, predetermined competing-news shocks are used as instruments. The preferred instrument is major foreign disaster news pressure. It counts days in the post-event window that overlap with the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. The exclusion of Latin America and the Caribbean is intended to remove shocks that could directly affect Latino immigrant communities through homeland ties rather than through news-cycle competition only. The identifying assumption is that these shocks move media coverage of the enforcement action without directly changing the relative foot-traffic response in high-FB versus low-FB neighborhoods. Section 4.3 and the Online Appendix report the corresponding balance and robustness checks.

Out-of-state U.S. disasters also serve as a complementary instrument. This measure captures domestic disaster-news pressure in states other than the ICE event state. The foreign- and U.S.-disaster instruments identify different local margins of news-cycle access. The IV estimates should therefore be read as local effects for events whose public attention is shifted by the corresponding source of competing news pressure.


## 4 Results

The response is shown in event time first, pooling the 51 event panels so that the average path, its pre-trends, and the scaling of the response with coverage are all visible. The cross-event analysis then presents the paper’s central estimates, with OLS and instrumented effects of coverage side by side, using competing news pressure from major foreign disasters to isolate variation in the attention an event receives. Timing and design-validity checks address reverse causality and design artifacts, and the remaining sections trace the geography, venues, and spending margins of the response.

### 4.1 The response in event time

Figure 1 pools the 51 event panels and traces the response in event time. Panel A plots the level path  $\gamma_\tau$ , the high-FB versus low-FB gap in weekly visits (foot traffic) for an event of average attention, and Panel B plots the interaction path  $\delta_\tau$ , the additional gap per standard deviation of the Spanish-language coverage spike. Each is shown for the full sample and for the 23 events with no overlapping-catchment contamination and statistically flat pre-trends. Three features stand out. First, the pooled level path shows little average response. The path  $\gamma_\tau$  moves around zero after the event in both samples, and its late-window average is  $-0.003$  log points in each. This pooled estimate is not the main object of the paper, since Section 4.2 studies the unweighted distribution of event-level estimates. Second, the full-sample pre-event level path is not flat, drifting down from about  $+0.013$  at  $\tau = -8$ , while the clean subsample’s pre-event path stays near zero. The full-sample level path is therefore read as descriptive, and a formal sensitivity analysis of the pooled path to pre-trend violations appears in the Online Appendix. Third, the interaction path is flat and near zero through the entire pre-period in both samples and turns negative only after the event, reaching about  $-0.02$  log points per standard deviation of coverage in the late window. The clean-subsample interaction is nearly identical to the full-sample estimate, so the coverage effect is not driven by contaminated catchments or by the events with non-flat pre-trends.

Figure 1: The pooled response in event time




*Notes.* Pooled distributed-lag DiD across the 51 event panels with event  $\times$  CBG and event  $\times$  calendar-week fixed effects, standard errors clustered at the CBG level. High-FB (low-FB) CBGs are those in the top (bottom) decile of Latin-American foreign-born share within the event's catchment. Panel A plots the level path  $\gamma_\tau$ , the high-FB versus low-FB gap in weekly visits (foot traffic) for an event of average attention. Panel B plots the interaction path  $\delta_\tau$ , the additional gap per standard deviation of the Mediacloud Spanish-language ICE coverage spike, standardized at the event level on the full sample. The clean subsample contains the 23 events with no overlapping-catchment contamination and statistically flat pre-trends ( $p > 0.10$ ). Shaded bands are 95 percent confidence intervals. The reference week  $\tau = -1$  is omitted and marked at zero.

Figure 2 estimates the attention effect separately at each event time, regressing the per-event coefficients on the coverage spike with each event weighted by the precision of

its estimate. The attention effect is indistinguishable from zero in every pre-event week and emerges only after the event, stabilizing near  $-0.02$  log points per standard deviation over the late window. The dispersion evidence has the same timing. If the post-event heterogeneity were simply a continuation of pre-existing differences across event catchments, dispersion would look similar at matched distances before and after the reference week. Instead, dispersion is larger after the event at each matched distance (Online Appendix). The heterogeneity that the rest of the paper studies therefore appears after the enforcement event, and its alignment with coverage has no counterpart before the event.

Figure 2: The attention effect over event time



*Notes.* Each point is the coefficient from a precision-weighted regression of the 51 per-event event-study coefficients  $\gamma_{\tau}^e$  on the standardized Spanish-language coverage spike, estimated separately at each event time  $\tau$ , with Knapp–Hartung 95 percent confidence intervals. Each event is weighted by the inverse of its coefficient’s CBG-clustered sampling variance. The reference week  $\tau = -1$  is zero for every event by construction and is marked separately. Estimates are pointwise across  $\tau$ . The companion dispersion analysis and the underlying numbers appear in the Online Appendix.

## 4.2 The attention effect

The average neighborhood response to an enforcement event is close to zero, but this average masks substantial cross-event heterogeneity. After accounting for sampling error in the event-level estimates, about 70 percent of the observed cross-event variance remains as true variation in the underlying response, and the implied prediction interval for a new event runs from a roughly 5.5 percent decline to a 4.8 percent increase. The Online Appendix

reports the full random-effects decomposition and a forest plot of the 51 event estimates. The question is what explains where in that range an event falls.

Table 1 reports the paper’s central estimates. Public attention strongly predicts the variation, and operational scale does not. Column (1) regresses each event estimate on the post-event Spanish-language ICE news spike, weighting each event by the precision of its first-step estimate (a random-effects meta-regression). The spike explains around 30 percent of the between-event variance. Columns (2) and (3) report the OLS counterpart, and columns (4) and (5) instrument the spike with foreign-disaster news pressure (Section 3.5). Across estimators, a one standard deviation larger coverage spike predicts a 1.8 to 2.7 percent larger foot-traffic decline, and the instrumented estimates are, if anything, slightly larger than OLS. Adding the log arrest count leaves the attention coefficient essentially unchanged, while the arrest coefficient is small and statistically insignificant in every specification.<sup>8</sup> The sign of the event estimate is ordered by attention. Negative estimates concentrate among high-attention events and positive estimates among low-attention events. The positive cases are therefore events where little public attention produced little chilling, rather than evidence that enforcement raises local activity.

Table 1: The attention effect on neighborhood foot traffic

|                            | Precision-weighted     | OLS                    |                        | 2SLS (news-pressure IV) |                        |
|----------------------------|------------------------|------------------------|------------------------|-------------------------|------------------------|
|                            | (1)                    | (2)                    | (3)                    | (4)                     | (5)                    |
| Spanish-ICE spike (per SD) | −0.0184***<br>(0.0051) | −0.0222***<br>(0.0062) | −0.0241***<br>(0.0064) | −0.0256***<br>(0.0082)  | −0.0267***<br>(0.0081) |
| log(1 + Arrests)           |                        |                        | −0.0053<br>(0.0042)    |                         | −0.0055<br>(0.0039)    |
| First-stage $F$            |                        |                        |                        | 52.7                    | 51.2                   |
| $R^2_{\text{meta}}$        | 30%                    |                        |                        |                         |                        |
| Events                     | 51                     | 51                     | 49                     | 51                      | 49                     |

*Notes:* The dependent variable is the per-event late-window (weeks 4–8) foot-traffic estimate from Equation (1), in log points. The attention spike is the post-event peak in Mediacloud US Spanish-language ICE coverage, standardized to the 51-event sample (one SD = 0.024 in coverage-share units), so all columns are per standard deviation of coverage. Column (1) is a random-effects meta-regression weighting each event by the inverse sampling variance of its estimate, with Knapp–Hartung standard errors;  $R^2_{\text{meta}}$  is the share of between-event variance in true effects that the spike explains. Columns (2)–(3) are OLS with HC1 standard errors. Columns (4)–(5) instrument the spike with foreign disaster news pressure, the number of post-event days overlapping the first week of a GDACS red-alert disaster outside the United States, Latin America, and the Caribbean (Section 3.5), with robust standard errors. The arrests specifications have  $N = 49$  because two events lack reported arrest counts. Stars: \*  $p < 0.10$ , \*\*  $p < 0.05$ , \*\*\*  $p < 0.01$ .

<sup>8</sup>Arrests are a coarse proxy for operational scale. As a result, a small coefficient does not imply that scale does not matter, only that the dimension captured by arrest counts does not predict the response.

The pattern is not specific to the headline measure, and it survives standard checks. The English-language GDELT news spike has the same sign and statistical significance and explains about 19 percent of the variance, and the US national Google Trends spike gives the same answer on the 2024–2026 subsample (Online Appendix). The effect survives controlling for operation type and sector and is robust to corrections for measurement error in arrests, making it unlikely that the attention effect simply proxies for operational scale. The post-event attention result is not explained by public news attention in the four weeks before the raid. Precision-weighted estimates, event-level bootstrap intervals, pre-event coefficients, and placebo estimates give the same conclusion. The Online Appendix reports these checks, the per-unit specification table, and an early-window table showing that the attention effect builds over the first eight post-event weeks.

### 4.3 Instrumenting attention with news pressure

Realized coverage is not randomly assigned. Features of an operation that make it salient to editors could also make it salient to residents, so the OLS association could overstate the causal effect of coverage. This concern is addressed with an instrumental-variables strategy following the news-pressure logic of ?. In short, the idea is that predetermined national news-cycle conditions can shift attention to an enforcement event without altering the operation itself.

The preferred instrument is major foreign disaster news pressure. It counts the number of days in the post-event window  $\tau \in [0, 8]$  that fall within the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. These disasters are predetermined with respect to the enforcement event and to varying degrees they occupy national news attention. Because the outcome is a within-catchment differential between heavily Latino-immigrant neighborhoods and otherwise similar neighborhoods, a national news shock that moves attention equally in both groups nets out of this differential, so the residual threat is a home-region disaster that differentially affects the treated population. Latin America and the Caribbean are therefore excluded, since disasters there could directly affect Latino immigrant communities through family ties, remittances, or Spanish-language news demand. Out-of-state U.S. disasters are also reported as a complementary margin of domestic disaster-news pressure. These instruments need not identify the same local average treatment effect. Each isolates events whose enforcement attention is moved by a particular source of competing news pressure.

Columns (4) and (5) of Table 1 report the second-stage estimates, and Table 2 reports the first stage. Foreign-disaster pressure has the expected crowd-out sign, with a first-stage

$F$  near 53. The second-stage estimates are similar to, and if anything slightly larger in magnitude than, the OLS estimates, consistent with classical measurement error in the attention spike attenuating the OLS association toward zero. The overidentified specification that adds out-of-state U.S. disasters gives a similar result, although the U.S.-disaster instrument is uninformative in isolation (Online Appendix). Reported arrests do not predict the foot-traffic response in any specification, so the instrument cannot transmit to the outcome through an arrest channel. The estimate is also unchanged when arrests are included as a control. The IV identifies the effect of coverage for mid-visibility events at the margin of the news cycle. These are not the largest, most visible operations. The Online Appendix reports the full per-unit OLS-versus-IV comparison and the robustness and balance checks.

Table 2: First-stage relationship between news pressure and attention

|                                  | Mediacloud Spanish-ICE |                        | English-GDELT          |                        |
|----------------------------------|------------------------|------------------------|------------------------|------------------------|
|                                  | (1) Foreign            | (2) US disaster        | (3) Foreign            | (4) US disaster        |
| Foreign disaster days (std.)     | −0.0168***<br>(0.0023) |                        | −0.0105***<br>(0.0014) |                        |
| Cross-state disaster days (std.) |                        | −0.0078***<br>(0.0015) |                        | −0.0031***<br>(0.0009) |
| First-stage $F$                  | 52.66                  | 28.44                  | 53.12                  | 12.85                  |
| $R^2$                            | 0.482                  | 0.103                  | 0.520                  | 0.046                  |
| $N$                              | 51                     | 51                     | 51                     | 51                     |

*Notes:* The dependent variable is the indicated post-event attention spike. Each column reports a single-instrument first stage with the instrument standardized. The foreign-disaster measure counts days in the ICE event’s post-event window that fall within the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. The US-disaster measure counts major natural disasters in a different US state from the ICE event. Heteroskedasticity-robust standard errors in parentheses. Stars: \*  $p < 0.10$ , \*\*  $p < 0.05$ , \*\*\*  $p < 0.01$ .

#### 4.4 Timing and design validity

Because coverage and foot traffic are measured after the same event, the estimated effect could in principle reflect coverage responding to the neighborhood response rather than the reverse. Timing restricts that account, and Table 3 reports three tests. First, the effect is re-estimated using only coverage from the first two post-event weeks, which predates the late-window response by at least three weeks and so cannot be caused by it. This ordered effect is about half the full-window effect and is essentially unchanged by an arrests control. Second, the reverse direction is tested directly. The early foot-traffic response does not predict late-window coverage conditional on early coverage. Third, a weekly cross-lag

panel with event and event-time fixed effects shows the same asymmetry. Lagged coverage predicts the next week’s response differential, while the lagged response does not predict the next week’s coverage. Temporal precedence rules out feedback from the measured response to the measured coverage, though not every common cause. The arrests control addresses operational severity, the leading candidate common cause, and the news-pressure instrument addresses common causes in the news environment.

Table 3: Temporal ordering of attention and the foot-traffic response

|                                                                                                         |  |          |
|---------------------------------------------------------------------------------------------------------|--|----------|
| <i>Panel A: Cross-event ordering (<math>N = 51</math> events)</i>                                       |  |          |
| Late response ( $\tau \in [4, 8]$ ) on early attention ( $\tau \in [0, 1]$ )                            |  | -0.0103  |
| (KH SE)                                                                                                 |  | (0.0051) |
| KH $p$                                                                                                  |  | 0.0476   |
| ...controlling for arrests (operational scale)                                                          |  | -0.0098  |
| (KH SE)                                                                                                 |  | (0.0055) |
| KH $p$                                                                                                  |  | 0.0775   |
| ...controlling for the early response ( $\tau \in [0, 2]$ )                                             |  | -0.0035  |
| (KH SE)                                                                                                 |  | (0.0035) |
| KH $p$                                                                                                  |  | 0.3152   |
| Late attention ( $\tau \in [4, 8]$ ) on early response ( $\tau \in [0, 2]$ )                            |  | -0.0392  |
| (SE)                                                                                                    |  | (0.1256) |
| $p$                                                                                                     |  | 0.7563   |
| <i>Panel B: Weekly cross-lag, event and event-time FE (<math>N = 459</math> event-weeks, 51 events)</i> |  |          |
| $\gamma_{e,\tau}$ on $\text{attention}_{e,\tau-1}$ (and $\gamma_{e,\tau-1}$ )                           |  | -0.0586  |
| (SE)                                                                                                    |  | (0.0295) |
| $p$                                                                                                     |  | 0.0524   |
| $\text{Attention}_{e,\tau}$ on $\gamma_{e,\tau-1}$ (and $\text{attention}_{e,\tau-1}$ )                 |  | 0.0728   |
| (SE)                                                                                                    |  | (0.0848) |
| $p$                                                                                                     |  | 0.3947   |

*Notes:* Attention is the weekly US Spanish-language Mediacloud ICE coverage ratio aligned to each event’s Monday anchor, as a deviation from the event’s pre-period baseline (mean over  $\tau \in [-8, -5]$ );  $\gamma_{e,\tau}$  is the per-event distributed-lag DiD coefficient (high- vs. low-foreign-born CBGs, reference  $\tau = -1$ ). All regressors and outcomes are standardized, so coefficients are per SD. Panel A rows 1–2 are precision-weighted regressions (random-effects meta-regressions) weighted by the CBG-clustered sampling variance of the late-window mean, with Knapp–Hartung inference; row 3 is OLS with HC1 standard errors. Panel B pools post-event weeks ( $\tau \in [0, 8]$ ) with event and event-time fixed effects and standard errors two-way clustered by event and ISO calendar week. Where the estimated  $\gamma$  appears as a regressor, sampling noise attenuates its coefficient toward zero, so the early-response control is imperfect. Temporal precedence rules out backward-in-time causation, not common causes; the arrests control addresses the leading common-cause candidate (operational severity). The Panel B lagged-dependent-variable specifications with unit fixed effects in a short panel are subject to Nickell bias.

The heterogeneity is not an artifact of the event-study design varying across contexts. Per-event design quality is uncorrelated with the attention spike. I measure design quality by the flatness of pre-event trends, the number of block groups in the catchment, the precision of the estimate, and exposure to overlapping nearby events. The attention effect is unchanged when these design characteristics are added as controls (Table 4). Restricting the sample to the 34 events with no overlapping-catchment contamination, and further to the 23 events with both no contamination and statistically flat pre-trends, leaves the share of variation attributable to true cross-event differences near 60 percent and preserves the attention effect, although precision falls in the smallest subsample. The effect is also a within-wave phenomenon rather than an artifact of the difference between the 2018–2019 and 2025–2026 enforcement eras. Estimated on the 2025–2026 wave alone, which contains 44 of the 51 events, it is if anything stronger. As Section 4.1 shows, the pooled level path is descriptive rather than identifying since its late-window estimate is sensitive to modest departures from parallel pre-trends.

Table 4: Design validity: per-event design quality is orthogonal to attention

| <i>Panel A: Design metric regressed on the standardized Spanish-ICE spike (HC1)</i>                          |     |           |        |           |                     |
|--------------------------------------------------------------------------------------------------------------|-----|-----------|--------|-----------|---------------------|
| Design metric                                                                                                | $N$ | slope/SD  | $t$    | $p$       | $R^2$               |
| Pre-trend joint $F$ (non-flatness)                                                                           | 51  | -217.6642 | -1.04  | 0.301     | 0.047               |
| pre-event coef                                                                                               | 51  | -0.0056   | -1.31  | 0.191     | 0.045               |
| placebo response                                                                                             | 51  | -0.0060   | -1.25  | 0.210     | 0.035               |
| # CBGs in catchment                                                                                          | 51  | -9.3960   | -0.24  | 0.810     | 0.001               |
| Per-event SE (imprecision)                                                                                   | 51  | -0.0005   | -0.11  | 0.915     | 0.000               |
| # contaminating events                                                                                       | 51  | 0.1667    | 1.04   | 0.296     | 0.025               |
| Frac. catchment shared                                                                                       | 51  | 0.0256    | 0.51   | 0.611     | 0.004               |
| <i>Panel B: Attention coefficient in the precision-weighted regression, with vs. without design controls</i> |     |           |        |           |                     |
|                                                                                                              | $N$ | coef/SD   | SE     | $t$       | $p$                 |
| Spanish-ICE spike (per SD), no controls                                                                      | 51  | -0.0184   | 0.0051 | -3.65     | 0.001               |
| + design controls                                                                                            | 51  | -0.0185   | 0.0054 | -3.42     | 0.001               |
| <i>Panel C: Heterogeneity on clean subsamples</i>                                                            |     |           |        |           |                     |
| Subsample                                                                                                    | $N$ | $I^2$     | $\tau$ | spike $p$ | $R^2_{\text{meta}}$ |
| All events                                                                                                   | 51  | 69.7%     | 0.0287 | 0.001     | 30.1%               |
| No contamination                                                                                             | 34  | 61.1%     | 0.0309 | 0.027     | 14.8%               |
| No contam. & flat pre-trends                                                                                 | 23  | 62.8%     | 0.0301 | 0.051     | 16.6%               |

*Notes:* Panel A regresses each per-event design-quality metric on the standardized Spanish-ICE news spike with HC1 standard errors; a null slope means design quality does not vary systematically with attention. Panel B shows the precision-weighted regression attention coefficient is unchanged when standardized design metrics (pre-trend  $F$ , # CBGs, # contaminating events) are added as controls. Panel C recomputes the heterogeneity statistics and the attention precision-weighted regression on progressively cleaner subsamples (no overlapping-catchment contamination; additionally non-rejected flat pre-trends at  $p > 0.10$ ).  $\tau$  is in log points.

## 4.5 National vs. local attention

Holding the search platform and query fixed and varying only geography, national Google Trends predicts larger foot-traffic declines than metro-level (DMA) Trends, and DMA Trends adds little once national attention is included. National rather than local attention appears to carry the signal, and reported arrests remain null. The national-vs.-DMA Google Trends horse race for the 2024–2026 subsample appears in the Online Appendix.

## 4.6 Mechanisms

The evidence so far shows that media attention predicts the size of the neighborhood response. I next ask what kind of response it produces. If direct disruption is the main channel,

the effect should be strongest close to the worksite and in the raided sector. If perceived risk is the main channel, it should extend across geography and across sectors. When the event study is estimated separately by distance band and by sector, the negative response is not confined to the immediate worksite area, appearing in the 5–15 km and 15–30 km bands and remaining negative though smaller in the 30–50 km band. It is also not concentrated in the raided sector. The same-sector estimate is not negative, while other-sector POIs and both consumer- and nonconsumer-facing POIs decline. The response therefore does not look like a same-sector demand shock or a purely retail channel. Table 5 reports the distance- and sector-split estimates.

Table 5: Mechanisms

| Sample                                                              | Events | Mean estimate | Coef. on Spanish-ICE spike |
|---------------------------------------------------------------------|--------|---------------|----------------------------|
| <i>Panel A. Geographic reach, mutually exclusive distance rings</i> |        |               |                            |
| 0–5 km from worksite                                                | 12     | +0.031        | −0.022<br>(0.017)          |
| 5–15 km from worksite                                               | 29     | −0.013        | −0.026***<br>(0.007)       |
| 15–30 km from worksite                                              | 36     | +0.002        | −0.023*<br>(0.012)         |
| 30–50 km from worksite                                              | 40     | −0.009        | −0.009<br>(0.006)          |
| <i>Panel B. Same-sector and broad-sector activity</i>               |        |               |                            |
| Same broad sector as raided worksite                                | 20     | +0.003        | +0.012<br>(0.017)          |
| Other sectors                                                       | 49     | −0.011        | −0.022***<br>(0.007)       |
| Consumer-facing POIs                                                | 51     | −0.007        | −0.018***<br>(0.006)       |
| Nonconsumer POIs                                                    | 48     | −0.003        | −0.029***<br>(0.008)       |

*Notes.* Each row reports a cross-event OLS regression of a subgroup-specific late-window estimate on the standardized Mediacloud Spanish-ICE attention spike. Standard errors are heteroskedasticity-robust and shown in parentheses. Panel A defines high-FB and low-FB CBGs once within the full 50 km event catchment and then estimates the event-study separately by distance ring. Panel B estimates the event-study separately by POI bucket. Same-sector POIs have the same broad NAICS-2 code as the raided worksite. Consumer-facing POIs are NAICS-2 codes 44, 45, 72, and 81. Stars denote \*  $p < 0.10$ , \*\*  $p < 0.05$ , \*\*\*  $p < 0.01$ .

Tests of the timing of coverage point the same way. Weekly news peaks, national search, and an eight-week English-news area-under-the-curve measure predict the response, while

one-day peaks and days-above-threshold measures are weaker. The mechanism evidence is most consistent with a broad perceived-risk response among heavily Latino-immigrant neighborhoods.

## 4.7 Civic vs. Commercial Response

The broad activity types most affected by the enforcement actions can also be identified. Equation (1) is re-estimated separately for two POI bins. The “civic/community” bin includes NAICS-3 codes 712 and 611, which cover parks, museums, and educational services. The “commercial” bin includes NAICS-2 codes 44, 45, and 72, which cover retail and food services. NAICS-3 = 813 (religious organizations and civic associations) is omitted from the civic bin because the category is heavily contaminated by community-bank branches rather than civic organizations.

Twelve events have both a civic and a commercial estimate. Across these events, civic visits in heavily-Latino CBGs fall by an average of 0.120 log points, while the commercial decline is smaller and not distinguishable from zero. The within-event civic-minus-commercial differential averages  $-0.111$  log points and is statistically significant. The activity that falls most sharply is therefore institutional contact at schools, parks, and museums, rather than purely commercial activity. The civic sample is small, and I cannot determine whether the gap is itself moderated by attention. The point estimates run opposite to that hypothesis, but the attention subgroups contain only six events each and are not statistically distinguishable. The civic result is therefore reported as an average contrast rather than an attention-moderated effect. The Online Appendix reports the inference and venue-coverage limits.

## 4.8 Spending

Equation (1) is re-estimated using  $\log(1 + \text{weekly spending})$  from the SafeGraph SpendPatterns panel. Forty-three events have enough high- and low-FB CBG coverage to estimate a spending response. Individual-event estimates are much noisier than for foot traffic, and the aggregate is statistically indistinguishable from zero. Despite the noise, spending falls more in events with larger Spanish-ICE and English-GDELT news spikes. Above-median-attention events show a differential decline of about 3 percent in heavily Latino-immigrant neighborhoods over the following eight weeks, while below-median-attention events show none. Because the panel records only card spending at covered merchants, this understates the total economic cost of the response.

## 5 Conclusion

This paper asks why ICE enforcement events generate large neighborhood responses in some cases and little response in others. Across 51 single-worksite events in 26 states, the average response is near zero, but the effect varies far more across events than chance allows, and about 70 percent of the cross-event variation reflects true differences in the response rather than estimation noise. That variation is structured by public attention. Events that become visible in national Spanish-language news, English-language news, and web search generate larger declines in foot traffic in heavily Latino-immigrant neighborhoods, and attention explains roughly 30 percent of the cross-event variation while arrest counts explain essentially none.

Enforcement operates through an information environment as well as through direct detention. A raid disrupts a worksite, but it also sends a public signal about enforcement risk, and whether that signal reaches people outside the worksite depends on news coverage, search, protest, rumor, and the broader attention cycle. The foreign-disaster instrumental-variables design supports this reading for events at the margin of news-cycle access, with out-of-state U.S. disasters serving as a complementary margin. It does not say all enforcement news shocks are identical or that the largest events obey the same margin, but predetermined news pressure moves coverage, and that movement predicts the neighborhood response.

The venue decomposition also changes the welfare interpretation. Commercial activity contracts modestly. This is the part of the response most naturally connected to retail-spending projections. The larger response appears in civic and institutional venues that include schools, parks, museums, and related public-facing institutions. Since the foot traffic panel does not cover churches or hospitals, this result is best read as evidence of institutional withdrawal in the venues observed, with an important measurement gap in other civic settings. The cost of enforcement visibility is therefore not only lost transactions. It is also reduced participation in public and institutional life.

The most visible events, including the Los Angeles cluster, are unlikely to be IV compliers because their attention spikes are large enough to persist through any reasonable amount of competing coverage. The spending analysis also has limited power. The venue decomposition is silent on foot traffic associated with religious congregations or healthcare venues, both of which are central to the January 2025 sensitive-locations rescission and to the broader immigration-fear literature. The civic estimate is consequently restricted to schools, parks, and museums. Additionally, the event inventory is built from public sources rather than a complete administrative list.

Taken together, the consequences of enforcement depend not only on the operation but

also on the information environment it creates. Media coverage can turn a localized action into a broader signal of risk, and communities respond by reducing activity well beyond the worksite, the raided sector, and ordinary commerce. The same lesson applies wherever a salient event’s economic footprint runs through public attention rather than the event itself.

## **Declaration of generative AI and AI-assisted technologies in the manuscript preparation process**

During the preparation of this work, the author used OpenAI’s ChatGPT and Codex to assist with coding, analysis checks, proofreading, and manuscript editing. After using these tools, the author reviewed and edited the content as needed and takes full responsibility for the content of the article.

## References