

Information Shocks and ICE Enforcement

Greg C. Wright

University of California, Merced

gwright4@ucmerced.edu

May 27, 2026

Abstract

ICE enforcement actions reduce foot traffic in some immigrant neighborhoods, but the response is far from uniform. This paper asks why activity falls sharply after some enforcement events and little after others. We study 51 single-worksite ICE enforcement events across 26 U.S. states from 2018 to 2026. For each event, we estimate a difference-in-differences model comparing foot traffic in heavily Latino-immigrant neighborhoods with otherwise similar neighborhoods before and after the event. The size of the response varies widely across events and is predicted by public attention, not by operation size. A post-event spike in US Spanish-language news coverage of ICE strongly predicts larger declines in foot traffic, while the number of arrests has no predictive power. English-language news and search-volume measures show the same pattern. An instrumental-variables design using news crowd-out from major foreign disasters isolates the effect due to attention on enforcement actions. We find that once national attention is accounted for, local attention adds little. Commercial activity contracts modestly, while civic and institutional visits to schools, parks, museums, and related venues fall more. The findings imply that when enforcement actions attract media attention, agencies can limit negative community spillovers by clearly communicating the scope and limits of those actions.

Keywords: immigration enforcement, chilling effects, foot traffic.

JEL Codes: J15, R23, K37.

Preliminary and incomplete. Comments welcome.

Acknowledgments: I thank the UC Merced Community and Labor Center for support and Amada Armenta and Albert Kochaphum for helpful conversations.

1 Introduction

Immigration enforcement can impact the lives of individuals beyond those who are detained. In part, this is because the community response may depend as much on the information environment surrounding an event as on the event itself. A workplace raid that goes largely unnoticed may have little effect beyond the worksite, whereas a smaller raid that becomes highly visible may lead residents to avoid public places. A long literature in economics and sociology documents these “chilling effects,” which can lead to reduced use of healthcare, schooling, and public benefits due to fear of detention [Watson, 2014, Asad, 2023]. The intensification of U.S. Immigration and Customs Enforcement (ICE) activity under the second Trump administration provides a striking real-time test. For example, Lester et al. [2026] document an 8–10 percent foot-traffic decline and a 20–25 percent spending decline in the most heavily Latin American-born neighborhoods of Los Angeles–Orange County in the weeks following ICE’s May 14, 2025 announcement of a major regional enforcement operation.

This raises a question that analysis of a single event cannot answer. Is the chilling effect due to the operation itself, or its visibility to the surrounding community? To the extent that direct disruption is important, arrests and operation size will best predict the community impact. If it is the broader *perception* of risk that is important, then the impact will depend on whether news of the event reaches people who were not directly targeted. An ICE raid clearly affects both margins at once. It detains a limited number of individuals at a specific site while also generating news coverage, social-media activity, rumor, and protest.

This paper uses cross-event variation to separate these channels. We assemble 51 ICE enforcement events from April 2018 through January 2026 spanning 26 states. The sample includes worksite raids from both the first and second Trump administrations including a small cluster of high-profile 2018–2019 actions and a much larger cluster of 2025–2026 worksite operations. For each event we estimate a distributed-lag difference-in-differences specification using Advan cellphone foot-traffic data, comparing Census Block Groups (CBGs) in the top decile of Latin-American foreign-born (FB) share to those in the bottom decile, and we extract a summary estimate of the average foot-traffic response.¹ We then explore the features of the event that predict those responses.

The strongest predictor of foot traffic is the public attention garnered by the event, not its operational scale. We exploit post-event spikes in Spanish-language news-volume queries for “ICE”, constructed from a collection of US Spanish-language outlets with daily coverage across the sample period. Events that receive more Spanish-language attention generate larger declines in foot traffic, and the same pattern appears in English-language news and Google search data. Arrest counts, by contrast, do not predict the subsequent size of the neighborhood response.

The pattern survives several checks that sharpen the interpretation. It holds when attention and arrests are included together, when noisier event estimates are down-weighted, when we ac-

¹This is the specification used by Lester et al. [2026] for their single LA event, which we implement across each of our events for comparison.

count for pre-event media attention, and when we focus on within-event changes rather than levels. We also exploit news of major foreign disasters, which crowd out public attention on enforcement actions, as an instrument for the attention spike, with out-of-state U.S. disasters providing a complementary version. These sources of news pressure strongly predict the amount of attention given to enforcement events and are plausibly unrelated to the enforcement process itself. Balance checks show no systematic relationship with pre-event foot-traffic trends or placebo responses, and the estimates are robust to a variety of controls.

Both national Google search Trends and national Spanish-language news predict the local decline in foot traffic, while metro-level search and local Spanish-language news add little once national attention is accounted for. We also show that the effect of this coverage is quite broad. Specifically, we show that the community response extends beyond the immediate worksite and beyond the industry of the raided site.

We then decompose the response by venue type and show that both commercial activity and participation in public and institutional life are affected. We find that foot traffic around schools, parks, museums, and historical sites falls more than foot traffic at commercial sites (e.g., retail and food services).² The welfare cost of visible enforcement is therefore not limited to lost commercial activity but also includes reduced participation in civic life.

An important policy implication is that ambiguity itself can be costly. If residents respond to the attention garnered by enforcement as a signal of broader risk, then enforcement agencies and local officials can reduce harmful spillovers by communicating the scope of an operation, what it does and does not imply, and whether routine public and institutional settings remain outside the action.

A broader economics literature uses changes in interior-enforcement intensity and federal enforcement programs, especially Secure Communities, to estimate the effects of immigration enforcement on safety-net participation, family decisions, childcare markets, healthcare use, and labor-market outcomes [Watson, 2014, Amuedo-Dorantes et al., 2018, 2020, Alsan and Yang, 2024, Ali et al., 2024, Herring and Barnow, 2025, East et al., 2023]. This work shows that enforcement can alter behavior and economic outcomes among people who are not themselves detained. A related sociology literature describes the everyday mobility and institutional-avoidance strategies of mixed-status Latino families [Asad, 2023]. Our paper differs by holding the enforcement event fixed and studying the public signal it generates. The object is not whether a larger enforcement regime changes behavior, but why two raids with similar operational severity can produce very different neighborhood responses.

Several recent papers are related. Amuedo-Dorantes and Antman [2022] show that ICE removals and raid awareness reduce immigrant labor-market engagement. Ciano and García-Jimeno [2026] study consumption responses to learned enforcement risk, while De Balanzó et al. [2026], Hernandez [2026], and Cox and East [2026] estimate economic effects of the 2025 enforcement wave using spending, mobility, and labor-market data. These papers show that enforcement can move economic

²Religious organizations and hospitals are not present in the data, see Section 4.6.

activity when risk becomes visible or learned. Our paper asks a different but complementary question. Conditional on a single-worksite enforcement event occurring, does the public signal around that event explain why some neighborhoods withdraw much more than others? This shifts the object from enforcement exposure itself to the information environment surrounding otherwise similar enforcement events.

Our contribution is to connect that enforcement literature to evidence on perceived risk and public attention. The recent COVID-mobility literature [Goolsbee and Syverson, 2021, Baker et al., 2020, Chetty et al., 2024] establishes that perceived health risk drove substantial economic activity changes independently of formal mandates. We bring this logic into the immigration-enforcement domain, finding that the relevant proxy for perceived risk is what the community is actually paying attention to, not what the underlying operation looks like to the agency carrying it out. Methodologically, prior work instruments for or uses quasi-experimental variation in enforcement itself. The news-pressure instrument we use instead isolates the attention channel within a given enforcement event, an approach borrowed from Eisensee and Strömberg [2007] in the media-effects literature and, to our knowledge, new to this immigration-enforcement literature.

The remainder of the paper proceeds as follows. Section 2 describes the foot-traffic, spending, news and trends, and event-inventory data. Section 3 describes the empirical strategy. Section 4 presents the main results. Section 5 concludes.

2 Data

The analysis combines nine data sources. These are an inventory of single-worksite ICE enforcement events spanning the first and second Trump administrations, cellphone foot-traffic data, SafeGraph consumer-spending data, American Community Survey (ACS) demographics, Google Trends search interest, Mediacloud US Spanish-language news-volume data, GDELT 2.0 English-language news-volume data, Global Disaster Alert and Coordination System (GDACS) disaster alerts, and Nielsen Designated Market Area (DMA) definitions. We describe each in turn.

2.1 Event inventory

We compile a list of single-worksite ICE enforcement events from news sources, agency press releases, and litigation records. The inventory covers two distinct enforcement waves from April 2018 through August 2019 and January 2025 through January 2026 (the second term’s worksite operations).³ The inventory restricts attention to events that satisfy four criteria. First, the event date must fall within these two windows. Second, the event must correspond to a specific worksite or commercial address. Third, the worksite must be in one of 26 states for which we have CBG-level Latin American foreign-born data from the ACS. Finally, the worksite lat/lon must be geocoded

³The first Trump term’s worksite operations included Bean Station TN, Salem OH Fresh Mark, Mt. Pleasant IA Precast, Allen TX CVE Technology, Sumner TX Load Trail, and the August 2019 Mississippi poultry operation, which hit seven plants across five companies on a single day and which we include as a single Morton-MS-anchored compound event.

either to the named business address or, when that is not feasible, to the city or county centroid.⁴ These screens identify 55 single-worksite ICE-led events. Four have insufficient catchment POI coverage to estimate the event-study specification, leaving a main analysis sample of 51 events spanning 26 states.

A natural concern is that a public event inventory could be selected toward visible raids. The ideal sampling frame would be a complete administrative list of single-worksite ICE enforcement actions, but no such list is publicly available for these two windows. We therefore built the inventory from several discovery channels rather than national news alone, including agency releases, litigation records, immigration-policy trackers, state and local reporting, and prior event compilations. The resulting single-worksite inventory includes a substantial low-attention tail. Of the 55 single-worksite ICE-led events, 32 are coded medium or low publicity and 29 have an English-GDELT news spike at or below 0.01 (Section 2.6). Main-sample attrition is also not concentrated among visible cases. The four POI-coverage exclusions include three medium-publicity events and one high-publicity event. These facts do not establish a complete universe of obscure raids, but they make it unlikely that the headline attention gradient is mechanically generated by retaining only highly visible events. Appendix Table 10 reports the source-type audit.

Table 1 summarizes the sample flow from the single-worksite ICE-led inventory to the estimable event panel and to the subsamples used in each robustness check. The same 51-event panel is the unit of analysis for the foot-traffic results throughout the paper. Smaller numbers reflect data-availability restrictions on specific covariates. The complete event-level enumeration is provided in Appendix Table 9.

⁴Four events are explicitly coded as city-centroid locations, one additional event uses the earlier city-query fallback, and one event uses a county-level fallback.

Table 1: Sample-flow from full event inventory to estimation subsamples

Definition	Events
Single-worksite ICE-led events satisfying inventory screens	55
– Single-worksite ICE-led but no estimable late-window estimate (insufficient catchment POI coverage) ^a	–4
Main analysis sample (single-worksite ICE-led with estimable late-window estimate)	51
<i>Estimation subsamples:</i>	
Events with reported arrest counts (severity test)	49
Events in 2024–2026 subsample (Google Trends national)	44
Events with DMA-level Google Trends coverage	41
Events with at least one in-state Mediacloud Spanish-language source	31
Events with estimable per-event civic estimate	11

Notes: “Events” is the unit. The 51-event panel is the main analysis sample for the per-event distributed-lag DiD and the headline cross-event horse race. Smaller numbers below reflect availability of specific second-stage covariates rather than re-estimation. The Mediacloud Spanish-language national series and the English-GDELT series are defined for all 51 events. Trends, DMA Trends, and in-state Spanish sources cover smaller subsamples (see Sections 2.5 and 2.6). The civic estimate requires sufficient civic-NAICS POI density in the high-FB (foreign-born) CBGs of an event’s catchment. Only 11 of 51 events meet this threshold (see Section 4.6).

2.2 Foot traffic

We obtain weekly visitor counts at point-of-interest (POI) level from Advan Research’s *Foot Traffic/Weekly Patterns Plus* dataset [Advan Research, 2025], derived from anonymized opt-in mobile devices. Each POI is geocoded to a latitude and longitude (lat/lon), assigned a NAICS code, and assigned to the CBG within which its lat/lon falls. For each event, we collect all Advan POIs within a 50 km radius of the event’s worksite anchor and assemble a POI $\times$ week panel spanning ± 12 weeks of the event date. POIs whose nearest CBG centroid is more than 3 km away (typically in states outside our ACS coverage) are dropped.

The resulting multi-event panel contains approximately 69 million POI-week observations and 3.3 million distinct POI–event pairs across 1.5 million unique POIs. Within each event’s catchment, we aggregate POI visits to CBG-week totals and restrict to CBGs containing at least 10 POIs. This restriction concentrates the analysis on retail-active CBGs.

2.3 Consumer spending

To measure spending, we use SafeGraph’s *SpendPatterns* dataset [SafeGraph, 2022], which reports monthly POI-level aggregates of consumer spending from a large opt-in financial-institution panel. Each row contains an array of daily spending values that we aggregate to weeks, producing a POI $\times$ week spending panel comparable to the foot-traffic data. We then aggregate POI-week to CBG-week using the same catchment definitions as the foot-traffic build.

2.4 ACS exposure measure

We define treatment as CBGs in the top decile of share of residents born in Latin America, which we designate as “high exposure”, while CBGs in the bottom decile serve as the comparison group. We use the 2018–2022 5-year ACS and calculate the foreign-born share from Latin America at the tract level and allocate it to all CBGs within each tract.

The deciles are computed within each event’s catchment, so the threshold for “top decile” varies across events. In heavily-Latino metros, the top-decile cutoff is roughly 32 percent foreign-born from Latin America, while in less-Latino metros the threshold is much lower in absolute terms. This within-event normalization means our cross-event comparison is between “most-Latino” and “least-Latino” CBGs in each respective metro.

2.5 Google Trends salience

For each event, we measure public salience using weekly Google Trends search interest for three keywords (“ICE raid”, “redada”, “deportación”) at three geographies. These are the U.S. national level, the event’s state, and its Nielsen Designated Market Area (DMA). Trends rescales each query to peak-100 within a single geography and timeframe, so cross-geography magnitudes are not directly comparable in levels. We therefore use within-geography pre-vs-post changes (peak minus baseline) rather than absolute units. The pre-event baseline window is the four-week average over $\tau \in [-8, -5]$ weeks relative to the event Monday. We restrict the Trends measures to the 2024–2026 enforcement wave since 2018–2019 events are too noisy. The 2018–2019 events therefore use GDELT (Section 2.6) for the attention measure, and the Trends-based correlations and Spanish-keyword robustness are reported on the 2024–2026 subsample. DMA-level Trends are available for a 30-DMA subset of those events (small media markets including Biloxi–Gulfport, MS and Burlington–Plattsburgh, VT returned insufficient search volume). The three geographies provide a natural ladder for testing where the chilling-relevant attention signal lives.

2.6 News-volume series

We use two daily news-volume series spanning both enforcement waves on a consistent scale. The primary series measures attention in the Spanish-language news environment that the affected immigrant population consumes. A complementary series measures attention in the English-language US news supply.

Mediacloud US Spanish-language news. We query the Mediacloud Online News Archive [Media Cloud, 2025] for daily story counts in its curated US Spanish-language collection (407 sources, including national outlets such as Univision, Telemundo, La Opinión, El Diario, and EFE alongside state and local Spanish-language publications). The Mediacloud directory is the most actively maintained inventory of US Spanish-language online news and has dense daily coverage back to January 2018 (95–250 indexed stories per day for our headline query). The headline query

is the noun “ICE” restricted to this collection. We normalize by the daily total of all stories in the collection to produce a share-of-coverage time series directly comparable across years. We also pull complementary Spanish keywords (“redada”, “deportación”, “deportaciones”, “redada migratoria”, “agentes de ICE”) and report robustness across these terms. In addition to the national Spanish-language series, we construct a state-local Spanish-news series for each event by restricting the same query to the Mediacloud sources tagged as published in the event’s home state.

GDELT English-language US news. As a corroborating measure constructed on a different platform, language, and methodology, we use the GDELT 2.0 Global Knowledge Graph project’s Document API [Leetaru and Schrod, 2013]. We query the API for daily article counts matching the phrase “ICE raid” restricted to US English-language news sources, in normalized “Volume Intensity” units (article-count divided by daily total monitored articles). The result is a single national daily time series spanning January 2018 through March 2026. GDELT closely tracks the Spanish-language Mediacloud series at the per-event spike level. We report GDELT-based numbers throughout as a robustness check.

2.7 Competing-news instruments

To construct news-pressure instruments, we use three predetermined event calendars. The preferred instrument uses the Global Disaster Alert and Coordination System (GDACS) event feed [Global Disaster Alert and Coordination System, 2026]. We obtain red-alert sudden disasters, including earthquakes, tropical cyclones, floods, volcanic eruptions, and wildfires, that occur outside the United States, Latin America, and the Caribbean. The Latin America and Caribbean exclusion is deliberate because homeland shocks could directly affect Latino immigrant communities through family ties, remittances, or Spanish-language news demand rather than only through national news-cycle competition. For each ICE event, the instrument counts the number of days in the post-event window that overlap with the first seven days of one of these foreign-disaster alerts.

We also construct a complementary news-pressure measure using a hand-coded calendar of major U.S. natural disasters in states other than the ICE event state. The foreign- and U.S.-disaster instruments identify different margins of news-cycle access, so we interpret the IV estimates as local effects for events whose attention is moved by the corresponding source of competing news pressure.

3 Empirical Strategy

3.1 Per-event distributed-lag DiD

For each event e in our 51-event sample (single-worksite ICE enforcement events from April 2018 through January 2026, see Section 2.1), we estimate the same distributed-lag difference-

in-differences model:

$$\log(1 + y_{c,t}^e) = \alpha_c^e + \lambda_t^e + \sum_{\tau \neq -1} \gamma_\tau^e (\text{HighFB}_c^e \times \mathbf{1}[t - T_0^e = \tau]) + \varepsilon_{c,t}^e, \quad (1)$$

where $y_{c,t}^e$ is weekly visitor count (foot traffic) at CBG c in week t , restricted to the catchment of event e . T_0^e is the Monday of the week containing event e 's date. The dummy HighFB_c^e equals one for CBGs in the top decile of share Latin American foreign-born within event e 's catchment and zero for CBGs in the bottom decile (middle deciles are dropped). α_c^e and λ_t^e are CBG and calendar-week fixed effects, both event-specific because Equation (1) is estimated separately for each event. Standard errors are clustered at the CBG level. The event-time window is $\tau \in [-8, +8]$ weeks.

3.2 Identifying variation

The design exploits two layers of variation, stacked. First, we exploit variation within each event, specifically high-FB vs. low-FB CBGs in the same catchment. Equation (1) is the within-catchment difference-in-differences specification we use. The CBG fixed effect α_c^e absorbs everything time-invariant about a CBG (population, retail density, baseline foot traffic level). The calendar-week fixed effect λ_t^e absorbs everything that hits all CBGs in the catchment in a given week, including the event-week itself and any national news cycle that affects all CBGs in the catchment equally. What survives the two-way fixed effects is the *differential* foot-traffic response of heavily-Latin-American-foreign-born CBGs versus the otherwise-similar low-FB CBGs in the same metro in the same calendar weeks. The chilling effect is identified entirely off heterogeneity by Latin-American foreign-born share within the catchment. Each event yields a late-window estimate.

Next, we exploit variation across events, specifically variation in the magnitude of the estimated foot-traffic response to enforcement news. We use the 51 event estimates and regress them on the event-level attention measures described in Section 3.4. The variation here is cross-event differences in how big an attention spike each enforcement event generated. A Camarillo Glass House operation produces a far larger national Spanish-language news spike than a Mt. Pleasant IA Precast raid, and the second stage asks whether that cross-event dispersion in spike size predicts cross-event dispersion in the estimate.

3.3 Information-shock vs. event severity measures

To test whether the chilling response is driven by information or by operational scale, we construct two categories of cross-event measures. First, we construct four complementary information-shock proxies. The headline measure is the post-event spike in the Mediacloud US Spanish-language news-volume series for the query "ICE" (Section 2.6), measured as the share of all stories in the 407-source US Spanish-language collection that mention the term in a given week. This measure has dense daily coverage on a single consistent share-of-coverage scale across our full 2018–2026

panel and isolates attention in the Spanish-language news environment that the affected immigrant population consumes. We compute the spike on the same windows as the other measures (peak over $\tau \in [0, 8]$ minus baseline over $\tau \in [-8, -5]$).

Our cross-event design does not require every raid to be a surprise. Some enforcement initiatives may be partially anticipated, especially broad regional operations or operations that unfold over several days. The question is whether the realized coverage generated by each event explains why otherwise similar raids produce different neighborhood responses. We therefore interpret the OLS horse race as a heterogeneity and mechanism exercise, and use the instrumental-variables design below to isolate plausibly exogenous variation in the media coverage that is generated by the surrounding news cycle. Section 4.2 reports specifications that add the immediate pre-event spike over $\tau \in [-4, -1]$ as a control.

We complement the headline measure with three corroborating measures. The first is the analogous spike in GDELT 2.0’s national US English-language news-volume series for the query “ICE raid” (Section 2.6). The second and third are two Google Trends measures on the 2024–2026 subsample. These measures differ from the news-volume measures since they come from a different source and are based on revealed search demand rather than journalist-produced coverage. The first Trends measure is the post-event national Google Trends spike.

$$\text{Spike}_e^{\text{Nat}} = \max_{\tau \in [0, 8]} \text{Trends}_{e, \tau}^{\text{Nat}} - \frac{1}{4} \sum_{\tau = -8}^{-5} \text{Trends}_{e, \tau}^{\text{Nat}}, \quad (2)$$

where $\text{Trends}_{e, \tau}^{\text{Nat}}$ is national-US weekly search interest for “ICE raid” in the week τ relative to the Monday of event e . The series has complete coverage for all 44 events in the 2024–2026 subsample. The second measure, $\text{Spike}_e^{\text{DMA}}$, is identical but is the *post-event Google Trends spike* at the DMA level, the local TV/media market in which the event occurred.

Since Google Trends rescales each query within each geography and timeframe, both Trends spikes are within-geography pre-vs-post changes such that cross-event magnitudes within a single column are directly comparable across 2024–2026 events. We do not extend the Trends measures back to 2018 since those years are noisy (in part because ICE-raid search interest was near zero during the Biden interval, making cross-era rescaling unreliable). The 2018–2019 events therefore use only the Spanish-language Mediachannel and English-language GDELT news-volume measures, which use share-of-coverage units that are stable across the panel. The national and DMA Trends spikes are positively correlated, and the two news-volume measures co-move strongly across the full 51-event panel, confirming that the same event-level news cycle propagates through both information environments. The Mediachannel and national-Trends measures also co-move strongly on the 2024–2026 subsample where both are defined.

With respect to the severity of the enforcement event, we use two measures of underlying operational scale. These are the count of arrests reported by ICE at the worksite, Arrests_e , and the same measure in log form, $\log(1 + \text{Arrests}_e)$, to accommodate the right-skewed distribution. The

arrests measure has coverage of 49 of 51 fitted events.⁵

3.4 Horse-race specification

To assess whether attention rather than severity predicts chilling, we run a cross-event OLS regression of the per-event late-window coefficient on a measure of attention and, in the severity specification, the log-arrests control:

$$\beta_{\text{late}}^e = \pi_0 + \pi_S \text{Spike}_e + \pi_A \log(1 + \text{Arrests}_e) + u_e, \quad (3)$$

where Spike_e is the post-event spike in one of the four information-shock measures defined above. Heteroskedasticity-robust (HC1) standard errors are reported. The columns of Table 2 report the headline Mediacloud Spanish-ICE specification alone and with the arrests control on the full 51-event panel, a standalone arrests specification, and single-regressor specifications using English-GDELT and national Google Trends as corroborating measures on the same windows. The Trends specification uses the 2024–2026 subsample. Comparing the single-regressor estimates across sources shows whether the same attention gradient appears outside the headline Spanish-language measure.

3.5 News-pressure instruments

The OLS horse race treats realized post-event attention as the explanatory variable. To isolate variation in attention generated by the surrounding news cycle, we use predetermined competing-news shocks as instruments. The preferred instrument is major foreign disaster news pressure. It counts days in the post-event window that overlap with the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. The exclusion of Latin America and the Caribbean is intended to remove shocks that could directly affect Latino immigrant communities through homeland ties rather than through news-cycle competition only. The identifying assumption is that these shocks move enforcement action media coverage without directly changing the relative foot-traffic response in high-FB versus low-FB neighborhoods. Section 4.4 and Appendix C report the corresponding balance and robustness checks.

We also use out-of-state U.S. disasters as a complementary instrument. This measure captures domestic disaster-news pressure in states other than the ICE event state. The foreign- and U.S.-disaster instruments identify different local margins of news-cycle access. The IV estimates should therefore be read as local effects for events whose public attention is shifted by the corresponding source of competing news pressure.

⁵The two events with missing arrest counts in our inventory are Newark NJ (an unnamed business) and Baldwin County AL (an unnamed construction site).

3.6 Overlap robustness

We run a contamination-robustness restriction. For each event e , we compute the number of *other* events in our sample whose date falls within ± 8 weeks of e and whose CBG catchment overlaps with e 's. Events with zero such contaminating neighbors are designated “clean.” Thirty-four of 51 events satisfy this restriction. We re-estimate the horse race on the clean subsample as a robustness check.

3.7 Spending

We re-estimate Equation (1) on the SafeGraph SpendPatterns weekly spending panel. The per-event spending estimates are an order of magnitude noisier than the foot-traffic estimates (median SE 0.142 vs. approximately 0.03 for foot traffic), and the cross-event sample mean is statistically indistinguishable from zero. We treat the spending result as a power-limited sensitivity rather than a primary finding. Details are in Section 4.7.

4 Results


We begin with the event-level estimates from the distributed-lag design and use OLS to ask which event features predict larger neighborhood responses. We then implement the instrumental-variables strategy using variation in news volume from major foreign disasters, with out-of-state U.S. disasters as a complementary margin. These news shocks crowd out, and therefore attenuate, the volume of news associated with enforcement actions.

4.1 Event chilling dispersion

We begin with the cross-event distribution of late-window (weeks 4 through 8) estimates for foot traffic. Across our 51 events, 29 produce negative point estimates (57 percent) and 22 produce positive estimates. The cross-event mean is -0.008 log points. The median is -0.003 , and the mean is not statistically distinguishable from zero across events. After Benjamini-Hochberg multiple-testing correction, 17 of 51 events have a statistically significant effect and 11 of those are negative (chilling) and 6 are positive. The 11 negative events cluster at high national-attention spikes. Their median Spanish-ICE spike is 0.060 and their median English-GDELT spike is 0.034. The 6 positive events cluster at lower spikes, with medians of 0.020 and 0.006 respectively.⁶ Figure 1 shows how the event estimates relate to the English-GDELT attention spikes.

⁶Formal distribution tests reject equality between the two groups for both the Spanish-ICE and English-GDELT measures.

Figure 1: Per-event late-window estimate vs English-GDELT national news spike, by BH- $q < 0.05$ significance group



Notes. Each marker is one of the 51 fitted events. Blue circles denote events with a BH-corrected per-event significance of $q < 0.05$ and a negative late-window estimate (chilling). Red squares denote BH-significant events with a non-negative late-window estimate. Grey dots denote non-significant events. English-GDELT spikes are higher among the negative-significant group than the positive-significant group. The equivalent figure built on the headline Mediacloud Spanish-ICE measure (Section 4.2) gives the same qualitative pattern.

4.2 The information-shock horse race

Table 2 reports the main horse-race result. Columns (1)–(2) anchor on the Mediacloud US Spanish-language news-volume spike for the query “ICE” (407-source Spanish-language collection, Section 2.6). Column (3) reports the standalone arrests specification. Columns (4)–(5) replace the headline measure with two corroborating attention proxies built on different platforms. These are the English-language GDELT national news-volume spike and the US national Google Trends spike for “ICE raid.”

Table 2: Information shock horse race. Late-window estimate on attention and severity

	(1) Spa. ICE	(2) + arr.	(3) Arr. only	(4) GDELT only	(5) Trends only
MC Spa. ICE spike	−0.9193*** (0.2559)	−0.9987*** (0.2649)			
English GDELT spike				−1.3382*** (0.4382)	
National Trends spike $\times 10^{-3}$					−0.5254*** (0.1576)
$\log(1 + \text{Arrests}_e)$		−0.0053 (0.0042)	−0.0028 (0.0049)		
Constant	+0.0193** (0.0096)	+0.0386** (0.0184)	+0.0015 (0.0172)	+0.0114 (0.0088)	+0.0110 (0.0086)
N	51	49	49	51	44
R^2	0.227	0.260	0.006	0.173	0.215

Notes. OLS regressions of the per-event late-window estimate on the post-event attention spike ($\text{peak}[\tau \in [0, 8]]$ minus $\text{baseline}[\tau \in [-8, -5]]$, event-week aggregation) and $\log(1 + \text{Arrests}_e)$. The headline measure in columns (1)–(2) is the Mediacloud US Spanish-language news ICE-spike (407-source collection, share-of-coverage). Column (3) reports the standalone severity specification. Columns (4)–(5) report the same single-regressor attention specification with English-GDELT (full 51-event panel) and US national Google Trends ‘ICE raid’ (44-event 2024–2026 subsample) as corroborating attention measures. The headline measure closely comoves with the corroborating measures, so the single-regressor attention columns should be read as cross-source robustness rather than independent channels. Heteroskedasticity-robust (HC1) standard errors in parentheses. Stars denote * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

In short, a larger Spanish-ICE spike predicts a larger decline in foot traffic. Adding $\log(1 + \text{Arrests}_e)$ leaves the Spanish-ICE coefficient essentially unchanged, while arrests alone explain very little of the cross-event variation. English-GDELT and national Google Trends show the same negative impact on foot traffic.

Table 3: Early-window robustness

	Weeks 0–2	Weeks 1–3	Weeks 4–8
<i>Panel A. Event-estimate distribution</i>			
Mean estimate	−0.001	−0.004	−0.008
Median estimate	0.000	−0.003	−0.003
Share negative	0.51	0.55	0.57
<i>Panel B. Cross-event attention gradient</i>			
Pearson r , Spanish-ICE spike	−0.263	−0.375	−0.477
OLS coef., Spanish-ICE spike	−0.323*	−0.524***	−0.919***
	(0.174)	(0.191)	(0.256)
OLS coef., Spanish-ICE spike, with arrests	−0.388**	−0.592***	−0.999***
	(0.181)	(0.202)	(0.265)
<i>Panel C. Corroborating English-news measure</i>			
Pearson r , English-GDELT spike	−0.227	−0.333	−0.416
OLS coef., English-GDELT spike	−0.466	−0.774**	−1.338***
	(0.302)	(0.333)	(0.438)
N	51	51	51

Notes. Each column summarizes the same per-event distributed-lag DiD path over a different post-event window. Week 0 is the calendar week containing the enforcement event and can include pre-event days. The weeks 1–3 window excludes the event week. The weeks 4–8 column is the headline persistent-response window used in Table 2. OLS rows report single-regressor HC1 standard errors in parentheses, except for the “with arrests” row, which also includes $\log(1 + \text{Arrests}_e)$ as a control. The arrests-control regressions use 49 events. Stars denote * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 3 repeats the exercise across three post-event windows. The negative association is already detectable within the first three post-event weeks but grows monotonically in magnitude, statistical significance, and across-event R^2 as the event window extends to weeks 4–8. The same build-up appears for the English-GDELT measure, consistent with a chilling response that accumulates rather than spikes immediately at the event date.

Several robustness checks support our interpretation of the observed effect. Spanish-language news, English-language news, and national search measures move together across events. The post-event attention result is not explained by public news attention in the four weeks before the raid. Precision-weighted estimates and event-level bootstrap intervals give the same conclusion. The average pre-event coefficient over weeks −8 through −2 is +0.012 log points, compared with a late-window mean of −0.008. Finally, a placebo estimate constructed from pre-event weeks averages +0.001, with correlations of +0.15 and +0.17 with the Spanish-ICE and English-GDELT attention spikes. These checks suggest that the main pattern is not a continuation of a common pre-event decline in high-FB neighborhoods. Appendix B reports these checks. Appendix Figure 2 also shows a pooled event-time interaction specification.

4.3 National vs. local attention

We next ask whether the primary effect of media coverage on foot traffic is national or local. Table 4 reports the levels version of the same-source national-vs.-DMA Google Trends horse race for the 2024–2026 subsample. This comparison holds the search platform and query fixed and varies only the geography of measured attention.

Table 4: Local vs. National Attention to Enforcement Events

	(1) Nat.	(2) + arr.	(3) + DMA	(4) DMA	(5) DMA + arr.
National Trends spike $\text{Spike}_e^{\text{Nat}}$	−0.000525*** (0.000158)	−0.000564*** (0.000167)	−0.000621*** (0.000228)		
DMA Trends spike $\text{Spike}_e^{\text{DMA}}$			+0.000074 (0.000215)	−0.000326* (0.000174)	−0.000330* (0.000184)
$\log(1 + \text{Arrests}_e)$		−0.001275 (0.005263)			+0.000589 (0.006522)
Constant	+0.010989 (0.008648)	+0.015012 (0.017686)	+0.012604 (0.008953)	+0.006794 (0.008799)	+0.005279 (0.021078)
N	44	42	41	41	39
R^2	0.215	0.232	0.239	0.085	0.079

Notes: OLS regressions of the per-event late-window estimate on national and DMA-level Google Trends spikes for “ICE raid” and, in columns (2) and (5), $\log(1 + \text{Arrests}_e)$. Both Trends measures are post-event spikes, defined as peak search interest over $\tau \in [0, 8]$ minus the baseline average over $\tau \in [-8, -5]$. The Google Trends sample is restricted to the 2024–2026 events for which weekly search data are internally consistent. The DMA measure is missing for a small set of events whose media markets have insufficient search volume. Heteroskedasticity-robust (HC1) standard errors in parentheses. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

National Trends predicts larger declines in foot traffic and remains negative when DMA Trends is included. DMA Trends has the expected negative sign when entered alone, but it adds little once national attention is accounted for. Operational severity remains null in the same exercise.

4.4 Instrumenting attention with news pressure

The relationship between media coverage and foot traffic raises the concern that because both are measured after the same enforcement action, they may be confounded. We address this with an instrumental-variables strategy following the news-pressure logic of [Eisensee and Strömberg \[2007\]](#). In short, the idea is that predetermined national news-cycle conditions can shift attention to an enforcement event without altering the operation itself.

The preferred instrument is major foreign disaster news pressure. It counts the number of days in the post-event window $\tau \in [0, 8]$ that fall within the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. These disasters are predetermined with respect to the enforcement event and to varying degrees they occupy national news attention. We exclude Latin America and the Caribbean because shocks in those regions

could directly affect Latino immigrant communities through family ties, remittances, or Spanish-language news demand. We also report out-of-state U.S. disasters as a complementary margin of domestic disaster-news pressure. These instruments need not identify the same local average treatment effect. Each isolates events whose enforcement attention is moved by a particular source of competing news pressure.

Table 5: First-stage relationship between news pressure and attention

	Mediacloud Spanish-ICE		English-GDELT	
	(1) Foreign	(2) US disaster	(3) Foreign	(4) US disaster
Foreign disaster days (std.)	−0.0168*** (0.0023)		−0.0105*** (0.0014)	
Cross-state disaster days (std.)		−0.0078*** (0.0015)		−0.0031*** (0.0009)
First-stage F	52.66	28.44	53.12	12.85
R^2	0.482	0.103	0.520	0.046
N	51	51	51	51

Notes: The dependent variable is the indicated post-event attention spike. Each column reports a single-instrument first stage with the instrument standardized. The foreign-disaster measure counts days in the ICE event’s post-event window that fall within the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. The US-disaster measure counts major natural disasters in a different US state from the ICE event. Heteroskedasticity-robust standard errors in parentheses. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 5 reports the first stage directly. Foreign-disaster pressure has the expected crowd-out sign. Events whose post-event windows overlap more with major foreign disasters receive smaller enforcement-attention spikes in both Spanish-language and English-language news. Out-of-state U.S. disasters give the same sign.

Table 6: Instrumental-variables estimation of the attention $\rightarrow$ chilling channel

	(1) OLS	(2) Foreign disasters	(3) US disasters (cross-state)	(4) Both IVs (over-id)
<i>Panel A: Mediacloud Spanish-ICE spike (endogenous)</i>				
Attention coefficient	−0.919*** (0.256)	−1.059*** (0.338)	−0.911 (0.796)	−1.052*** (0.340)
<i>t</i> -statistic	−3.59	−3.13	−1.14	−3.09
First-stage <i>F</i> on instruments	—	52.66	28.44	25.78
Sargan over-id (<i>p</i>)	—	—	—	0.830
<i>Panel B: English-GDELT spike (endogenous, corroborating)</i>				
Attention coefficient	−1.338*** (0.438)	−1.699*** (0.565)	−2.272 (1.779)	−1.682*** (0.563)
<i>t</i> -statistic	−3.05	−3.01	−1.28	−2.99
First-stage <i>F</i> on instruments	—	53.12	12.85	26.07
Sargan over-id (<i>p</i>)	—	—	—	0.748
<i>N</i>	51	51	51	51

Notes: Each panel reports the coefficient on the post-event news-attention spike in a regression of per-event β_{late}^e on attention, with no other covariates. Column (1) is OLS. Columns (2)–(4) are 2SLS specifications using predetermined news-pressure instruments. The foreign-disaster instrument is the number of days in the post-event window $\tau \in [0, 8]$ that fall within the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. The US-disaster instrument counts major natural disasters in a different US state from the ICE event. Column (4) uses both instruments jointly. First-stage statistics are robust Wald *F* tests for the excluded instruments. Heteroskedasticity-robust standard errors in parentheses. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 6 reports the IV estimates against the OLS reference. The second-stage estimates are similar to the OLS estimates. The foreign-disaster IV is strong on its own, and the overidentified specification using both disaster instruments gives a similar result. The IV is consistent with a causal interpretation of the effect of media coverage on foot traffic for events whose attention is mitigated by competing news pressure. Appendix C reports the robustness and balance checks.

4.5 Mechanisms

The evidence so far shows that media attention predicts the size of the neighborhood response to enforcement actions. We next ask what kind of response that attention produces. We start by further testing the main hypothesis: if direct disruption is the main channel, the attention effect should be strongest close to the worksite and in the same broad sector as the raided employer; if perceived risk is the main channel, the response should extend across a wider geography and across sectors.

Table 7: Mechanism diagnostics

Sample	Events	Mean estimate	Coef. on Spanish-ICE spike
<i>Panel A. Geographic reach, mutually exclusive distance rings</i>			
0–5 km from worksite	12	+0.031	−0.022 (0.017)
5–15 km from worksite	29	−0.013	−0.026*** (0.007)
15–30 km from worksite	36	+0.002	−0.023* (0.012)
30–50 km from worksite	40	−0.009	−0.009 (0.006)
<i>Panel B. Same-sector and broad-sector activity</i>			
Same broad sector as raided worksite	20	+0.003	+0.012 (0.017)
Other sectors	49	−0.011	−0.022*** (0.007)
Consumer-facing POIs	51	−0.007	−0.018*** (0.006)
Nonconsumer POIs	48	−0.003	−0.029*** (0.008)

Notes. Each row reports a cross-event OLS regression of a subgroup-specific late-window estimate on the standardized Mediacloud Spanish-ICE attention spike. Standard errors are heteroskedasticity-robust and shown in parentheses. Panel A defines high-FB and low-FB CBGs once within the full 50 km event catchment and then estimates the event-study separately by distance ring. Panel B estimates the event-study separately by POI bucket. Same-sector POIs have the same broad NAICS-2 code as the raided worksite. Consumer-facing POIs are NAICS-2 codes 44, 45, 72, and 81. Stars denote * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Panel A of Table 7 estimates the event-study separately by distance from the worksite. The inner-ring estimates are noisy because only 12 events have enough high- and low-FB CBGs within 5 km. Even so, the negative effect is not confined to the immediate worksite area. It appears in the 5–15 km and 15–30 km bands and remains negative, although smaller, in the 30–50 km band. In cumulative-radius versions, the coefficient is −0.021 within 15 km, −0.030 within 30 km, and −0.023 in the full 50 km catchment.

Panel B asks whether the response is concentrated in the raided sector. Same-sector estimates are available for only 20 events, largely because agricultural, industrial, and construction worksites often have few visitor-facing POIs in the same broad sector. In this sample, the effect is not negative. In contrast, other-sector POIs show a negative gradient close to the headline estimate. Consumer-facing POIs show the expected negative gradient, but nonconsumer POIs do as well. The response therefore does not look like a same-sector demand shock or a purely retail and restaurant channel.

Tests of the timing of news coverage point in the same direction. Weekly news peaks, national search, and an eight-week English-news area-under-the-curve measure all predict the response, while one-day peaks and simple days-above-threshold measures are weaker. These measures co-

move across events, so we cannot cleanly separate peak salience from sustained salience. Still, the pattern suggests that the relevant object is the post-event news cycle over weeks, not just a single daily headline. Taken together, the mechanism evidence is most consistent with a broad perceived-risk response among heavily Latino-immigrant neighborhoods.

4.6 Civic vs. Commercial Response

We can also identify the broad activity types that are most affected by the enforcement actions. We re-estimate Equation (1) separately for two POI bins. The “civic/community” bin includes NAICS-3 codes 712 and 611, which cover parks, museums, and educational services. The “commercial” bin includes NAICS-2 codes 44, 45, and 72, which cover retail and food services. We omit NAICS-3 = 813 (religious organizations and civic associations) from the civic bin because the category is heavily contaminated by community-bank branches rather than civic organizations.

Table 8: Late-window estimate by POI category $\times$ baseline tier

Baseline	POI bin	n	Mean estimate	Share negative
Low	Civic/community ($n_3 \in \{611, 712\}$)	5	−0.029	80%
Low	Commercial ($n_2 \in \{44, 45, 72\}$)	19	−0.007	63%
High	Civic/community	6	−0.137	83%
High	Commercial	21	+0.001	62%
<i>Within-event civic-minus-commercial differential</i>				
Low	Within-event diff	5	−0.026	—
High	Within-event diff	6	−0.136	—

Notes: Per-event distributed-lag DiD on the expanded 51-event sample, with the visitor outcome restricted to POIs in the civic or commercial bin. Baseline tier defined by state-Trends median split ($\text{Sal}_e^{\text{state}} \geq \text{median}$); events without state-Trends coverage (2018–2019 events and events in states not in the Trends DMA panel) are excluded from the tier-based grouping. Civic bin: parks and museums (NAICS-3 = 712), educational services (611). Commercial bin: retail (NAICS-2 = 44, 45) and food services (72). NAICS-3 = 813 (religious organizations and civic associations) is excluded from the civic bin because a desk-research audit of CA POIs found that 110 of 118 Advan POIs in 813 are labeled “Independent Community Bankers of America” but in fact correspond to individual community bank branch addresses (true NAICS = 522, banking) rather than civic organizations; see Appendix F. Only events with ≥ 3 POIs in the civic bin per CBG can support estimation; the resulting civic subsample contains 11 events (5 low-baseline, 6 high-baseline). The commercial bin requires ≥ 8 POIs per CBG and covers a larger subsample.

The civic result is necessarily small-sample because only 11 events have enough civic POI coverage to support estimation. Among the six high-baseline events with a civic estimate, the average civic estimate is −0.137 and the average within-event civic-minus-commercial differential is −0.136. In the five low-baseline civic events, the corresponding civic estimate is −0.029 and the within-event differential is −0.026. Civic visits in heavily-Latino CBGs fall more than commercial visits in the same event catchments, although the civic sample is small. Thus, the activity that falls most sharply is institutional contact at schools, parks, museums, and similar civic venues,

rather than purely commercial activity. Appendix E reports the venue-coverage limits and the small-sample checks on whether the civic differential itself varies with the type of media coverage.

4.7 Spending

We re-estimate Equation (1) using $\log(1 + \text{weekly spending})$ from the rebuilt daily-unpacked SafeGraph SpendPatterns panel. The spending panel contains 46 of the 51 main-sample events with any SafeGraph rows, and 43 have enough high- and low-FB CBG coverage to estimate the event-level spending response. Individual-event spending estimates are much noisier than foot-traffic estimates, reflecting thinner POI coverage. The aggregate cross-event spending estimate is statistically indistinguishable from zero in all sample restrictions.

Despite the noise, the spending estimates move in the same qualitative direction as the foot-traffic estimates. Spending responses are more negative in events with larger Spanish-ICE news spikes, and similarly for English-GDELT spikes.

5 Conclusion

This paper asks why ICE enforcement events generate large neighborhood responses in some cases and little response in others. Across 51 single-worksite events in 26 states, I find that events that become visible in national Spanish-language news, English-language news, and web search generate larger declines in foot traffic in heavily Latino-immigrant neighborhoods. Arrest counts do not explain the same variation once the attention on the event is measured directly.

The interpretation is that enforcement operates through an information environment as well as through direct detention. A raid disrupts a particular worksite, but it also sends a public signal about enforcement risk. Whether that signal reaches people outside the worksite depends on news coverage, search activity, protest, rumor, and the broader attention cycle around the event. The instrumental-variables design based on major foreign disasters supports this reading for events at the margin of news-cycle access, with out-of-state U.S. disasters providing a complementary margin. The design does not say that all enforcement news shocks are identical, or that the largest events in the sample are governed by the same margin. It shows that predetermined variation in news pressure moves enforcement coverage, and that this movement predicts the neighborhood response.

The venue decomposition also changes the welfare interpretation. Commercial activity contracts modestly, and the commercial response is the part most naturally connected to retail-spending projections. The larger response appears in civic and institutional venues that include schools, parks, museums, and related public-facing institutions. Since the foot traffic panel does not cover churches or hospitals, this result is best read as evidence of institutional withdrawal in the venues we observe, with an important measurement gap in other civic settings. The cost of enforcement visibility is therefore not only lost transactions. It is also reduced participation in public and institutional life.

Importantly, the most visible events, including the Los Angeles cluster, are unlikely to be

IV compliers because their attention spikes are large enough to persist through any reasonable amount of competing coverage. The spending leg of the analysis also has limited power. The venue decomposition is silent on foot traffic associated with religious congregations or healthcare venues, both of which are central to the January 2025 sensitive-locations rescission and to the broader immigration-fear literature. The civic estimate is consequently restricted to schools, parks, and museums. Additionally, the event inventory is built from public sources rather than a complete administrative list.

Taken together, the results show that the consequences of enforcement depend not only on the operation itself, but also on the information environment it creates. Media coverage can turn a localized enforcement action into a broader signal of risk, and communities respond by reducing activity in ways that extend beyond the worksite, the raided sector, and ordinary commercial activity. This suggests a practical policy implication. Enforcement agencies and local officials cannot control the news cycle, but they can reduce avoidable spillovers by communicating clearly about the scope of an operation, what it does and does not imply, and whether routine civic and institutional settings remain outside the action.

References

- Advan Research. Foot traffic / weekly patterns plus [dataset], 2025.
- Umair Ali, Jessica H. Brown, and Chris M. Herbst. Secure communities as immigration enforcement: How secure is the child care market? *Journal of Public Economics*, 233:105101, 2024. doi: 10.1016/j.jpubeco.2024.105101.
- Marcella Alsan and Crystal S. Yang. Fear and the safety net: Evidence from secure communities. *Review of Economics and Statistics*, 106(6):1427–1441, 2024. doi: 10.1162/rest_a.01250.
- Catalina Amuedo-Dorantes and Francisca Antman. De facto immigration enforcement, ICE raid awareness, and worker engagement. *Economic Inquiry*, 60(1):373–391, 2022. doi: 10.1111/ecin.13041.
- Catalina Amuedo-Dorantes, Esther Arenas-Arroyo, and Almudena Sevilla. Immigration enforcement and economic resources of children with likely unauthorized parents. *Journal of Public Economics*, 158:63–78, 2018. doi: 10.1016/j.jpubeco.2017.12.004.
- Catalina Amuedo-Dorantes, Esther Arenas-Arroyo, and Chunbei Wang. Is immigration enforcement shaping immigrant marriage patterns? *Journal of Public Economics*, 190:104242, 2020. doi: 10.1016/j.jpubeco.2020.104242.
- Asad L. Asad. *Engage and Evade: How Latino Immigrant Families Manage Surveillance in Everyday Life*. Stanford University Press, 2023.
- Scott R. Baker, Robert A. Farrokhnia, Steffen Meyer, Michaela Pagel, and Constantine Yannelis. How does household spending respond to an epidemic? consumption during the 2020 COVID-19 pandemic. *Review of Asset Pricing Studies*, 10(4):834–862, 2020.
- Raj Chetty, John N. Friedman, and Michael Stepner. The economic impacts of COVID-19: Evidence from a new public database built using private sector data. *Quarterly Journal of Economics*, 139(2):829–889, 2024.
- Alberto Ciancio and Camilo García-Jimeno. Anticipating state action: Risk perceptions and consumption under immigration enforcement. Discussion Paper 18446, IZA Institute of Labor Economics, 2026.
- Elizabeth Cox and Chloe N. East. Labor market impacts of ICE activity in trump 2.0. Working Paper 35129, National Bureau of Economic Research, 2026.
- Uma De Balanzó, Núria Rodríguez-Planas, and Jennifer Roff. Immigration enforcement visibility and consumer spending. Discussion Paper 18620, IZA Institute of Labor Economics, 2026.

- Chloe N. East, Annie Laurie Hines, Philip Luck, Hani Mansour, and Andrea Velasquez. The labor market effects of immigration enforcement. *Journal of Labor Economics*, 41(4):957–996, 2023. doi: 10.1086/721152.
- Thomas Eisensee and David Strömberg. News droughts, news floods, and U.S. disaster relief. *Quarterly Journal of Economics*, 122(2):693–728, 2007.
- Global Disaster Alert and Coordination System. Gdacs event feed [database]. <https://www.gdacs.org>, 2026. Accessed via GDACS API.
- Austan Goolsbee and Chad Syverson. Fear, lockdown, and diversion: Comparing drivers of pandemic economic decline 2020. *Journal of Public Economics*, 193:104311, 2021.
- Exequiel Hernandez. ICE-ing the economy: Immigration enforcement under trump 2.0 and local economic activity. Technical report, SSRN, 2026.
- Brad Herring and Burt S. Barnow. Spillover effects of intensified immigration enforcement on healthcare use among likely-authorized Hispanic immigrants. *Working Paper*, 2025.
- Kalev Leetaru and Philip A. Schrodtt. GDELT: Global data on events, location, and tone, 1979–2012. In *International Studies Association Annual Convention*, San Francisco, CA, 2013.
- T. William Lester, Matthew Wilson, and Eli Knaap. The “chilling effect” of ice enforcement: Evidence from high-frequency mobility and spending data in the los angeles region. Working paper, January 26, 2026, 2026.
- Media Cloud. Media cloud online news archive [database]. <https://search.mediacloud.org>, 2025. Accessed via Python API client (`mediacloud>=4.3.0`); US Spanish Language Collection, id 196136637.
- SafeGraph. Spend patterns [dataset], 2022.
- Tara Watson. Inside the refrigerator: Immigration enforcement and chilling effects in Medicaid participation. *American Economic Journal: Economic Policy*, 6(3):313–338, 2014.

A Single-worksite event inventory

The compiled single-worksite inventory contains 55 ICE-led events satisfying the event-date, geography, and geocoding screens described in Section 2.1. Table 9 lists the 51 events in the main analysis sample and the four single-worksite events that lack sufficient catchment POI coverage to fit the per-event distributed-lag DiD. Inclusion in the main sample requires at least 10 high-FB and 10 low-FB CBGs with sufficient POI coverage in the catchment. The four catchment-coverage exclusions are data-availability-driven rather than attention-correlated.

Table 9: Single-worksite ICE-led event inventory

Date	State	Event / worksite	Arrests
Main sample ($n = 51$)			
2018-04-05	TN	Southeastern Provision	97
2018-05-09	IA	MPC Enterprises (Midwest Precast Concrete)	32
2018-06-05	OH	Corso’s Flower & Garden Center (Sandusky + Castalia)	114
2018-06-19	OH	Fresh Mark	146
2018-08-28	TX	Load Trail LLC	159
2019-04-03	TX	CVE Technology Group	280
2019-08-07	MS	Mississippi poultry operation (7 plants, 5 companies)	680
2025-01-23	NJ	Newark	
2025-02-03	PA	Philadelphia	7
2025-02-12	TX	Los Fresnos	2
2025-02-18	NY	Tupper Lake Pine Mill	9
2025-02-26	MS	3 J Underground LLC	7
2025-02-26	MS	Gulf Coast Prestress Partners Ltd	18
2025-02-26	NJ	North Bergen	16
2025-02-27	PA	Jumbo Meat Market	4
2025-03-27	CA	Powder & Protective Coatings	5
2025-04-02	WA	Mt. Baker Roofing	37
2025-04-22	VT	Berkshire	8
2025-04-27	CO	Colorado Springs	100
2025-05-02	NY	Lynn-Ette & Sons Farms	14
2025-05-13	FL	Wildwood	33
2025-05-20	WA	Eagle Beverage	17
2025-05-21	AL		
2025-05-27	LA	Mirabeau Water Garden stormwater project	20
2025-05-29	FL	Perla	100

continued on next page

(Table 9 continued)

Date	State	Event / worksite	Arrests
2025-05-30	CA	Los Angeles	36
2025-05-30	CA	Buono Forchetta	3
2025-06-04	TX	Brownsville	25
2025-06-05	PA	Wyoming Valley Pallets Inc.	4
2025-06-06	CA	Los Angeles	44
2025-06-10	CA	Oxnard	35
2025-06-10	NE	Glenn Valley Foods	76
2025-06-11	PA	Five10 Flats restoration jobsite	17
2025-06-17	LA	Delta Downs Racetrack	84
2025-07-02	TX	Groomer's Seafood	12
2025-07-08	NJ	Alba Wine and Spirits	20
2025-07-10	CA	Glass House Farms	361
2025-07-10	TX	Taco Ole restaurant	18
2025-07-16	PA	Super Gigante	14
2025-08-20	NJ	Edison	29
2025-08-23	CT	Optimo Car Wash	7
2025-09-02	MO	Golden Apple Buffet	12
2025-09-04	GA	HL-GA Battery Company	475
2025-09-04	NY	Nutrition Bar Confectioners	57
2025-10-09	OH	Pancho's Tacos	5
2025-10-15	CT	Hamden	7
2025-10-19	ID	La Catedral Arena	105
2025-10-30	OR	Woodburn	35
2025-10-30	LA	Barrois Welding Services shipyard	25
2025-11-04	MA	Allston Car Wash	9
2025-11-14	AL	Nori Thai and Sushi	18
Excluded: insufficient catchment POI coverage ($n = 4$)			
2018-08-08	NE	O'Neill NE agricultural operation (3 sites)	
2025-06-04	NM	Outlook Dairy Farms	11
2026-01-15	MN	D.R. Horton	
2026-01-15	TX	El Paso	38

Table 10: Source-type audit for the single-worksite inventory

Primary cited source type	Inventory	Main	Med./low	Low spike
Prior compilation or tracker	23	22	11	14
State or local news	15	15	9	7
Agency or public record	10	9	8	4
National news	3	3	0	1
Business or address verification	2	2	2	1
Trade/compliance tracker	2	0	2	2
Total	55	51	32	29

Notes: The table classifies the primary citation recorded in the event inventory. “Inventory events” are the 55 single-worksite ICE-led events satisfying the screens in Section 2.1. “Main sample” counts the 51 events with sufficient catchment POI coverage to estimate the event-level distributed-lag DiD. “Low news spike” denotes events with an English-GDELT post-event spike at or below 0.01. The table is an audit of the public-source inventory rather than a claim that the inventory is an administrative universe of all obscure raids.

B Additional result diagnostics

B.1 News measures, inference, and local attention

The cross-language and cross-source validity of the headline measure can be quantified directly. The Mediacloud Spanish-ICE event spike closely tracks the English-GDELT ICE-raïd spike across the full event panel and closely tracks the US national Google Trends ICE-raïd spike where Trends data are available. The English-GDELT spike also tracks Spanish-language search spikes for “redada” and “deportación” in the 2024–2026 subsample. Event-level ICE news cycles propagate consistently through English news supply, Spanish news supply, and English and Spanish search demand on essentially the same schedule.


The choice of search term within the Mediacloud Spanish corpus matters. The headline uses the noun “ICE,” which has by far the cleanest event-specific signal. Alternative phrasings such as “agentes de ICE” and “redada migratoria” give the same negative sign but weaker statistical power. “Deportación” is essentially null, consistent with the term tracking the Spanish-language political-policy news cycle rather than event-specific raid coverage.

The generated-regressor concern does not drive the result. In inverse-variance-weighted WLS, where each event is weighted by the inverse of its squared standard error from the distributed-lag DiD, the Spanish-ICE coefficient is smaller than the OLS estimate but remains statistically significant. Adding arrests leaves the result intact. An event-level bootstrap on the OLS specification also excludes zero.

The post-event attention measure also does not appear to be standing in for anticipatory media attention. For each news-volume series we construct an immediate pre-event spike, defined as the peak over $\tau \in [-4, -1]$ minus the same $\tau \in [-8, -5]$ baseline used in the headline post-event spike. The immediate pre-event Spanish-ICE spike is uncorrelated with both the late-window estimate and the average pre-event foot-traffic coefficient. Adding it to the headline horse race leaves the post-event Spanish-ICE coefficient unchanged, and the English-GDELT measure gives the same pattern.

We also construct an event-local Spanish-news series by restricting the Mediacloud query to outlets tagged as published in the event’s home state. The local-Spanish spike moves in the same direction as the national headline, but its predictive power is weaker and it adds little once national Spanish-language attention is included. This mirrors the same-platform national-vs.-DMA contrast within Google Trends. The predictive content of local-language local-news attention lies primarily in the part that tracks the national signal.

Figure 2: Pooled event-time attention interaction



Notes. Estimates from a single pooled regression on the stacked CBG-week panel.

$\log(1 + y) = \alpha_{c,e} + \lambda_{w,e} + \sum_{\tau \neq -1} [\gamma_\tau + \delta_\tau \text{Spike}_e^z] \cdot \text{HighFB}_c^e \cdot \mathbf{1}[t - T_0^e = \tau] + \varepsilon$, where Spike_e^z is the standardized event-level national Google Trends spike. The figure plots δ_τ . Standard errors are clustered at the CBG level.

Whiskers are $\pm 1.96 \times \text{SE}$. Reference week is $\tau = -1$. $N = 204,510$ CBG-week observations. The pre-event coefficients are not perfectly flat, so the figure is treated as descriptive rather than as the main identifying test.

B.2 Clean-event robustness

We compute a contamination metric for each event. The metric is the number of other in-sample events whose dates fall within ± 8 weeks of event e and whose CBG catchments overlap with e 's. Of the 51 main-sample events, 34 have zero contaminating neighbors. Of those 34, 24 fall in the 2024–2026 subsample for which the pre-event state-Trends baseline is internally consistent (Section 2.5). On this subsample the saturation-test correlation between the late-window estimate and the pre-event state-Trends baseline strengthens slightly relative to the full 2024–2026 Trends sample, in the same direction and with the same statistical strength. The leave-one-out diagnostic for the headline Spanish-ICE horse race keeps the coefficient close to the full-sample point estimate when any single event is dropped. The corresponding English-GDELT leave-one-out diagnostic leads to the same conclusion.

C Instrument robustness and balance

Table 11: Competing-news IV robustness: foreign disasters and controls

	(1) Foreign	(2) + sev./era	(3) + event ctrls	(4) Both IVs
Attention coef.	−1.059*** (0.338)	−1.123*** (0.341)	−0.993** (0.475)	−1.052*** (0.340)
<i>t</i> -statistic	−3.13	−3.29	−2.09	−3.09
First-stage <i>F</i>	52.66	77.83	24.80	25.78
Sargan over-id (<i>p</i>)	—	—	—	0.830
Foreign-disaster IV	Yes	Yes	Yes	Yes
US-disaster IV	No	No	No	Yes
log(1 + arrests) + 2025+	No	Yes	Yes	No
Publicity + sector ctrls	No	No	Yes	No
<i>N</i>	51	49	49	51

Notes: The dependent variable is per-event β_{late}^e and the endogenous regressor is the Mediacloud Spanish-ICE attention spike. Columns (1)–(3) use the foreign-disaster instrument alone. Column (4) uses the foreign-disaster and US-disaster instruments jointly. Severity/era controls are log(1 + arrests) and an indicator for events in 2025 or later. Event controls add the event-publicity flag and broad sector indicators for agriculture/dairy, restaurant/food, manufacturing/industrial, car wash, and construction. First-stage statistics are robust Wald *F* tests for the excluded instruments. Heteroskedasticity-robust standard errors in parentheses. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 11 shows that the foreign-disaster IV result is not an artifact of obvious event-level controls. Adding log(1 + arrests) and a second-Trump-era indicator leaves the estimate close to the headline foreign-disaster estimate. Adding event-publicity and broad-sector controls attenuates the estimate but preserves the sign and statistical strength. The overidentified specification with foreign disasters and out-of-state U.S. disasters also leads to the same conclusion.

Table 12: Balance check: foreign disaster news pressure and event characteristics

Covariate	N	Coef. on foreign disasters (cond.)	p
Log arrests	49	+0.296	0.034
Pre-event beta	51	-0.002	0.624
Backward placebo beta	51	+0.004	0.674
Cross-state disaster days (std.)	51	+0.486	0.003
Joint test of listed covariates	49	$F = 4.35$	0.006

Notes: Each row reports the coefficient on standardized foreign-disaster days in a regression of the listed predetermined event characteristic on foreign-disaster days, second-Trump-era, event-publicity, and broad-sector controls. The foreign-disaster measure counts days in the ICE event’s post-event window that fall within the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. The joint test reverses the regression and asks whether the listed covariates jointly predict foreign-disaster days after the same era, publicity, and sector controls. Heteroskedasticity-robust inference is used. The table is a balance diagnostic, not a proof of exclusion.

Table 12 addresses the most direct timing concern. The question is whether foreign-disaster exposure is proxying for observable differences in which events occur when. The timing measure is related to the U.S.-disaster measure and to reported arrests, which is why we report the specification with both instruments and controls. It is not related to pre-event foot-traffic trends or the placebo response.

D Exposure-measure robustness

The headline high-FB indicator uses share of Latin-American foreign-born population (ACS B05006) at the CBG level. This is a broad measure because it counts naturalized US citizens alongside non-citizens, all Latin-American national-origin groups alongside the highest-enforcement-risk Mexican and Central American populations, and all adults alongside the US-citizen children of foreign-born parents. To test whether the chilling response is concentrated in narrower subgroups most directly exposed to deportation risk, we refit the per-event distributed-lag DiD with three sharper exposure indicators. These are share Hispanic non-citizen adults, share foreign-born from Mexico, El Salvador, Guatemala, Honduras, or Nicaragua, and share US-citizen children under 18 with at least one foreign-born parent.

Table 13: Exposure-measure robustness, per-event estimate on Spanish-ICE spike by high-FB indicator

Exposure indicator	N	mean estimate	sd estimate	r vs base	coef	t	R^2
Latin-Amer. FB (baseline, B05006)	51	−0.011	0.044	+1.000	−0.92	−3.59	0.227
Hispanic non-citizen (B05003I)	51	−0.007	0.046	+0.440	−0.17	−0.41	0.006
Mex. + Cent. Amer. FB (B05006)	51	−0.006	0.043	+0.716	−0.58	−1.98	0.089
Mixed-status kids (B05009)	49	−0.001	0.040	+0.551	−0.29	−1.05	0.027

Notes: For each exposure indicator, we refit the per-event distributed-lag DiD (Equation 1) defining high-FB CBGs as the top-decile and low-FB CBGs as the bottom-decile within each event’s catchment by the indicated variable. “Mean estimate” and “sd estimate” are the cross-event distribution of the resulting late-window estimates. “ r vs base” is the cross-event correlation of the refined-indicator estimates with the baseline estimates. The last three columns report the single-regressor OLS of the late-window estimate on the Mediacloud Spanish-ICE spike with HC1 standard errors.

All three refinements produce the same negative sign as the baseline at smaller mean magnitudes and substantially weaker cross-event predictive power for the Spanish-ICE attention spike. The pattern suggests that the cross-event response operates at the broad community level rather than narrowly at the personal-deportation-risk level. This weakening is not mechanically driven by fewer POIs in the sharper top-decile cells. Top-decile POI density under the sharper indicators is within ± 6 percent of the baseline.

E Venue-coverage and civic-mechanism checks

Advan’s POI panel does not measurably capture two venue categories that are central to the sensitive-locations policy and to the broader immigration-fear literature. Religious organizations (NAICS-3 = 813) contain 4,025 POIs across the 26-state catchment universe, of which 3,817 (95 percent) are the Independent Community Bankers of America contamination documented in Appendix F. After stripping obvious mislabels, only a very small number of panel POIs correspond to real religious congregations. Healthcare venues are absent. NAICS-2 = 62 has zero POIs in the panel. The civic-versus-commercial decomposition therefore captures the schools and parks/museums response but is silent on churches and hospitals.

We also test whether the civic-vs.-commercial differential is itself moderated by the cross-event attention spike. With the English-GDELT spike, the within-event civic-minus-commercial differential is -0.153 in the low-spike group and -0.051 in the high-spike group. With the Mediacloud Spanish-ICE spike, the pattern is similar. The differential is -0.162 in the low-spike group and -0.060 in the high-spike group. The civic-vs.-commercial gap is wider in low-attention events than in high-attention events under both measures, the opposite of what an attention-moderated mechanism would predict.

F POI data-quality audits

F.1 Advan and SafeGraph POI classification

This section discusses non-trivial point-of-interest classification errors in the underlying SafeGraph POI universe. These include NAICS misclassification across sectors, POIs identified as active businesses that are in fact vacant or demolished structures, and visit-count series that are mathematically inconsistent with the POI’s physical footprint.

F.2 Geometric POI flags

We construct two independent quality flags for all Advan POIs in the catchments of every event in our sample, covering approximately 1.54 million distinct POIs across the event catchments used in the audit. Building-containment is tested against Microsoft Building Footprints, and freeway-proximity is tested against TIGER Primary and Secondary Roads. A POI receives the freeway-traffic flag if its lat/lon falls outside any building footprint and is within 50 meters of an interstate centerline. A POI receives the remote flag if it falls outside any building, more than 100 meters from the nearest building, and more than 200 meters from any S1100 or S1200 road.

Across the 26-state panel, 7,881 POIs receive the freeway-traffic flag (0.5 percent) and 20,324 receive the remote flag (1.3 percent), for a combined drop rate of about 1.8 percent. The flagged set catches the expected error mode. In California, the five most-visited flagged POIs are the Port of Los Angeles, the Port of Long Beach, an Irwindale aggregates quarry, the Long Beach port-adjacent industrial complex, and the Port of Los Angeles San Pedro terminal. The flagged POIs are disproportionately in low-Latino-foreign-born CBGs rather than high-Latino-foreign-born CBGs, so removing them shrinks the comparison group rather than the treatment group. Re-estimating the per-event late-window effect on the cleaned panel shifts the cross-event mean by 0.0004 log-points. Three of 51 events flip sign, all with absolute estimates below 0.01 in either direction. The mean absolute change per event is 0.0056 log-points. The information-shock finding does not depend on the small subset of POIs that fail the building-containment and freeway-proximity checks.

F.3 POI verification and sector recoding

We supplement the geometric check with hand-verification of 100 randomly sampled California POIs, stratified by Veridion firm-registry match quality. Of the 100, 96 correspond to a real business or commercial property at the listed address. The four exceptions are all residential single-family homes that Advan labels with a generic NAICS-category string, and all four fall in the no-Veridion-match tier where Advan errors should be concentrated. No POI in the sample was a demolished structure or a business that never existed.

A second verification of 500 California POIs stratified by NAICS-3 was designed to test whether classification noise could reshuffle the civic-versus-commercial buckets. The Advan-bucket $\times$ true-bucket cross-tabulation shows that the directional contamination that would erase the mechanism

finding is essentially zero. Only 1 of 200 civic-labeled POIs is actually commercial. The audit did identify a consequential bug in NAICS-3 = 813. In California, 110 of the 118 Advan POIs in NAICS-3 = 813 are labeled “Independent Community Bankers of America” but correspond to individual community-bank branch addresses. A further three NAICS-813 records are sports stadiums. We therefore drop NAICS-3 = 813 from the civic bin in Table 8. Including 813 attenuates the civic estimate toward zero by mixing in community-bank visitor activity whose own per-event response is closer to the commercial-bucket value.

F.4 Structural name-level checks

To move from spot-checks to panel-wide quantification, we conduct a verification exercise at the name-recurrence level. The 1.15M-POI 21-state catchment universe contains only 129,773 unique `location_name` values. The top 1,500 names cover 978,660 POIs. We classify each of the top names as a national chain with consistent NAICS, a generic NAICS-category-as-name label, a public-sector entity, a brand-at-host pattern in which a service brand sits inside a host retailer, or a trade-association-at-members pattern in which an association name is spread across member-business addresses.

The verification identifies 62 names with bulk-recodable structural miscoding, of which 44 require an actual NAICS-3 change. The 44 effective recoding rules cover 63,924 POIs (5.5 percent of the 21-state universe by POI count and 5.3 percent of the panel by visitor-week count). The largest pattern is money-transfer-agent POIs that live inside grocery and convenience-store hosts. For the civic-versus-commercial decomposition, however, the implication is limited. Only 148 POIs would leave the civic bucket under the recoding rules, and the audit identifies no name that would move POIs from civic to commercial or commercial to civic. Re-estimating the category decomposition after applying the full set of structural recoding rules moves the high-baseline civic estimate from -0.135 to -0.120 , while the commercial estimate moves from -0.001 to $+0.001$. The order-of-magnitude civic-versus-commercial gap is preserved.

F.5 Panel coverage

A separate concern is that businesses may be real but missing from the Advan POI universe altogether. If Advan systematically under-covers immigrant-serving businesses, the chilling regression would be attenuated downward in heavily Latino CBGs because the relevant POIs are not in the panel. We test this using Veridion as an external firm registry. For each 5-digit ZIP in the panel, we count Advan POIs and Veridion firms and compute the coverage ratio. Restricting to dense urban ZIPs and stratifying by foreign-born share quartile, the firm-weighted Advan/Veridion ratio is 0.28 in the lowest-FB quartile and 0.33 in the highest-FB quartile. The mild positive gradient does not support systematic undercoverage of immigrant neighborhoods.

Two Census establishment references corroborate the Veridion-based finding. At the county $\times$ NAICS-3 level, comparing Advan POIs to Census County Business Patterns, the establishment-weighted Advan/CBP coverage ratio rises monotonically from the lowest-FB quartile to the highest-

FB quartile. At the ZIP level, comparing against Census ZIP Business Patterns, the Advan/ZBP coverage ratio is essentially flat across foreign-born-share quartiles. Three independent external references therefore point the same direction. The panel-composition channel is a real concern in principle, but it is not the empirical source of the result.

F.6 Residential labels and visit-pattern mismatches

A further exercise targets residential single-family homes that Advan labels with a generic NAICS-category string. We construct a residence-mislabel flag combining a generic NAICS-category-as-name label with a strict-residential street suffix. The flag identifies 28,204 POIs (2.4 percent of the 21-state panel). A stronger version that also requires no Veridion firm within 100 meters identifies 17,456 POIs (1.5 percent). Spot-checks confirm high precision. Only 157 of the flagged residences fall in the civic bucket, and dropping them would shrink the civic bucket by less than 1 percent. The remainder of the flag concentrates in the commercial and other buckets and does not shift the commercial estimate off its near-zero baseline.

A final check addresses visit-count series that are inconsistent with what the POI claims to be, using only the visit data itself. For each POI we aggregate the Advan day-of-week and hour-of-day visit arrays over all of 2025, normalize each to a shape profile, and build a reference signature for every NAICS-3 with at least 30 well-measured POIs. A POI is flagged `flag_visit_pattern_mismatch` if both its day-of-week and hour-of-day profile correlate below 0.30 with its own sector’s reference signature. The flag identifies 75,963 POIs (6.6 percent of the panel). It is the noisiest quality check. The flagged set mixes genuine anomalies with legitimate businesses that keep atypical hours. We report it as a diagnostic rather than fold it into the headline cleaning, because at this rate and with non-trivial false positives it is too blunt to use as a panel filter. Its value is that it catches a class of error that no external firm registry can detect, a real address with a real NAICS label but an implausible visit series.