

Online Appendix to

JUE Insight: Attention and the Economic Response to Immigration Enforcement

Greg C. Wright
University of California, Merced

September 1, 2026

This Online Appendix collects the additional tables, figures, and diagnostics referenced in the main paper. It is provided for review and is not part of the typeset article. Sections and exhibits are numbered with an “S” prefix. Acronyms and variable definitions follow the main paper. References to “the main paper” point to *JUE Insight: Attention and the Economic Response to Immigration Enforcement*.

S1 Event inventory and summary statistics

Table S1 reports event-level summary statistics by enforcement wave. The compiled single-worksite inventory contains 55 ICE-led events satisfying the event-date, geography, and geocoding screens described in the main paper. Table S3 lists the 51 events in the main analysis sample and the four single-worksite events that lack sufficient catchment POI coverage to fit the per-event distributed-lag DiD. Inclusion in the main sample requires at least 10 high-FB and 10 low-FB CBGs with sufficient POI coverage in the catchment. The four catchment-coverage exclusions are data-availability-driven rather than attention-correlated.

Table S1: Event-level summary statistics, by enforcement wave

	Full	2025–2026	2018–2019
Arrests reported at worksite	70 (128)	46 (89)	215 (218)
Spanish-ICE news spike	0.0298 (0.0242)	0.0337 (0.0235)	0.0054 (0.0100)
English-GDELT news spike	0.0146 (0.0145)	0.0163 (0.0148)	0.0036 (0.0027)
Late-window estimate (log pts)	-0.008 (0.047)	-0.009 (0.047)	-0.003 (0.046)
High-FB CBGs in catchment	118 (167)	130 (175)	43 (69)
Low-FB CBGs in catchment	134 (165)	143 (172)	84 (100)
Events	51	44	7

Notes: Cells report the cross-event mean with the standard deviation in parentheses, for the full 51-event sample and separately by enforcement wave. News spikes are the post-event peak above the pre-event baseline in the share of stories mentioning the term (see the main paper). The late-window estimate is the per-event weeks-4–8 differential in log foot traffic. High-FB and low-FB CBGs are the top- and bottom-decile Latin-American-foreign-born block groups in each event’s catchment. Arrests are missing for two events.

Table S2: Sample-flow from full event inventory to estimation subsamples

Definition	Events
Single-worksite ICE-led events satisfying inventory screens	55
– Single-worksite ICE-led but no estimable late-window estimate (insufficient catchment POI coverage) ^a	–4
Main analysis sample (single-worksite ICE-led with estimable late-window estimate)	51
<i>Estimation subsamples:</i>	
Events with reported arrest counts (severity test)	49
Events in 2025–2026 event wave (Google Trends national)	44
Events with DMA-level Google Trends coverage	41
Events with at least one in-state Mediacloud Spanish-language source	31
Events with estimable per-event civic estimate	12

Notes: “Events” is the unit. The 51-event panel is the main analysis sample for the stacked foot-traffic design and the event-level robustness exercises. Smaller numbers below reflect availability of specific covariates rather than re-estimation. The Mediacloud Spanish-language national series and the English-GDELT series are defined for all 51 events. Trends, DMA Trends, and in-state Spanish sources cover smaller subsamples. The civic estimate requires sufficient civic-NAICS POI density in the high-FB CBGs of an event’s catchment.

Table S3: Single-worksite ICE-led event inventory

Date	State	Event / worksite	Arrests
Main sample ($n = 51$)			
2018-04-05	TN	Southeastern Provision	97
2018-05-09	IA	MPC Enterprises (Midwest Precast Concrete)	32
2018-06-05	OH	Corso’s Flower & Garden Center (Sandusky + Castalia)	114
2018-06-19	OH	Fresh Mark	146
2018-08-28	TX	Load Trail LLC	159
2019-04-03	TX	CVE Technology Group	280
2019-08-07	MS	Mississippi poultry operation (7 plants, 5 companies)	680
2025-01-23	NJ	Newark	
2025-02-03	PA	Philadelphia	7
2025-02-12	TX	Los Fresnos	2
2025-02-18	NY	Tupper Lake Pine Mill	9
2025-02-26	MS	3 J Underground LLC	7
2025-02-26	MS	Gulf Coast Prestress Partners Ltd	18
2025-02-26	NJ	North Bergen	16
2025-02-27	PA	Jumbo Meat Market	4
2025-03-27	CA	Powder & Protective Coatings	5
2025-04-02	WA	Mt. Baker Roofing	37
2025-04-22	VT	Berkshire	8
2025-04-27	CO	Colorado Springs	100
2025-05-02	NY	Lynn-Ette & Sons Farms	14
2025-05-13	FL	Wildwood	33
2025-05-20	WA	Eagle Beverage	17
2025-05-21	AL	Baldwin County construction site	

continued on next page

(Table S3 continued)

Date	State	Event / worksite	Arrests
2025-05-27	LA	Mirabeau Water Garden stormwater project	20
2025-05-29	FL	Perla	100
2025-05-30	CA	Los Angeles	36
2025-05-30	CA	Buono Forchetta	3
2025-06-04	TX	Brownsville	25
2025-06-05	PA	Wyoming Valley Pallets Inc.	4
2025-06-06	CA	Los Angeles	44
2025-06-10	CA	Oxnard	35
2025-06-10	NE	Glenn Valley Foods	76
2025-06-11	PA	Five10 Flats restoration jobsite	17
2025-06-17	LA	Delta Downs Racetrack	84
2025-07-02	TX	Groomer's Seafood	12
2025-07-08	NJ	Alba Wine and Spirits	20
2025-07-10	CA	Glass House Farms	361
2025-07-10	TX	Taco Ole restaurant	18
2025-07-16	PA	Super Gigante	14
2025-08-20	NJ	Edison	29
2025-08-23	CT	Optimo Car Wash	7
2025-09-02	MO	Golden Apple Buffet	12
2025-09-04	GA	HL-GA Battery Company	475
2025-09-04	NY	Nutrition Bar Confectioners	57
2025-10-09	OH	Pancho's Tacos	5
2025-10-15	CT	Hamden	7
2025-10-19	ID	La Catedral Arena	105
2025-10-30	OR	Woodburn	35
2025-10-30	LA	Barrois Welding Services shipyard	25
2025-11-04	MA	Allston Car Wash	9
2025-11-14	AL	Nori Thai and Sushi	18
<i>Excluded: insufficient catchment POI coverage ($n = 4$)</i>			
2018-08-08	NE	O'Neill NE agricultural operation (3 sites)	
2025-06-04	NM	Outlook Dairy Farms	11
2026-01-15	MN	D.R. Horton	
2026-01-15	TX	El Paso	38

S2 Additional result diagnostics

S2.1 Event-level robustness specification

The main paper estimates the pooled stacked design. As a robustness exercise, I also estimate the event-time design separately for each event. For event e , the event-level specification is

$$\log(1 + y_{c,t}^e) = \alpha_c^e + \lambda_t^e + \sum_{\tau \neq -1} \gamma_\tau^e (\text{HighFB}_c^e \times \mathbf{1}[t - T_0^e = \tau]) + \varepsilon_{c,t}^e. \quad (\text{S1})$$

The omitted category is $\tau = -1$. The late-window estimate β_{late}^e is the average of γ_τ^e over $\tau \in [4, 8]$, with its standard error computed from the CBG-clustered covariance matrix. I then relate those event-level estimates to news coverage and arrests:

$$\beta_{\text{late}}^e = \pi_0 + \pi_S \text{Spike}_e + \pi_A \log(1 + \text{Arrests}_e) + u_e. \quad (\text{S2})$$

Table S4 reports the event-level version of the main estimate, with precision-weighted, OLS, and IV estimates on the same per-standard-deviation scale.

Table S4: The effect of news coverage on neighborhood foot traffic

	Precision-weighted	OLS		2SLS (news-shock IV)	
	(1)	(2)	(3)	(4)	(5)
Spanish-ICE spike (per SD)	−0.0184*** (0.0051)	−0.0222*** (0.0062)	−0.0241*** (0.0064)	−0.0256*** (0.0082)	−0.0267*** (0.0081)
$\log(1 + \text{Arrests})$			−0.0053 (0.0042)		−0.0055 (0.0039)
First-stage F				52.7	51.2
True variance explained	30%				
Events	51	51	49	51	49

Notes: The dependent variable is the per-event late-window (weeks 4–8) foot-traffic estimate from Equation (S1), in log points. The news coverage spike is the post-event peak in Mediacloud US Spanish-language ICE coverage, standardized to the 51-event sample. Column (1) weights each event by the inverse sampling variance of its estimate, with Knapp–Hartung standard errors. Columns (2)–(3) are OLS with HC1 standard errors. Columns (4)–(5) instrument the spike with foreign-disaster news pressure. The arrests specifications have $N = 49$ because two events lack reported arrest counts. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

The log transformation of the outcome does not drive the result. Table S5 re-estimates the separate-event robustness design by Poisson pseudo-maximum likelihood on visit counts. The news-coverage coefficient remains negative and similar in magnitude.

Table S5: PPML robustness of the effect of news coverage


First-step outcome model	Precision-weighted	OLS	2SLS
log(1 + visits) (baseline)	−0.0184*** (0.0051)	−0.0222*** (0.0062)	−0.0256*** (0.0082)
Poisson (PPML), visit counts	−0.0156*** (0.0055)	−0.0157** (0.0064)	−0.0158 (0.0104)
Events	51		

Notes: Each cell reports the coefficient on the standardized Spanish-ICE news-coverage spike from the event-level robustness exercise, using per-event late-window estimates from the indicated outcome model. The PPML rows re-estimate every event by Poisson pseudo-maximum likelihood on raw visit counts with the same CBG and calendar-week fixed effects and CBG-clustered standard errors; PPML coefficients are semi-elasticities, comparable to the log-outcome estimates. The per-event PPML and log estimates correlate at $r = 0.71$ across events. No CBG-week observation in the estimation panels has zero visits, so the two outcome models share an identical estimation sample. Precision-weighted columns use Knapp–Hartung standard errors, OLS columns HC1, and 2SLS columns the foreign-disaster news-shock instrument with robust standard errors. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

S2.2 Average response in event time

Figure S1 reports the average high-FB versus low-FB foot-traffic path that corresponds to the coverage interaction shown in Figure 1 of the main paper. The average post-event response is small in both samples. The full-sample path drifts downward before the event, while the clean-subsample path is flat. This contrast shows why the paper focuses on how the response scales with coverage rather than on a uniform average effect.

Figure S1: Average foot-traffic response in event time



Notes. Pooled distributed-lag DiD with event $\times$ CBG and event $\times$ calendar-week fixed effects. The figure plots the high-FB versus low-FB foot-traffic gap for an event with average Spanish-language news coverage. The clean subsample has no overlapping-catchment contamination and flat pre-trends. Shaded bands are 95 percent confidence intervals based on standard errors clustered by event.

S2.3 News measures, inference, and local attention

The cross-language and cross-source validity of the headline measure can be quantified directly. The Mediacloud Spanish-ICE event spike closely tracks both the English-GDELT ICE-raid spike across the full event panel and the US national Google Trends ICE-raid spike where Trends data are available. The English-GDELT spike also tracks Spanish-language search spikes for “redada” and “deportación” in the 2025–2026 event wave. Event-level ICE news cycles propagate consistently through English news supply, Spanish news supply, and English and Spanish search demand on essentially the same schedule.

The choice of search term within the Mediacloud Spanish corpus matters. The headline uses the noun “ICE,” which has by far the cleanest event-specific signal. Alternative phrasings such as “agentes de ICE” and “redada migratoria” give the same negative sign but weaker statistical power. “Deportación” is essentially null, consistent with the term tracking the Spanish-language political-policy news cycle rather than event-specific raid coverage.

The predictive content is concentrated in sustained rather than transient coverage. Table S6 regresses the per-event late-window estimate on each of several coverage-timing measures in turn. The weekly peak, an eight-week area-under-the-curve measure, and national search each predict the response, while a single-day peak and counts of days on which coverage crosses a threshold are weaker, both in statistical significance and in the share of cross-event variance explained.

Table S6: Timing of news coverage and the foot-traffic response

Coverage-timing measure	N	Coef. per SD	R^2
<i>Panel A: Sustained / aggregate attention</i>			
National Google Trends spike	44	−0.022*** (0.007)	0.22
Spanish-ICE weekly peak	51	−0.022*** (0.006)	0.23
English-GDELT weekly peak	51	−0.019*** (0.006)	0.17
English-GDELT eight-week AUC	51	−0.019*** (0.006)	0.17
<i>Panel B: Transient / threshold measures</i>			
English-GDELT one-day peak	51	−0.009 (0.007)	0.03
Days above baseline + 1 SD	51	−0.011* (0.006)	0.06
Days above 2× baseline	51	−0.009 (0.006)	0.04

Notes: Each row is a separate cross-event OLS regression of the per-event late-window foot-traffic estimate on one standardized coverage-timing measure, so the coefficient is per standard deviation of that measure and a more negative value means a larger predicted foot-traffic decline. R^2 is the share of cross-event variance the single measure explains. Measures are built from the English-language GDELT ICE news-volume series unless noted: the weekly peak, the post-event eight-week area under the coverage curve (AUC), the single-day peak, and counts of days on which coverage exceeds its pre-event baseline by more than one standard deviation or twice the baseline. The national Google Trends spike is available for 44 events in the 2025–2026 wave; the other measures use the full 51-event sample. The English-GDELT weekly peak is the same series as the headline English-GDELT news spike used elsewhere in the paper, labeled here by its timing operationalization; it is shown once. Sustained and aggregate measures (Panel A) predict the response, while single-day and threshold-crossing measures (Panel B) are weaker. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

The generated-regressor concern does not drive the result. In inverse-variance-weighted WLS, where each event is weighted by the inverse of its squared standard error from the distributed-lag DiD, the Spanish-ICE coefficient is smaller than the OLS estimate but remains statistically significant. Adding arrests leaves the result intact. An event-level bootstrap on the OLS specification also excludes zero.

The post-event attention measure also does not appear to be standing in for anticipatory media attention. For each news-volume series an immediate pre-event spike is constructed, defined as the peak over $\tau \in [-4, -1]$ minus the same $\tau \in [-8, -5]$ baseline used in the headline post-event spike. The immediate pre-event Spanish-ICE spike is uncorrelated with both the late-window estimate and the average pre-event foot-traffic coefficient. Adding it to the headline horse race leaves the post-event Spanish-ICE coefficient unchanged, and the English-GDELT measure gives the same pattern.

An event-local Spanish-news series is also constructed by restricting the Mediacloud query to

outlets tagged as published in the event’s home state. The local-Spanish spike moves in the same direction as the national headline, but its predictive power is weaker and it adds little once national Spanish-language attention is included. This mirrors the same-platform national-vs.-DMA contrast within Google Trends. The predictive content of local-language local-news attention lies primarily in the part that tracks the national signal.

Table S7 reports the national-versus-DMA Google Trends comparison for the 2025–2026 event wave.

Table S7: Local vs. national coverage of enforcement events

	(1) Nat.	(2) + arr.	(3) + DMA	(4) DMA	(5) DMA + arr.
National Trends spike $\text{Spike}_e^{\text{Nat}}$	−0.000525*** (0.000158)	−0.000564*** (0.000167)	−0.000621*** (0.000228)		
DMA Trends spike $\text{Spike}_e^{\text{DMA}}$			+0.000074 (0.000215)	−0.000326* (0.000174)	−0.000330* (0.000184)
$\log(1 + \text{Arrests}_e)$		−0.001275 (0.005263)			+0.000589 (0.006522)
Constant	+0.010989 (0.008648)	+0.015012 (0.017686)	+0.012604 (0.008953)	+0.006794 (0.008799)	+0.005279 (0.021078)
N	44	42	41	41	39
R^2	0.215	0.232	0.239	0.085	0.079

Notes: OLS regressions of the per-event late-window estimate on national and DMA-level Google Trends spikes for “ICE raid.” Both spikes are peak search interest over $\tau \in [0, 8]$ minus the pre-event baseline. The sample is the 2025–2026 event wave with available Trends data. HC1 standard errors in parentheses. Stars:

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table S8 shows that the negative news-coverage relationship builds over the first eight post-event weeks.

Table S8: Early-window robustness

	Weeks 0–2	Weeks 1–3	Weeks 4–8
<i>Panel A. Event-estimate distribution</i>			
Mean estimate	−0.001	−0.004	−0.008
Median estimate	0.000	−0.003	−0.003
Share negative	0.51	0.55	0.57
<i>Panel B. Cross-event effect of news coverage</i>			
Pearson r , Spanish-ICE spike	−0.263	−0.375	−0.477
OLS coef., Spanish-ICE spike	−0.323* (0.174)	−0.524*** (0.191)	−0.919*** (0.256)
OLS coef., Spanish-ICE spike, with arrests	−0.388** (0.181)	−0.592*** (0.202)	−0.999*** (0.265)
<i>Panel C. Corroborating English-news measure</i>			
Pearson r , English-GDELT spike	−0.227	−0.333	−0.416
OLS coef., English-GDELT spike	−0.466 (0.302)	−0.774** (0.333)	−1.338*** (0.438)
N	51	51	51

Notes. Each column summarizes the same per-event distributed-lag DiD path over a different post-event window. Week 0 is the calendar week containing the enforcement event and can include pre-event days. The weeks 1–3 window excludes the event week. The weeks 4–8 column is the headline persistent-response window used in the main horse race. OLS rows report single-regressor HC1 standard errors in parentheses, except for the “with arrests” row, which also includes $\log(1 + \text{Arrests}_e)$ as a control. The arrests-control regressions use 49 events. Stars denote * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

S2.4 Geographic and sectoral mechanisms

Table S9 reports the distance- and sector-split estimates summarized in the main paper. Each row applies the event-level robustness design to the indicated POI subsample and relates the resulting late-window estimates to Spanish-language news coverage.

Table S9: Mechanisms

Sample	Events	Mean estimate	Coef. on Spanish-ICE spike
<i>Panel A. Geographic reach, mutually exclusive distance rings</i>			
0–5 km from worksite	12	+0.031	–0.022 (0.017)
5–15 km from worksite	29	–0.013	–0.026*** (0.007)
15–30 km from worksite	36	+0.002	–0.023* (0.012)
30–50 km from worksite	40	–0.009	–0.009 (0.006)
<i>Panel B. Same-sector and broad-sector activity</i>			
Same broad sector as raided worksite	20	+0.003	+0.012 (0.017)
Other sectors	49	–0.011	–0.022*** (0.007)
Consumer-facing POIs	51	–0.007	–0.018*** (0.006)
Nonconsumer POIs	48	–0.003	–0.029*** (0.008)

Notes. Each row regresses a subgroup-specific late-window estimate on the standardized Mediacloud Spanish-ICE news coverage spike. Panel A splits POIs by distance from the worksite. Panel B splits POIs by sector. HC1 standard errors in parentheses. Stars denote * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

S2.5 Clean-event robustness

A contamination metric counts other in-sample events within ± 8 weeks whose CBG catchments overlap with the focal event. Table S10 shows that news coverage is unrelated to pre-trend flatness, the number of CBGs in the catchment, estimate precision, or overlapping events. Adding these design measures as controls leaves the news-coverage coefficient unchanged. The effect also remains negative as the sample is restricted first to events without overlapping catchments and then to events that additionally have statistically flat pre-trends, although precision falls in the smallest subsample.

Table S10: Design validity and news coverage

<i>Panel A: Design metric regressed on the standardized Spanish-ICE spike (HC1)</i>					
Design metric	<i>N</i>	slope/SD	<i>t</i>	<i>p</i>	<i>R</i> ²
Pre-trend joint <i>F</i> (non-flatness)	51	-217.6642	-1.04	0.301	0.047
pre-event coef	51	-0.0056	-1.31	0.191	0.045
placebo response	51	-0.0060	-1.25	0.210	0.035
# CBGs in catchment	51	-9.3960	-0.24	0.810	0.001
Per-event SE (imprecision)	51	-0.0005	-0.11	0.915	0.000
# contaminating events	51	0.1667	1.04	0.296	0.025
Frac. catchment shared	51	0.0256	0.51	0.611	0.004
<i>Panel B: News-coverage coefficient with and without design controls</i>					
	<i>N</i>	coef/SD	SE	<i>t</i>	<i>p</i>
Spanish-ICE spike (per SD), no controls	51	-0.0184	0.0051	-3.65	0.001
+ design controls	51	-0.0185	0.0054	-3.42	0.001
<i>Panel C: Heterogeneity on clean subsamples</i>					
Subsample	<i>N</i>	<i>I</i> ²	τ	spike <i>p</i>	True var. expl.
All events	51	69.7%	0.0287	0.001	30.1%
No contamination	34	61.1%	0.0309	0.027	14.8%
No contam. & flat pre-trends	23	62.8%	0.0301	0.051	16.6%

Notes: Panel A regresses each per-event design-quality metric on the standardized Spanish-ICE news spike with HC1 standard errors. A null slope means design quality does not vary systematically with news coverage. Panel B shows that the precision-weighted news-coverage coefficient is unchanged when standardized design metrics (pre-trend *F*, # CBGs, # contaminating events) are added as controls. Panel C recomputes the heterogeneity statistics and the precision-weighted news-coverage regression on progressively cleaner subsamples (no overlapping-catchment contamination; additionally non-rejected flat pre-trends at $p > 0.10$). τ is in log points.

Table S11 reports the effect of news coverage estimated separately by enforcement wave. Within the 2025–2026 wave alone, which contains 44 of the 51 events, the effect is at least as strong as in the full sample, so the cross-event relationship is not an artifact of the structural differences between the two enforcement waves. The 2018–2019 wave has only seven events, too few for precision-weighted or IV inference, but its descriptive OLS slope has the same sign.

Table S11: Effect of news coverage by enforcement wave

Sample	<i>N</i>	Precision-weighted	OLS	IV
Full sample	51	−0.0184*** (0.0051)	−0.0222*** (0.0062)	−0.0256*** (0.0082)
2025–2026 wave	44	−0.0199*** (0.0053)	−0.0247*** (0.0069)	−0.0287*** (0.0076)
2018–2019 wave	7	—	−0.0149 (0.0220)	—
First-stage <i>F</i>				52.7 / 119.1

Notes: The dependent variable is the per-event late-window foot-traffic estimate. All coefficients are per standard deviation of the Mediacloud Spanish-ICE news-coverage spike, standardized within sample. The precision-weighted column weights each event by the inverse sampling variance of its estimate and uses Knapp–Hartung inference. The OLS and IV columns use HC1 and robust standard errors, respectively. The IV column instruments news coverage with foreign-disaster news pressure. The 2018–2019 wave has only seven events, so precision-weighted and IV estimates are not reported for that subsample. Its OLS estimate is descriptive. The first-stage *F* row reports the full-sample and 2025–2026 statistics. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

The effect is also unlikely to be an artifact of mismeasured operational scale. Table S12 shows that the news-coverage coefficient is essentially unchanged when arrests, operation type, and sector are added as controls. News coverage dominates arrests in a per-standard-deviation horse race, and the small arrest coefficient remains small relative to the coverage coefficient even under generous corrections for measurement error in arrests. Arrests are the only continuous operational-scale measure the shock inventory records, so it is controlled for directly rather than by building a scale index from a single variable. Operation type and sector separately capture composition.

Table S12: Operational scale and the effect of news coverage

	Coverage coef.	(KH SE)	p	True var. expl.	N
<i>Panel A: News-coverage coefficient as scale controls are added</i>					
Spanish-ICE spike	-0.0184	(0.0051)	0.0006	30.1%	51
+ log arrests	-0.0192	(0.0053)	0.0007	25.1%	49
+ arrests, type, sector	-0.0190	(0.0074)	0.0141	15.1%	49
English-GDELT spike	-0.0154	(0.0050)	0.0036	18.5%	51
+ log arrests	-0.0161	(0.0052)	0.0035	16.9%	49
<i>Panel B: Per-SD horse race, news coverage vs. log arrests</i>					
News coverage spike (per SD)	-0.0192	(0.0053)	0.0007	25.1%	49
Log arrests (per SD)	-0.0040	(0.0055)	0.4675	—	—
<i>Panel C: De-attenuated arrests coefficients</i>					
Reliability ρ	Arrests coef.	Ratio			
1.0 (no error)	-0.0030	0.16			
0.7	-0.0042	0.23			
0.5	-0.0059	0.32			
0.4	-0.0074	0.40			

Notes: Precision-weighted regressions of the per-event late-window estimate on standardized regressors, with Knapp–Hartung adjusted inference and inverse-variance weights from the clustered linear-combination standard errors. Arrests are the only continuous operational-scale measure recorded in the shock inventory: fines are unrecorded for every event, and operation end-dates are recorded for only 1 of 51 events (and reflect case resolution rather than operation length), so arrests are controlled for directly rather than building a scale index from one variable. Operation type/composition controls are worksite operation, sector agriculture, sector food, sector manufacturing, sector construction. Publicity is a visibility measure and is deliberately excluded from the scale block. Panel C reports the implied true single-regressor arrests coefficient b_{arr}/ρ under classical measurement error with reliability ρ ; even at $\rho = 0.5$ the implied coefficient is small relative to the news-coverage coefficient. Coverage coefficients are per standard deviation of the spike; the arrests specifications have $N = 49$.

Table S13 reports the same issue using alternative functional forms for reported arrests, including raw arrests and arrests scaled to the local catchment population. These measures explain little when entered alone, and adding them leaves the news-coverage coefficient stable. The two per-capita measures become negative when conditioned on coverage, so I treat them as secondary evidence rather than as proof that operational scale is irrelevant.

Table S13: Functional forms for reported arrests

Arrest proxy	Proxy only	R^2	Coverage coef	Proxy with coverage	N
Raw arrests	+0.0005 (0.0031)	0.000	-0.0245*** (0.0067)	-0.0052 (0.0038)	49
Log arrests	-0.0037 (0.0063)	0.006	-0.0241*** (0.0064)	-0.0068 (0.0054)	49
Arrests per 100k catchment pop.	-0.0008 (0.0028)	0.000	-0.0257*** (0.0067)	-0.0084** (0.0039)	49
Arrests per 100k high-FB pop.	-0.0011 (0.0029)	0.001	-0.0255*** (0.0066)	-0.0082* (0.0042)	49

Notes: Each row reports event-level OLS regressions where the dependent variable is the late-window foot-traffic estimate. All continuous regressors are standardized within the estimation sample, so coefficients are per standard deviation. “Proxy only” is the coefficient from a single-regressor specification. “Coverage coef” and “Proxy with coverage” come from a two-regressor specification that also includes the Mediacloud Spanish-ICE news-coverage spike. Catchment population is computed from unique CBGs in the 50 km Advan event catchment. High-FB CBG population uses CBGs in the top decile of Latin-American foreign-born share within that catchment. Heteroskedasticity-robust standard errors are in parentheses. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

The effect also does not reflect how events entered the inventory. Table S14 classifies each event by the type of its source citation. Events documented through news or official sources carry higher news coverage spikes than events documented through other channels. Discovery source does not predict the per-event estimate, and the effect of news coverage is, if anything, stronger under discovery-source fixed effects.

Table S14: Discovery-source robustness

<i>Panel A: News coverage spike and estimate by source type</i>			
Source type	Events	Mean spike	Mean β_{late}
news/web	32	0.0340	-0.0094
official	9	0.0349	+0.0050
other	10	0.0116	-0.0156
<i>Source $\rightarrow$ spike, joint F</i>		5.25	$p = 0.01$
<i>Source $\rightarrow \beta$, joint F</i>		1.63	$p = 0.21$
<i>Panel B: Effect of news coverage with source fixed effects</i>			
	Slope	(KH SE)	p
Baseline	-0.0184	(0.0051)	0.0006
+ discovery-source FE	-0.0228	(0.0051)	0.0001

Notes: $N = 51$. Each event is classified by the type of its inventory source citation (official government, news/web, advocacy/social, or other; categories smaller than 5 events are pooled into “other”). Panel A reports the mean news coverage spike and mean late-window estimate by source type and the joint F from regressing the standardized spike and the estimate on source dummies. Panel B is the precision-weighted news-coverage regression with and without discovery-source fixed effects. The source citation is a proxy for the discovery channel, not a measured discovery process; this is a selection probe, not an identifying control.

S3 Instrument robustness and balance

Table S15 shows that the foreign-disaster IV result is not an artifact of obvious event-level controls. Adding $\log(1 + \text{arrests})$ and a second-Trump-era indicator leaves the estimate close to the headline foreign-disaster estimate. Adding event-publicity and broad-sector controls attenuates the estimate

but preserves the sign and statistical strength. The overidentified specification with foreign disasters and out-of-state U.S. disasters also leads to the same conclusion.

Table S15: Competing-news IV robustness: foreign disasters and controls

	(1) Foreign	(2) + sev./era	(3) + event ctrls	(4) Both IVs
Coverage coef.	−1.059*** (0.338)	−1.123*** (0.341)	−0.993** (0.475)	−1.052*** (0.340)
<i>t</i> -statistic	−3.13	−3.29	−2.09	−3.09
First-stage <i>F</i>	52.66	77.83	24.80	25.78
Sargan over-id (<i>p</i>)	—	—	—	0.830
Foreign-disaster IV	Yes	Yes	Yes	Yes
US-disaster IV	No	No	No	Yes
log(1 + arrests) + 2025+	No	Yes	Yes	No
Publicity + sector ctrls	No	No	Yes	No
<i>N</i>	51	49	49	51

Notes: The dependent variable is per-event β_{late}^c and the endogenous regressor is the Mediacloud Spanish-ICE news coverage spike. Columns (1)–(3) use the foreign-disaster instrument alone. Column (4) uses the foreign-disaster and US-disaster instruments jointly. Severity/era controls are log(1 + arrests) and an indicator for events in 2025 or later. Event controls add the event-publicity flag and broad sector indicators for agriculture/dairy, restaurant/food, manufacturing/industrial, car wash, and construction. First-stage statistics are robust Wald *F* tests for the excluded instruments. Heteroskedasticity-robust standard errors in parentheses. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table S16 reports balance against operation size and pre-event outcomes. The foreign-disaster measure is unrelated to the pre-event coefficient and the backward placebo response. It is related to reported arrests and to the cross-state U.S.-disaster measure, which motivates the specifications with arrests, event controls, and both instruments.

Table S16: Balance check: foreign disaster news shocks and event characteristics

Covariate	<i>N</i>	Coef. on foreign disasters (cond.)	<i>p</i>
Log arrests	49	+0.296	0.034
Pre-event beta	51	−0.002	0.624
Backward placebo beta	51	+0.004	0.674
Cross-state disaster days (std.)	51	+0.486	0.003
Joint test of listed covariates	49	$F = 4.35$	0.006

Notes: Each row reports the coefficient on standardized foreign-disaster days in a regression of the listed event characteristic or pre-event outcome on foreign-disaster days, second-Trump-era, event-publicity, and broad-sector controls. The foreign-disaster measure counts days in the ICE event’s post-event window that fall within the first seven days of a GDACS red-alert sudden disaster outside the United States, Latin America, and the Caribbean. The joint test reverses the regression and asks whether the listed covariates jointly predict foreign-disaster days after the same era, publicity, and sector controls. Heteroskedasticity-robust inference is used. The table is a balance diagnostic, not a proof of exclusion.

Table S17 unpacks the joint balance test in Table S16. The rejection is driven almost entirely by the mechanical correlation between foreign-disaster timing and the cross-state U.S.-disaster measure, which is the second instrument rather than a confounder of foot traffic. The pre-event

coefficient and backward placebo are each individually unrelated to the instrument. The foreign-disaster second-stage estimate is stable when these measures and reported arrests are added as controls, moving from -1.06 to -1.04 . The balance result therefore reflects a controllable timing correlation rather than a channel that drives the estimate.

Table S17: IV balance: what drives the joint test, and conditional-IV survival

<i>Panel A: Loadings of balance covariates on the instrument (per SD)</i>				
	Coef	<i>t</i>	<i>p</i>	
Log arrests	+0.237	+1.26	0.208	
Pre-event β	-0.129	-1.14	0.255	
Backward placebo β	-0.081	-0.47	0.637	
Cross-state disaster days	+0.338	+3.05	0.002	
<i>Joint F (all four)</i>	4.35		0.006	
<i>Panel B: Foreign-disaster IV second stage, conditioning on the imbalance covariates</i>				
	Coef	(SE)	First-stage <i>F</i>	<i>N</i>
No controls (baseline)	-1.059	(0.338)	52.7	51
+ imbalance covariates	-1.056	(0.324)	62.0	49
+ imbalance, era, publicity	-1.040	(0.359)	92.4	49

Notes: Panel A regresses the standardized foreign-disaster instrument on the four covariates in Table S16, with the same era, publicity, and broad-sector controls. Coefficients are per standard deviation of each covariate. The joint F identifies which covariate carries the rejection. Panel B re-fits the foreign-disaster 2SLS estimate of the effect of coverage on foot traffic, adding reported arrests and the two pre-event outcomes as controls. The cross-state disaster measure is the second instrument rather than a control. The second-stage coefficient is stable across specifications, so the balance result reflects a controllable mechanical correlation rather than a violation that drives the estimate. The conditional specifications drop two events with missing arrest counts ($N = 49$). Heteroskedasticity-robust inference throughout.

Table S18 restricts the foreign-disaster instrument by world region and drops disasters in countries with the largest non-Latino diaspora populations in the United States. The estimate remains negative in these exercises. The table also reports the overidentified specification using both disaster instruments.

Table S18: Exclusion-restriction pressure tests for the foreign-disaster IV

	Coef.	(SE)	First-stage F	Disasters
<i>Panel A: Geographic restriction of the instrument</i>				
Baseline (ex-LatAm/Caribbean)	-1.059***	(0.338)	52.7	113
Asia-Pacific only	-1.240***	(0.402)	46.6	83
Africa / Europe / MENA only	-0.520	(0.857)	5.6	27
Drop largest US (non-Latino) diaspora origins	-0.889**	(0.439)	25.8	69
<i>Panel B: Overidentification and balance</i>				
Foreign + US disasters (over-id)	-1.052***	(0.340)	25.8	
Sargan over-id p	0.830			
Balance joint F (covariates $\rightarrow$ instrument)	$F = 4.35, p = 0.006$			
OLS reference	-0.919***	(0.256)		

Notes: The endogenous regressor is the Mediacloud Spanish-ICE news coverage spike and the outcome is per-event β_{late}^e . Panel A re-estimates the foreign-disaster IV restricting the instrument to disasters by region of the affected country. Of the 113 qualifying red-alert non-Latin-American sudden disasters, 83 (73%) are in the Asia-Pacific and the remainder elsewhere outside Latin America and the Caribbean, none with a material tie to the Mexican or Central-American-origin population. The diaspora row drops the largest US non-Latino foreign-born origin countries (the Philippines, India, China, Vietnam, Korea). Panel B reports the overidentified specification using the foreign-disaster and US-disaster instruments jointly with its Sargan test, and the joint test that the listed event characteristics and pre-event outcomes predict the instrument. First-stage statistics are robust Wald F tests. Heteroskedasticity-robust standard errors. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

S4 Exposure-measure robustness

The headline high-FB indicator uses the share of residents born in Latin America (ACS B05006) at the CBG level. This is a broad measure because it counts naturalized US citizens alongside non-citizens, all Latin-American national-origin groups alongside the highest-enforcement-risk Mexican and Central American populations, and all adults alongside the US-citizen children of foreign-born parents. To test whether the chilling response is concentrated in narrower subgroups most directly exposed to deportation risk, the per-event distributed-lag DiD is refit with three sharper exposure indicators. These are the share of Hispanic non-citizen adults, the share of residents born in Mexico, El Salvador, Guatemala, Honduras, or Nicaragua, and the share of US-citizen children under 18 with at least one foreign-born parent.

Table S19: Exposure-measure robustness, per-event estimate on Spanish-ICE spike by high-FB indicator

Exposure indicator	N	mean estimate	sd estimate	r vs base	coef	t	R^2
Latin-Amer. FB (baseline, B05006)	51	-0.008	0.047	+1.000	-0.92	-3.59	0.227
Hispanic non-citizen (B05003I)	51	-0.006	0.050	+0.413	-0.17	-0.41	0.006
Mex. + Cent. Amer. FB (B05006)	51	-0.004	0.047	+0.705	-0.58	-1.98	0.089
Mixed-status kids (B05009)	49	+0.001	0.042	+0.516	-0.29	-1.05	0.027

Notes: For each exposure indicator, the event-level robustness design is refit, defining high-FB CBGs as the top-decile and low-FB CBGs as the bottom-decile within each event’s catchment by the indicated variable. “Mean estimate” and “sd estimate” are the cross-event distribution of the resulting late-window estimates. “ r vs base” is the cross-event correlation of the refined-indicator estimates with the baseline estimates. The last three columns report the single-regressor OLS of the late-window estimate on the Mediacloud Spanish-ICE spike with HC1 standard errors.

All three refinements produce the same negative sign as the baseline at smaller mean magnitudes and substantially weaker cross-event predictive power for the Spanish-ICE news coverage spike. The pattern suggests that the cross-event response operates at the broad community level rather than narrowly at the personal-deportation-risk level. This weakening is not mechanically driven by fewer POIs in the sharper top-decile cells. Top-decile POI density under the sharper indicators is within ± 6 percent of the baseline.

S5 Venue-coverage and civic-mechanism checks

Advan’s POI panel does not measurably capture two venue categories that are central to the sensitive-locations policy and to the broader immigration-fear literature. Religious organizations (NAICS-3 = 813) contain 4,025 POIs across the 26-state catchment universe, of which 3,817 (95 percent) are mislabeled community-bank branches in the underlying classification. After stripping obvious mislabels, only a very small number of panel POIs correspond to real religious congregations. Healthcare venues are absent. NAICS-2 = 62 has zero POIs in the panel. The civic-versus-commercial decomposition therefore captures the schools and parks/museums response but is silent on churches and hospitals.

Table S20 reports the cross-event inference behind the civic-versus-commercial contrast in the main paper.

Table S20: Civic vs. commercial response, with inference

	<i>N</i>	Mean	(SE)	<i>t</i>
Civic estimate	12	−0.120**	(0.045)	−2.67
Commercial estimate	12	−0.010	(0.009)	−1.03
Civic − commercial	12	−0.111**	(0.040)	−2.79
High-coverage events	6	−0.060	(0.044)	
Low-coverage events	6	−0.162	(0.063)	

Notes: Sample is the 12 events with both a civic (NAICS-3 712 and 611: parks, museums, educational services) and a commercial (NAICS-2 44, 45, 72: retail and food services) per-event estimate. Each row reports a one-sample *t*-test across events of the indicated per-event estimate; the differential is computed within event before averaging. Inference is at the cross-event level and does not down-weight noisier event estimates. The coverage split uses the Mediacloud Spanish-ICE spike. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

I also test whether news coverage moderates the civic-vs.-commercial differential. With the English-GDELT spike, the within-event civic-minus-commercial differential is −0.153 in the low-coverage group and −0.051 in the high-coverage group. With the Mediacloud Spanish-ICE spike, the pattern is similar. The differential is −0.162 in the low-coverage group and −0.060 in the high-coverage group. The civic-vs.-commercial gap is wider in low-coverage events than in high-coverage events under both measures, the opposite of what a coverage-moderated mechanism would predict.

S6 Spending

Table S21 reports the card-spending response by news coverage. Per-event spending estimates come from the exact-geocoded SafeGraph SpendPatterns subsample. The 30 main-sample events with both a spending estimate and a news coverage spike enter this comparison, and the estimates are much noisier than the foot-traffic estimates. The cross-event aggregate is not distinguishable from zero. Descriptively, above-median-coverage events show a differential decline of about 3 percent while below-median-coverage events show a small increase of about 4 percent, but the cross-event slope of the spending estimate on the news coverage spike is not statistically significant, so the spending evidence is suggestive rather than conclusive.

Table S21: Spending response by news coverage

	N	Mean β_{late}	(SE)	Implied % decline
<i>Panel A: Cross-event aggregate</i>				
All events	30	+0.0069	(0.0314)	−0.7
<i>Panel B: By median Spanish-ICE coverage</i>				
Above-median coverage	15	−0.0278	(0.0285)	+2.7
Below-median coverage	15	+0.0415	(0.0556)	−4.2
<i>Panel C: Cross-event slope of β_{late} on the standardized spike (per SD)</i>				
Spanish-ICE spike	30	−0.0284	(0.0329)	
English-GDELT spike	30	−0.0285	(0.0339)	

Notes: Per-event late-window spending estimates β_{late}^e are estimated with the event-level robustness design on $\log(1 + \text{weekly spending})$ over the exact-geocoded SafeGraph SpendPatterns subsample, the 30 main-sample events with both a spend estimate and a Spanish-ICE news coverage spike. A positive β_{late} means high-FB spending rose relative to low-FB; the implied percent decline is $100(1 - \exp(\text{mean } \beta))$, so a negative mean β prints as a positive decline. Panel A reports the cross-event mean and its standard error; the aggregate is not distinguishable from zero. Panel B splits events at the median Spanish-ICE spike. Panel C regresses β_{late}^e on the standardized news coverage spike with HC1 standard errors; coefficients are per standard deviation of the spike. Per-event spend estimates are much noisier than for foot traffic, so these are order-of-magnitude patterns rather than precise effects. Stars: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.